



Buffer Overflow Vulnerabilities, Attacks, and Defenses



Software Security

Buffer Overflow Vulnerabilities, Attacks, and Defenses (Part I)

Professor Travis Peters
CSCI 476 - Computer Security
Spring 2020

*Some slides and figures adapted from Wenliang (Kevin) Du's
Computer & Internet Security: A Hands-on Approach (2nd Edition).
Thank you Kevin and all of the others that have contributed to the SEED resources!*

Today

Announcements

- **Reminder:** Lab 02 due Thursday (2/6) @ 3pm
- **Question:** Buffer Overflow CTF **or** Format String Vulns.
- **REU** (Research Experiences for Undergraduates)
 - **@ MSU:** Cybersecurity Algorithms - <http://www.bobcatsoftwarefactory.com/nsf-reu-2020>
 - **@ NYU:** HW and SW Security for Mobile Devices and Networks (see email sent earlier today)

Goals & Learning Objectives

- Buffer Overflow Vulnerabilities, Attacks, and Defenses
 - Review the layout of a program in memory
 - Understanding the stack layout
 - Vulnerable code
 - Challenges in exploiting buffer overflow vulnerabilities
 - Countermeasures



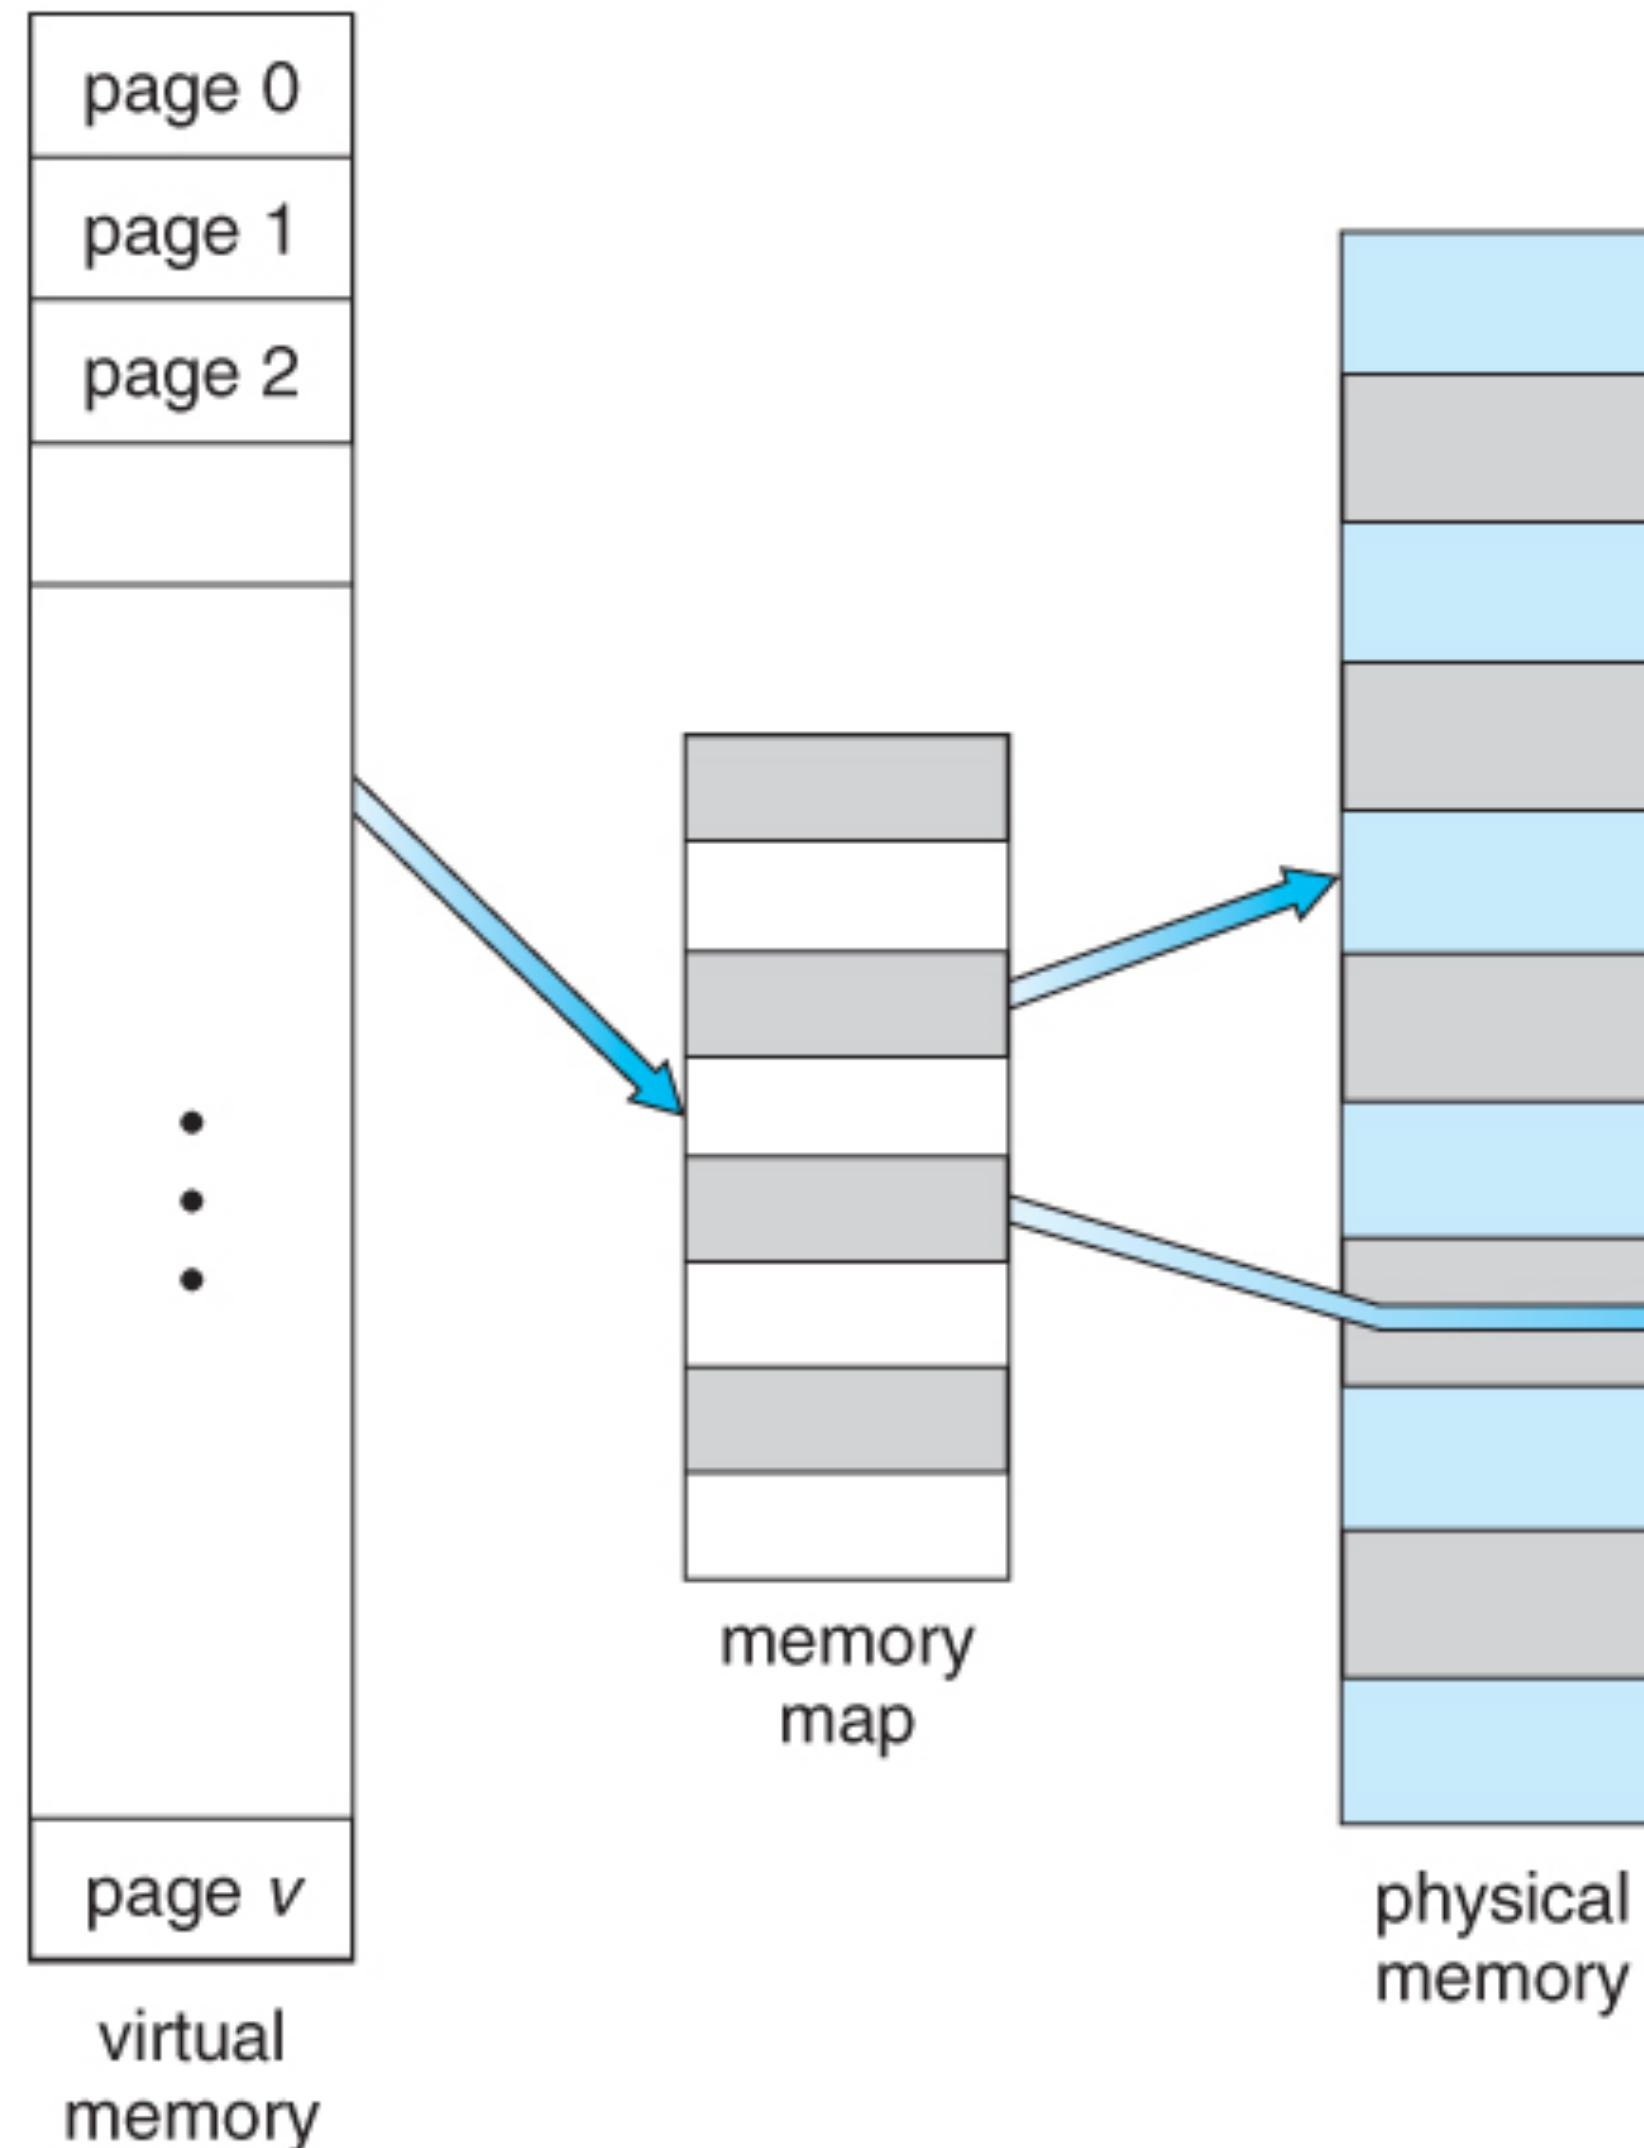


Foundations: Memory Layout, Stacks, and Function Invocation

Virtual Addresses vs. Physical Addresses

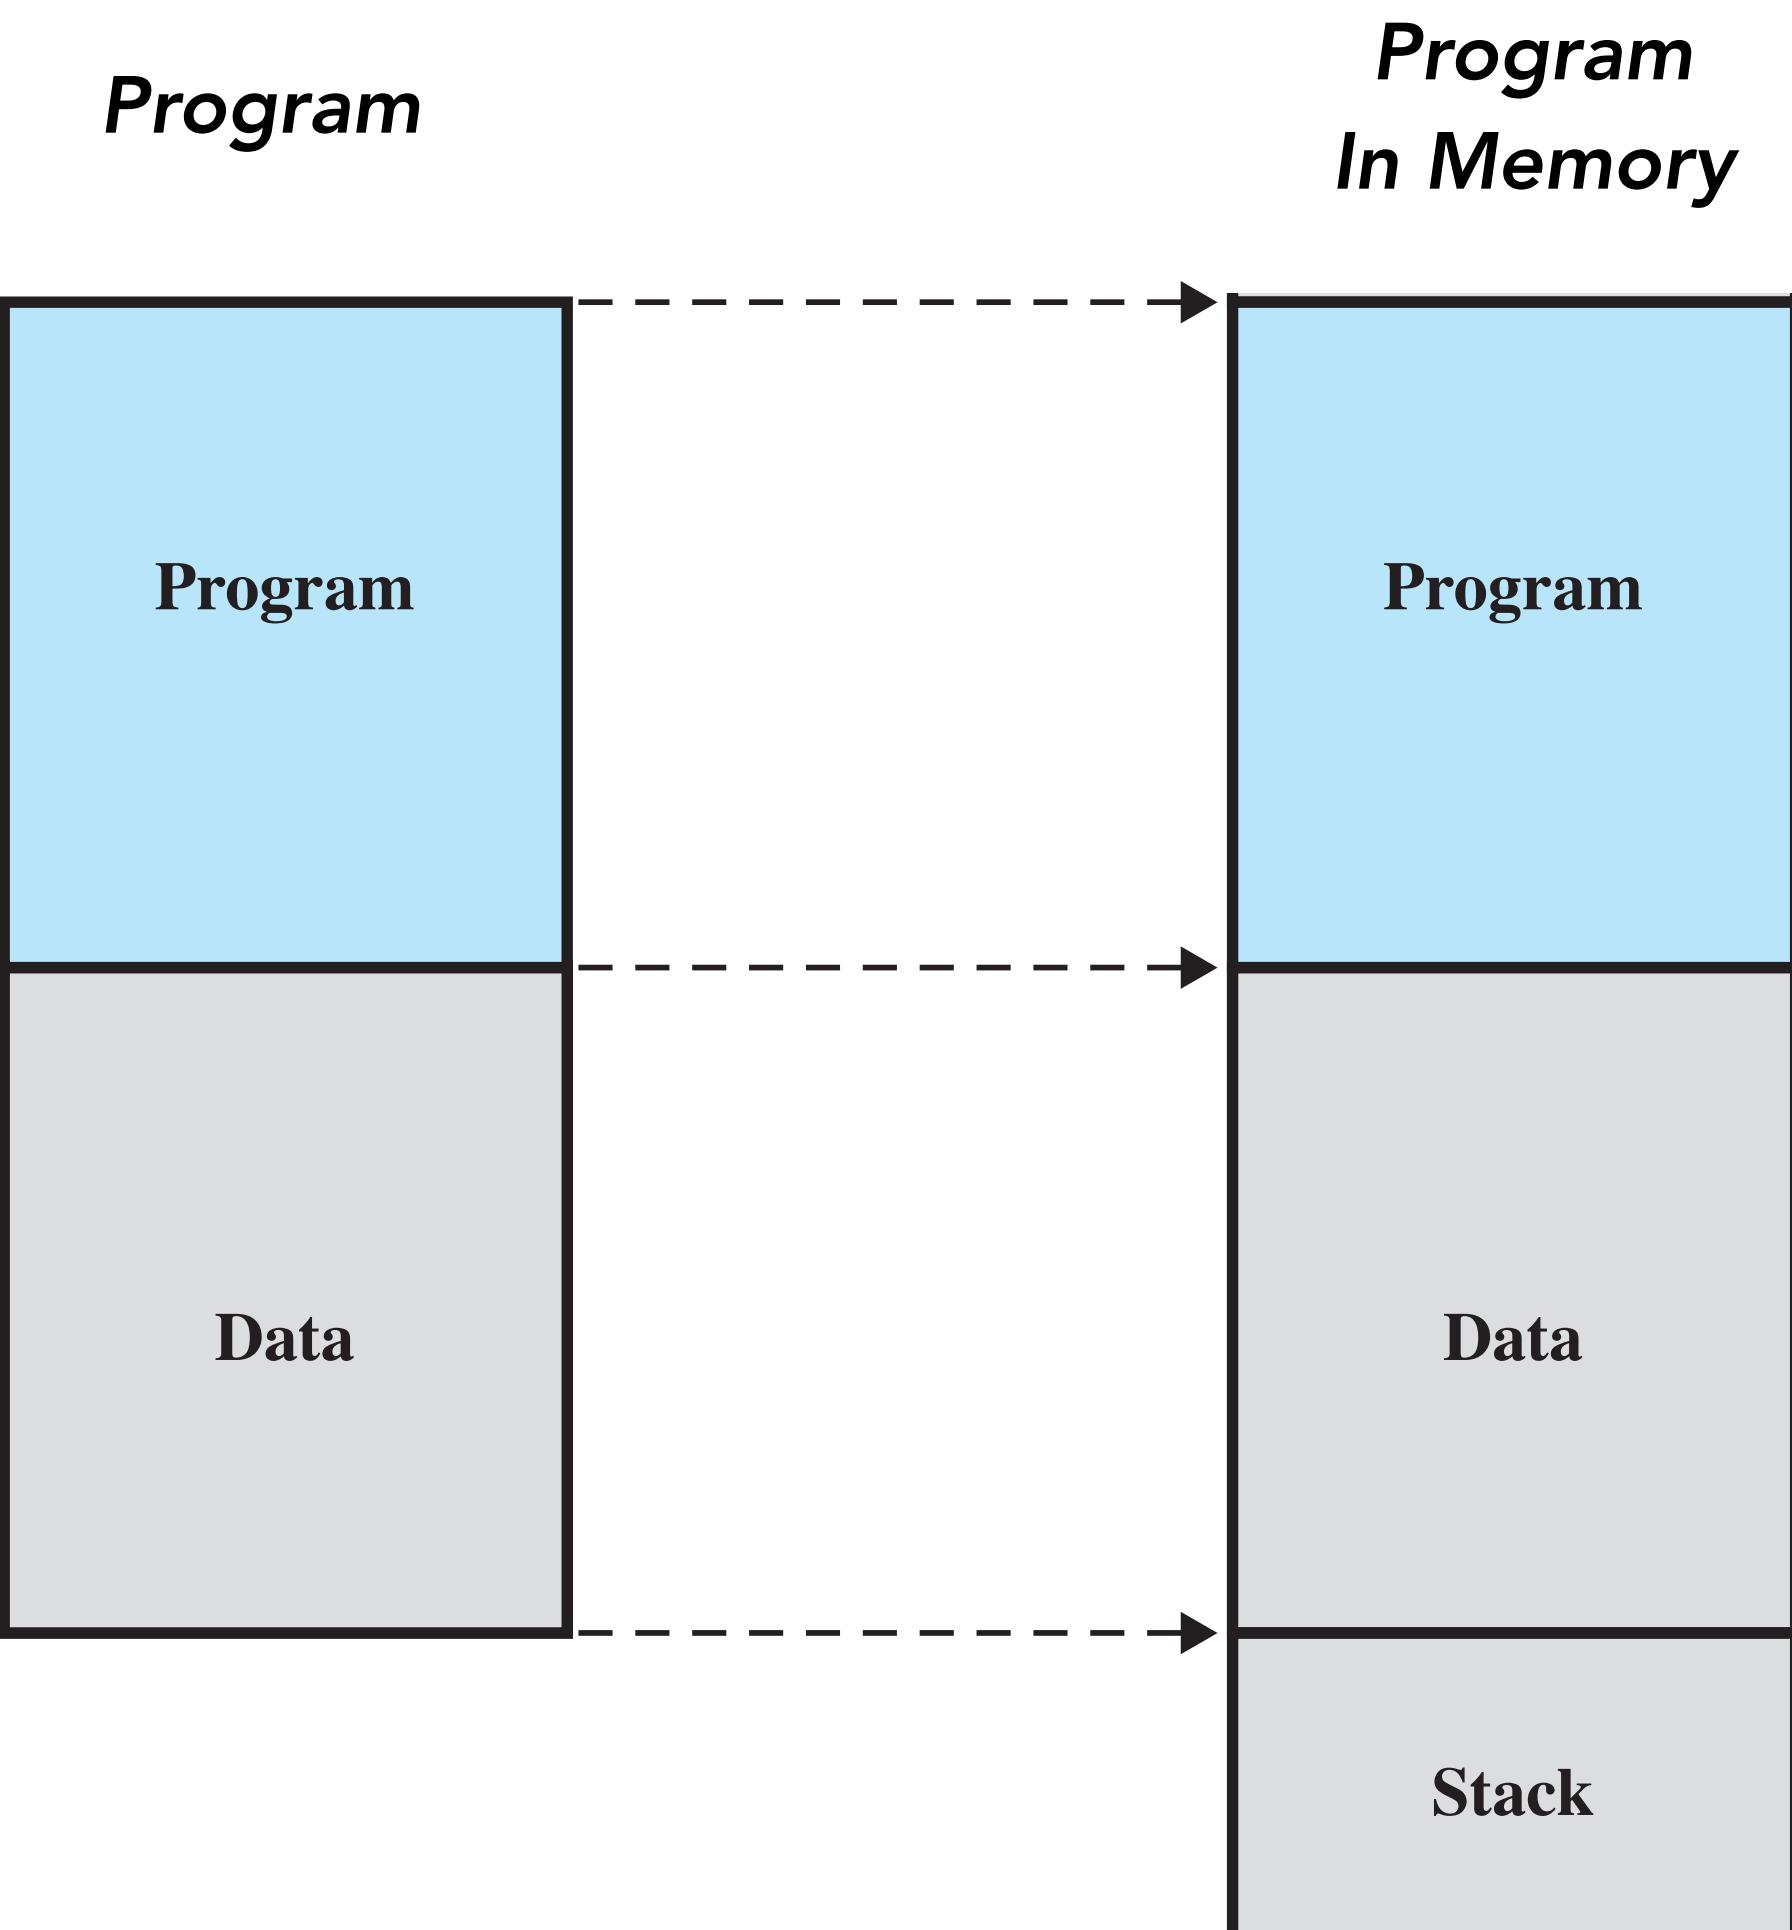
- Programs do not access physical memory directly
 - Process memory is mapped from a virtual address space to the physical address space
- **Virtual Memory**
- OS maintains a table (memory map) for each process

Why?



Loading a Program Into Memory

A standard layout for programs makes ***loading*** and ***managing*** programs easy!



The Memory Layout of a Running Program

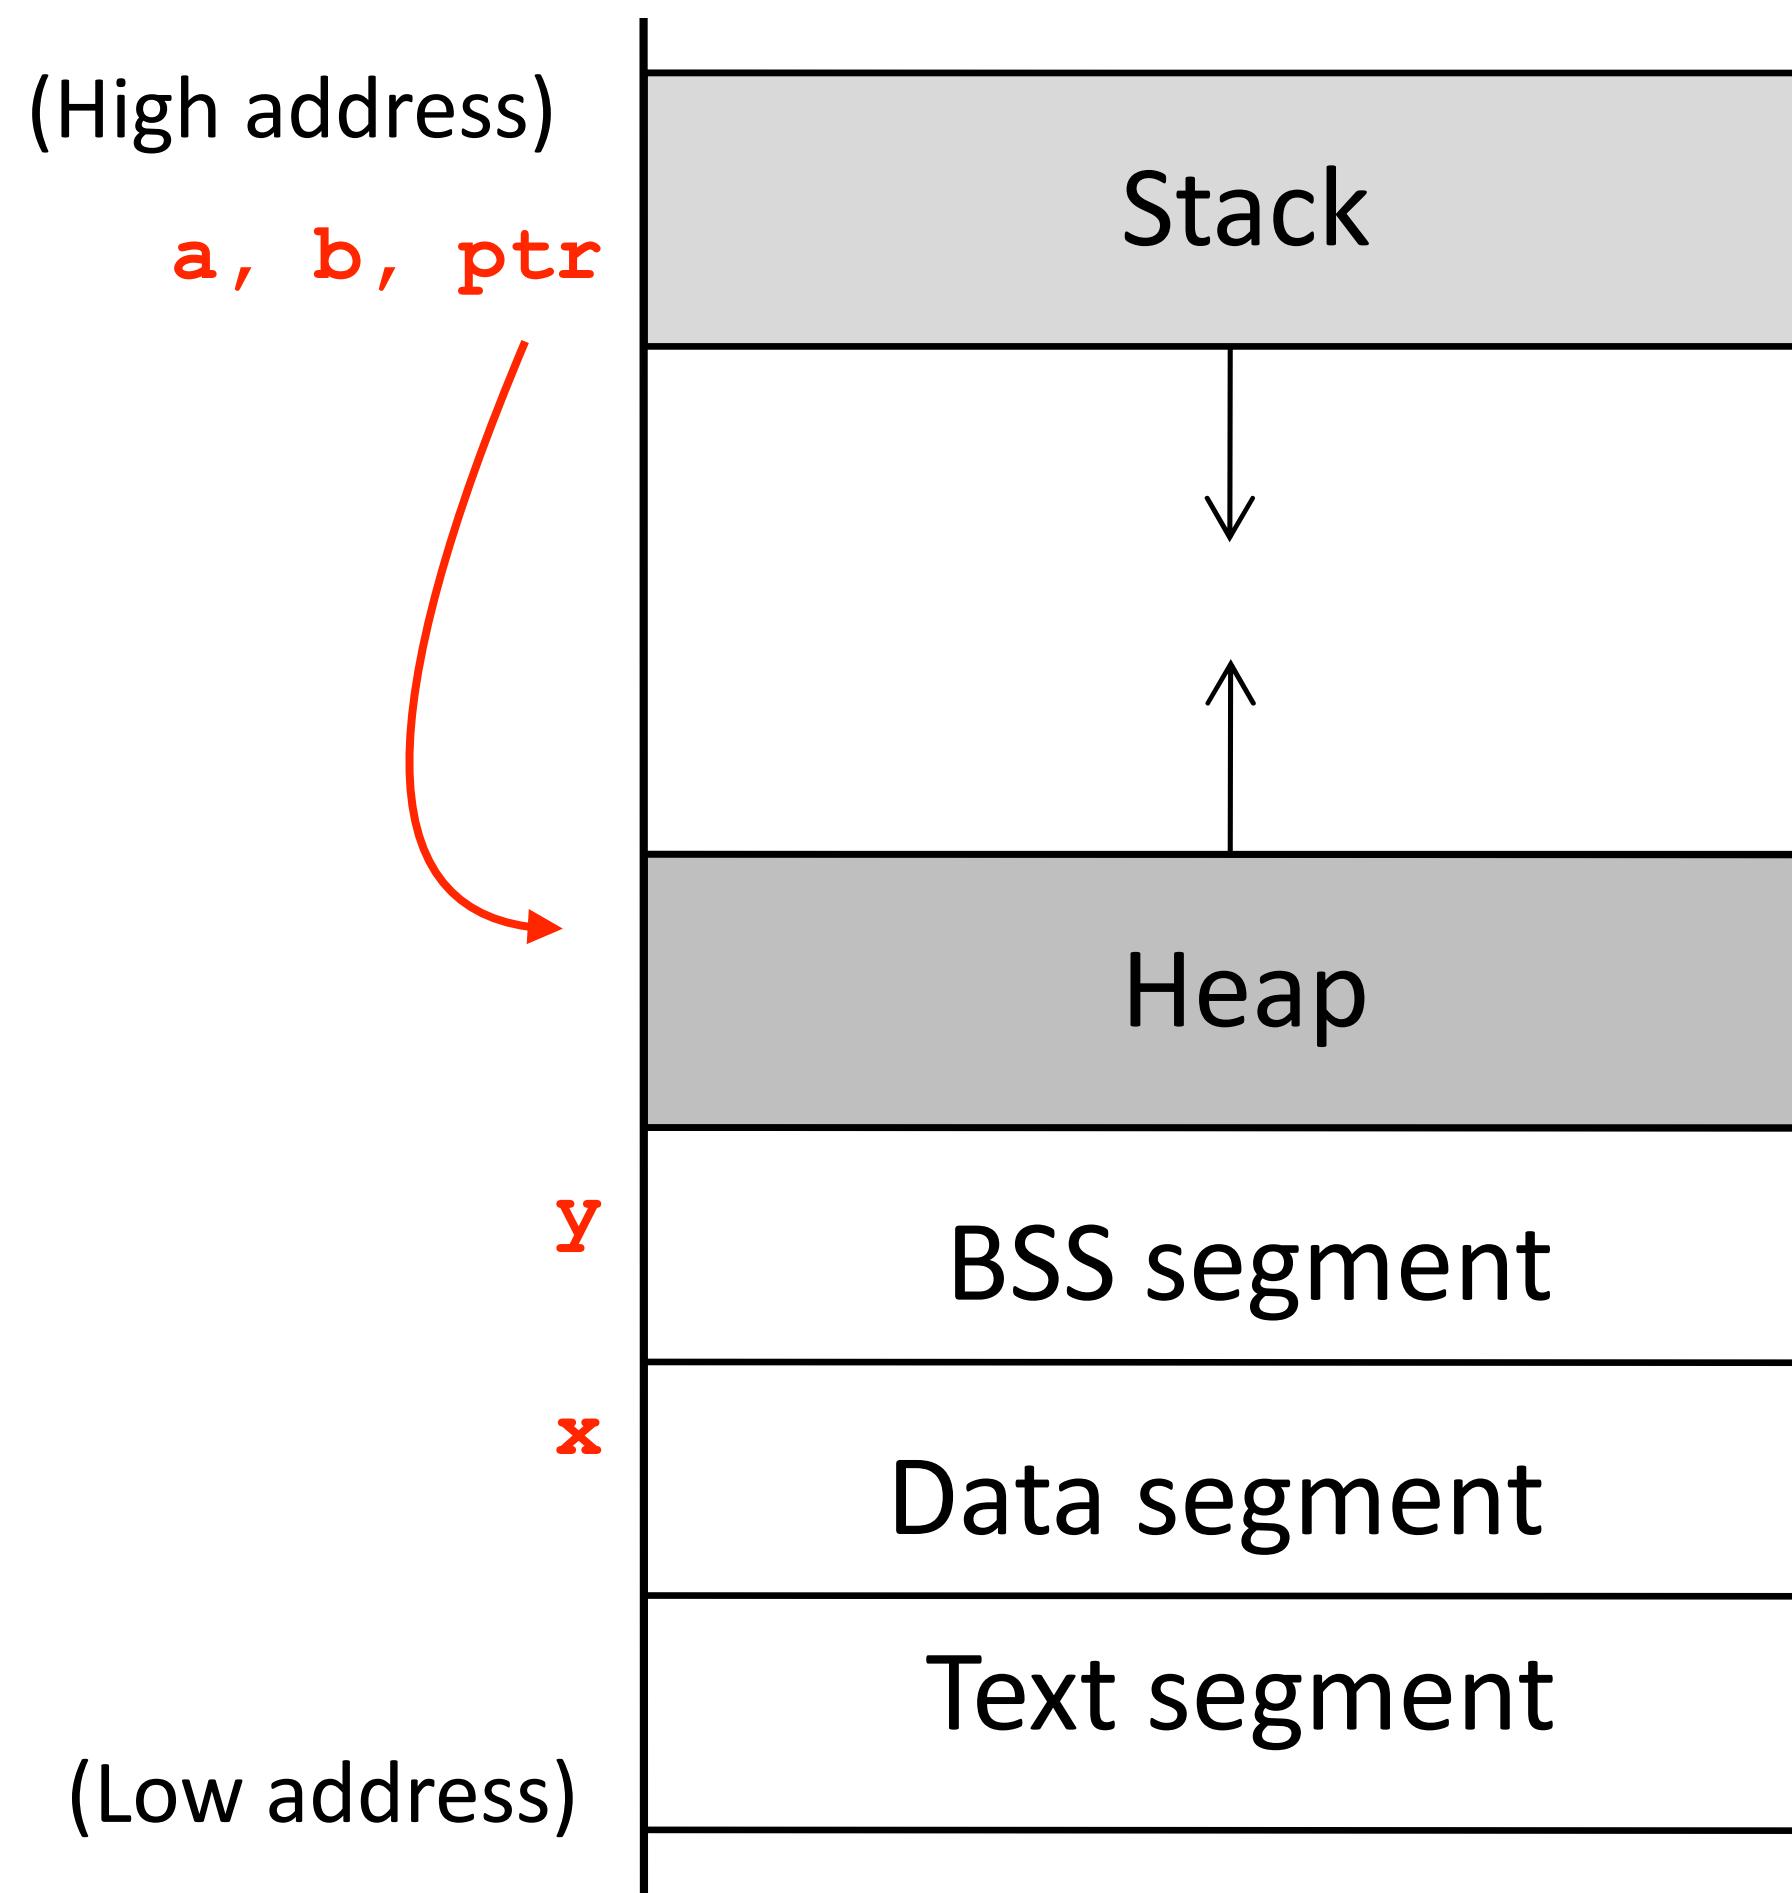
Activity!

Where in memory do the variables go for this program?!

```
int x = 100;
int main()
{
    int a = 2;
    float b = 2.5;
    static int y;

    int *ptr = (int *) malloc(2*sizeof(int));
    ptr[0] = 5;
    ptr[1] = 6;

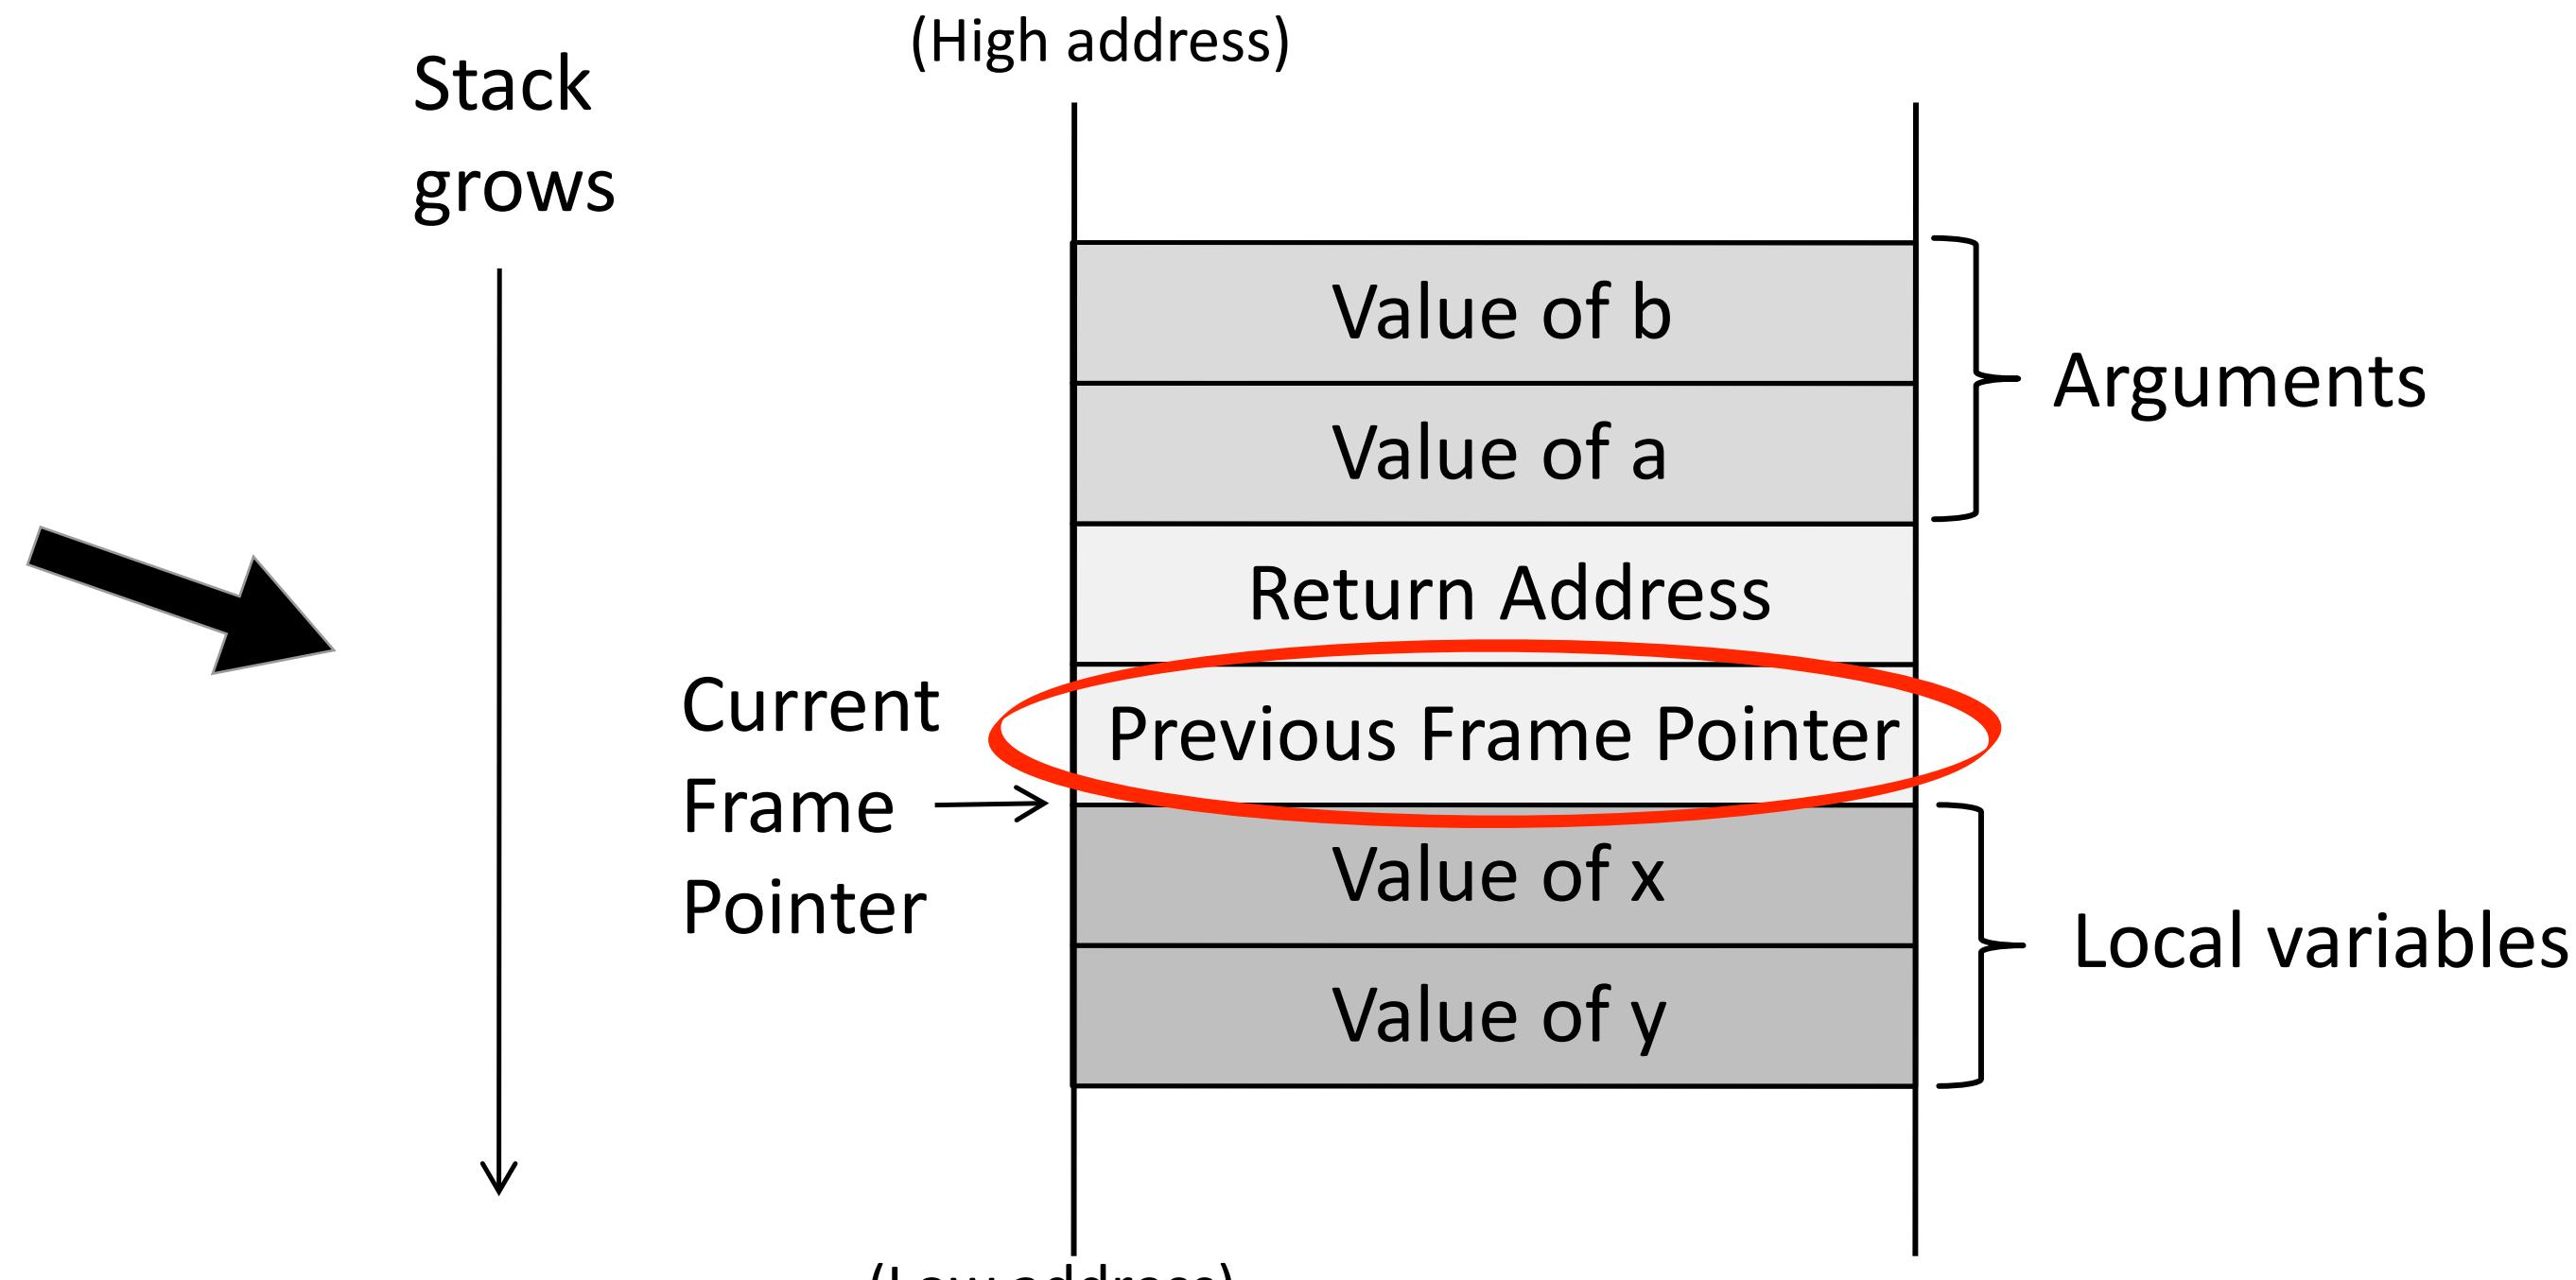
    free(ptr);
    return 1;
}
```



Stack Layout: Function Arguments, Local Variables, ...

```
void func(int a, int b)
{
    int x, y;

    x = a + b;
    y = a - b;
}
```



Solution: Frame Pointer (ebp)!

<code>movl 12(%ebp), %eax</code>	<code>; b is stored in %ebp + 12</code>
<code>movl 8(%ebp), %edx</code>	<code>; a is stored in %ebp + 8</code>
<code>addl %edx, %eax</code>	
<code>movl %eax, -8(%ebp)</code>	<code>; x is stored in %ebp - 8</code>

PROBLEM!

The compiler doesn't know where the stack frame will be in memory...

How do we know the address of function arguments/local variables?

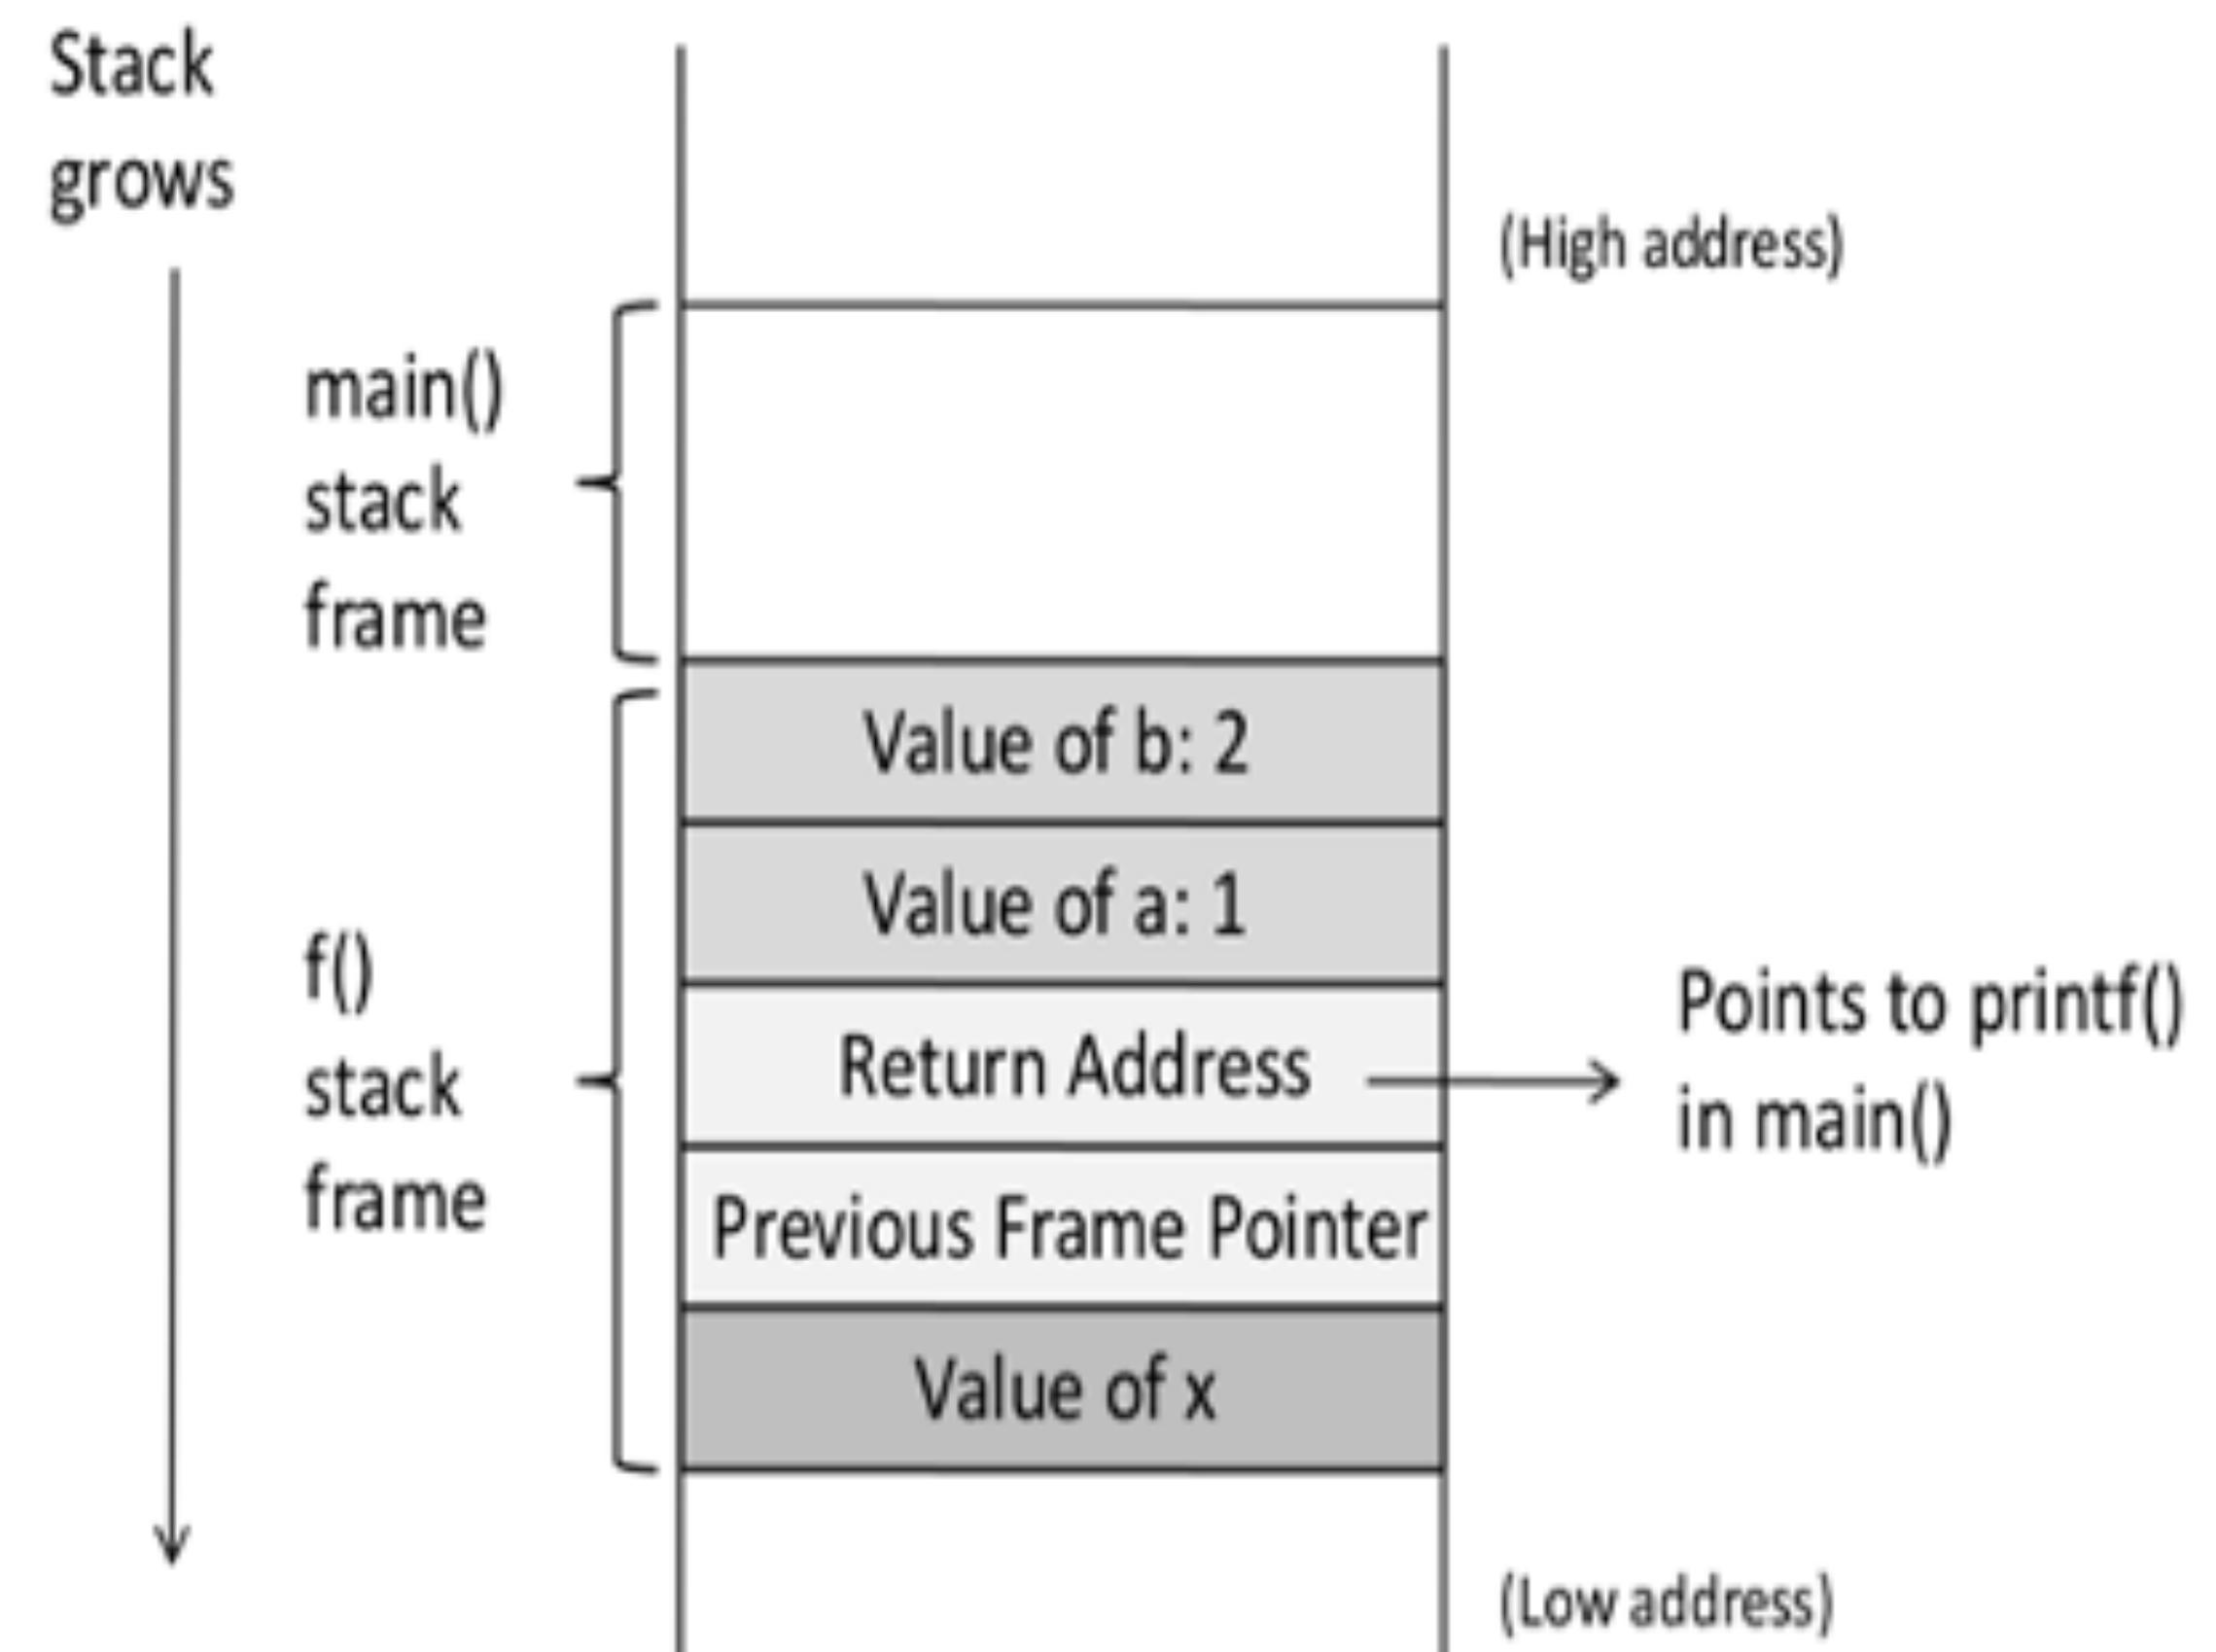
Example: A Function Call Stack

Activity!

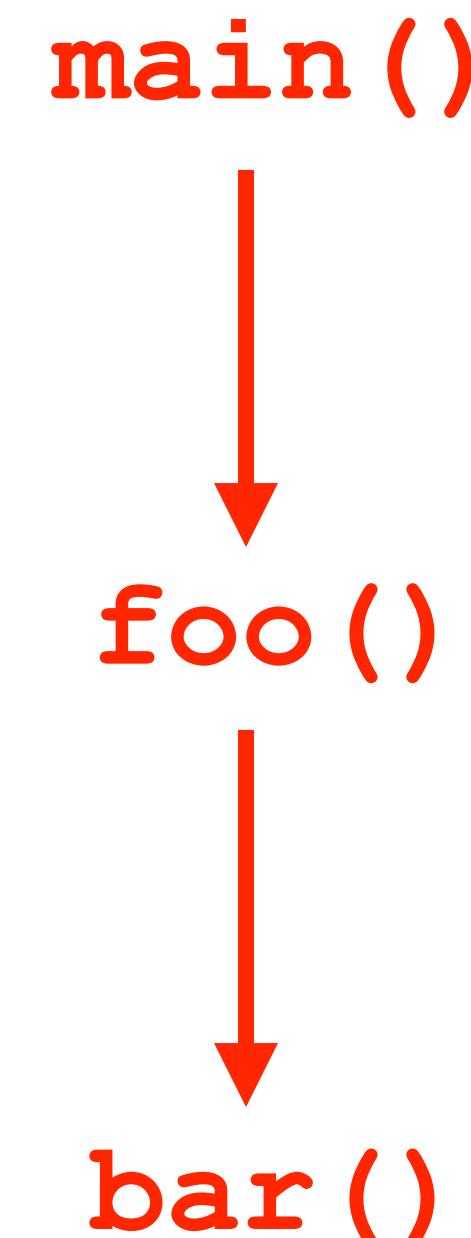
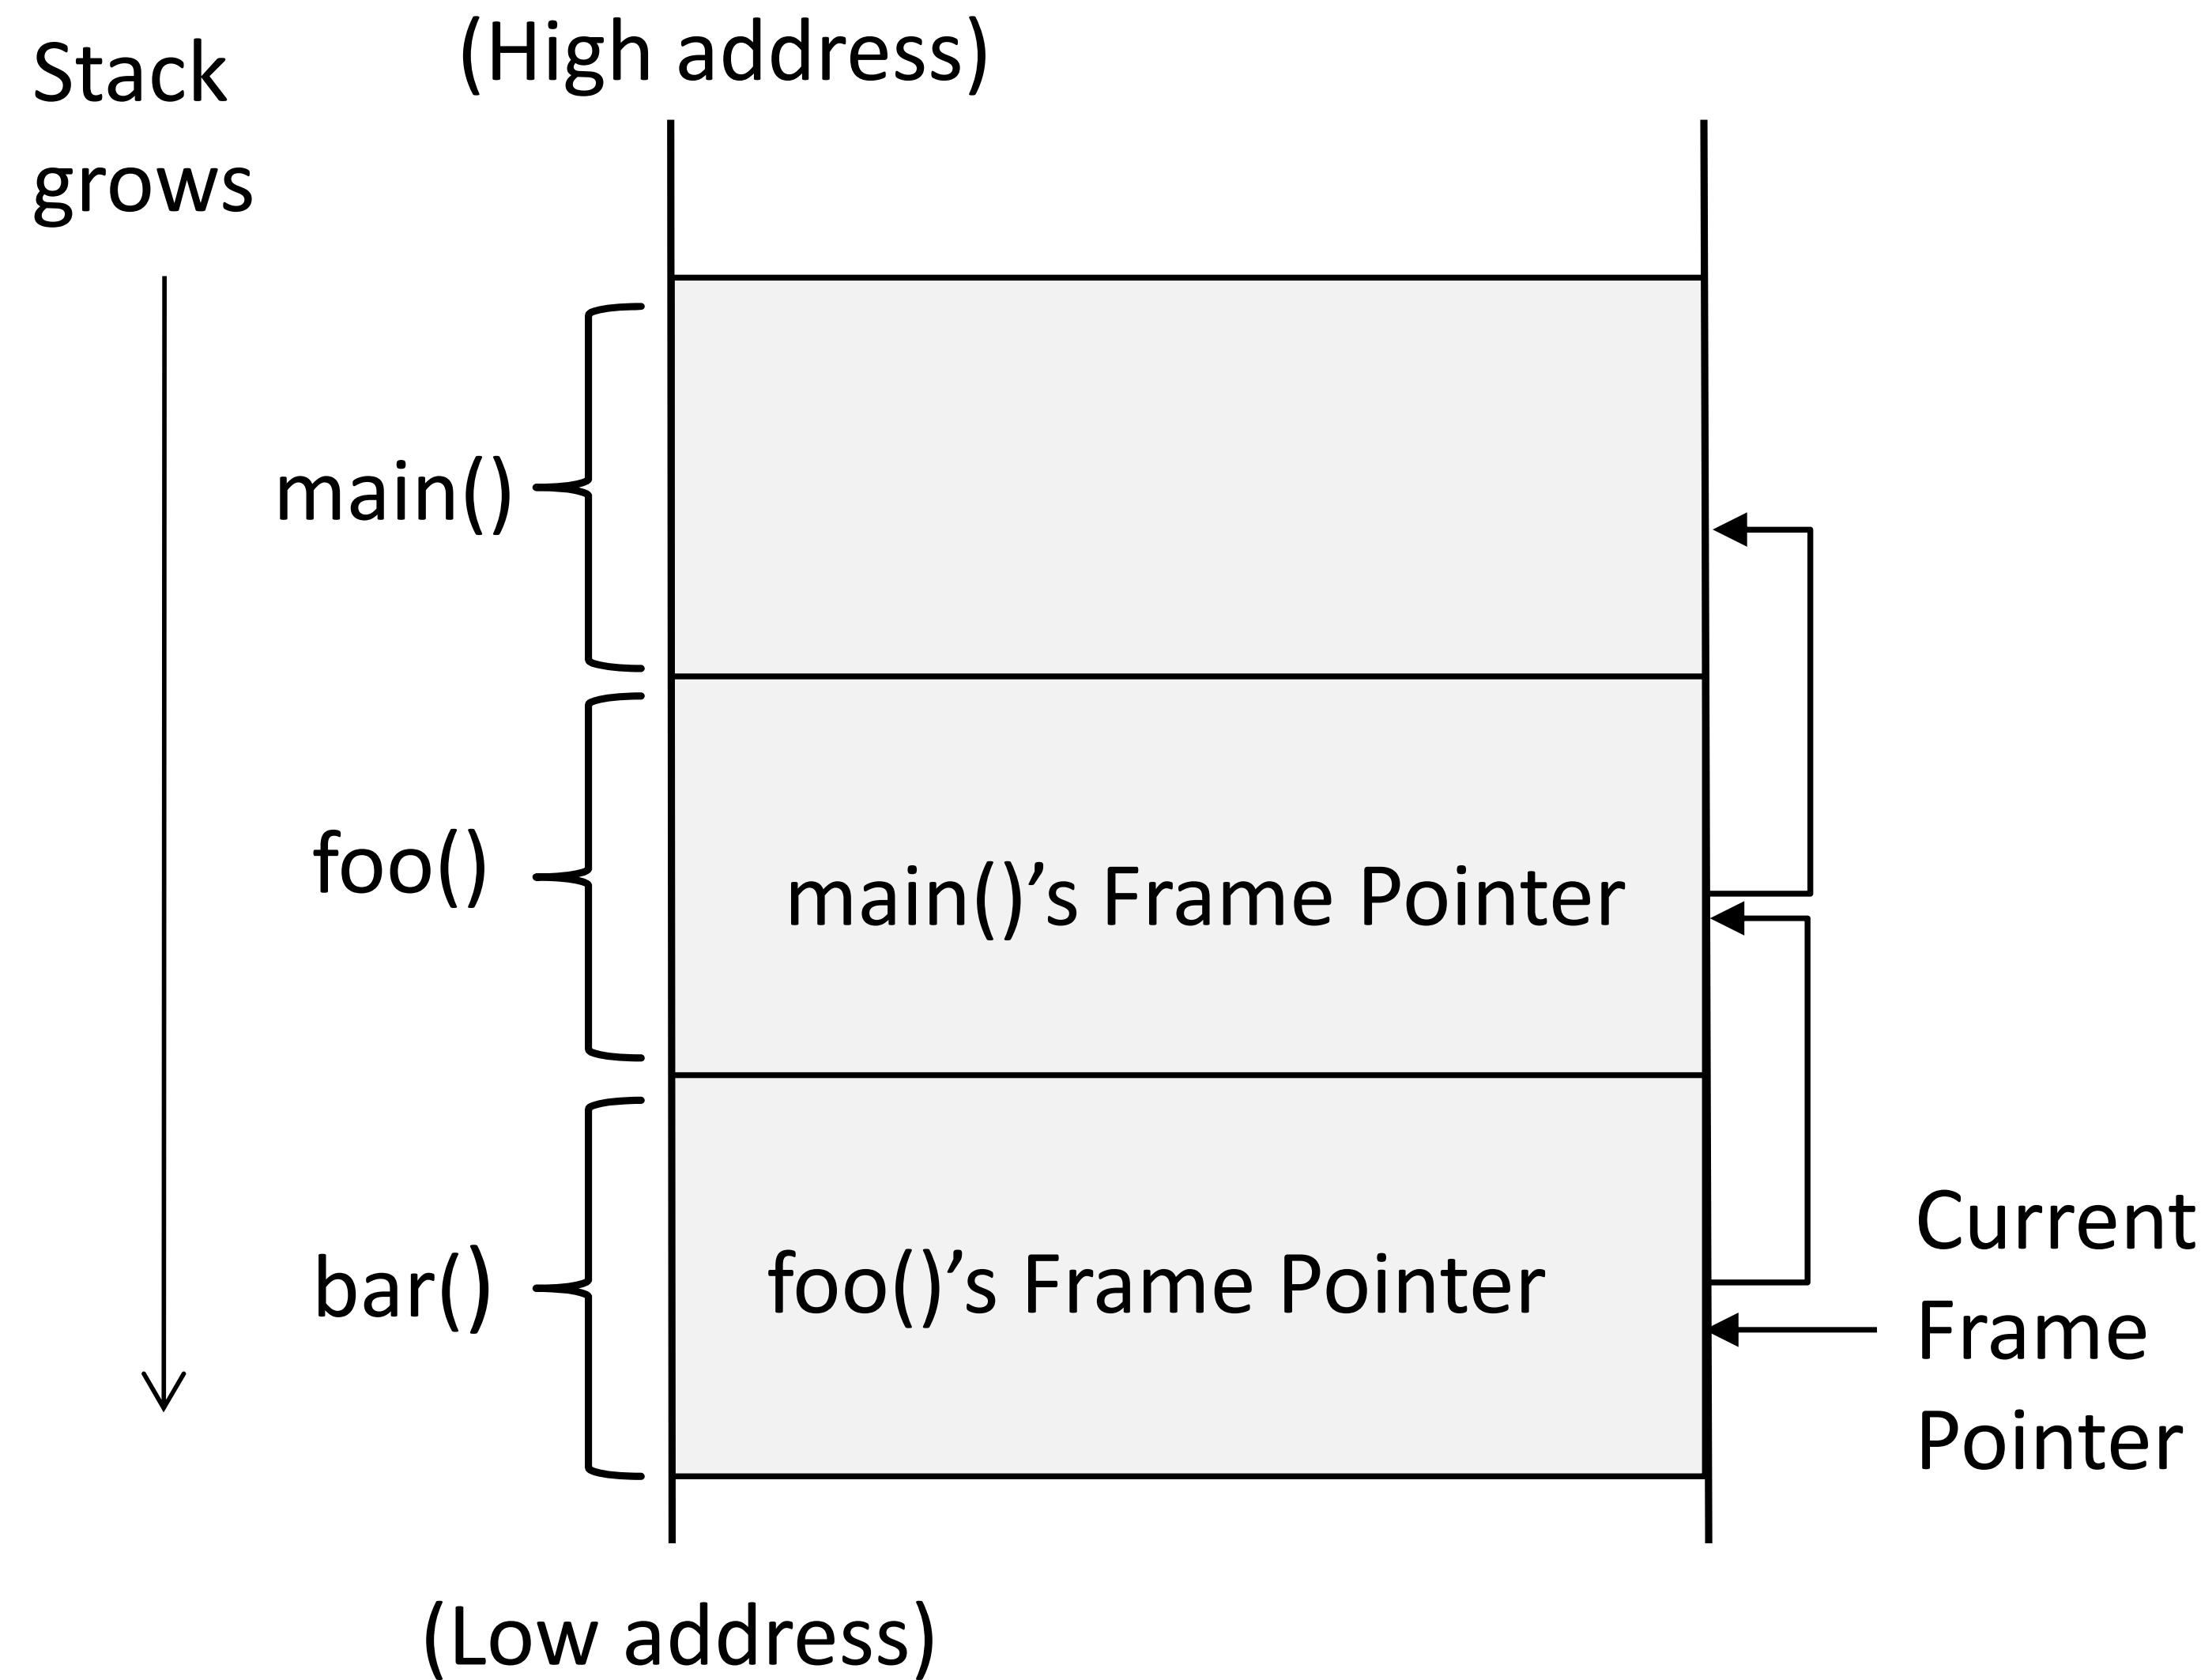
What will foo's stack frame look like?

```
void f(int a, int b)
{
    int x;
}

void main()
{
    f(1, 2);
    printf("hello world");
}
```



Stack Layout and a Function Call Chain



Current
Frame
Pointer

Note

Frame pointer (ebp) points to current frame;
ebp is different for each stack frame.
Need to save previous frame pointer



Now, let's consider an example vulnerable program...

A Vulnerable Program

- Reading 300 bytes of data from badfile.
- Storing the file contents into a str variable of size 400 bytes.
- Calling foo function with str as an argument.

Note: The badfile is created by the user and hence the contents are controlled by the user.

Call chain: main() → foo() → strcpy()

```
/* This program has a buffer overflow vulnerability. */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str)
{
    char buffer[100];

    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);

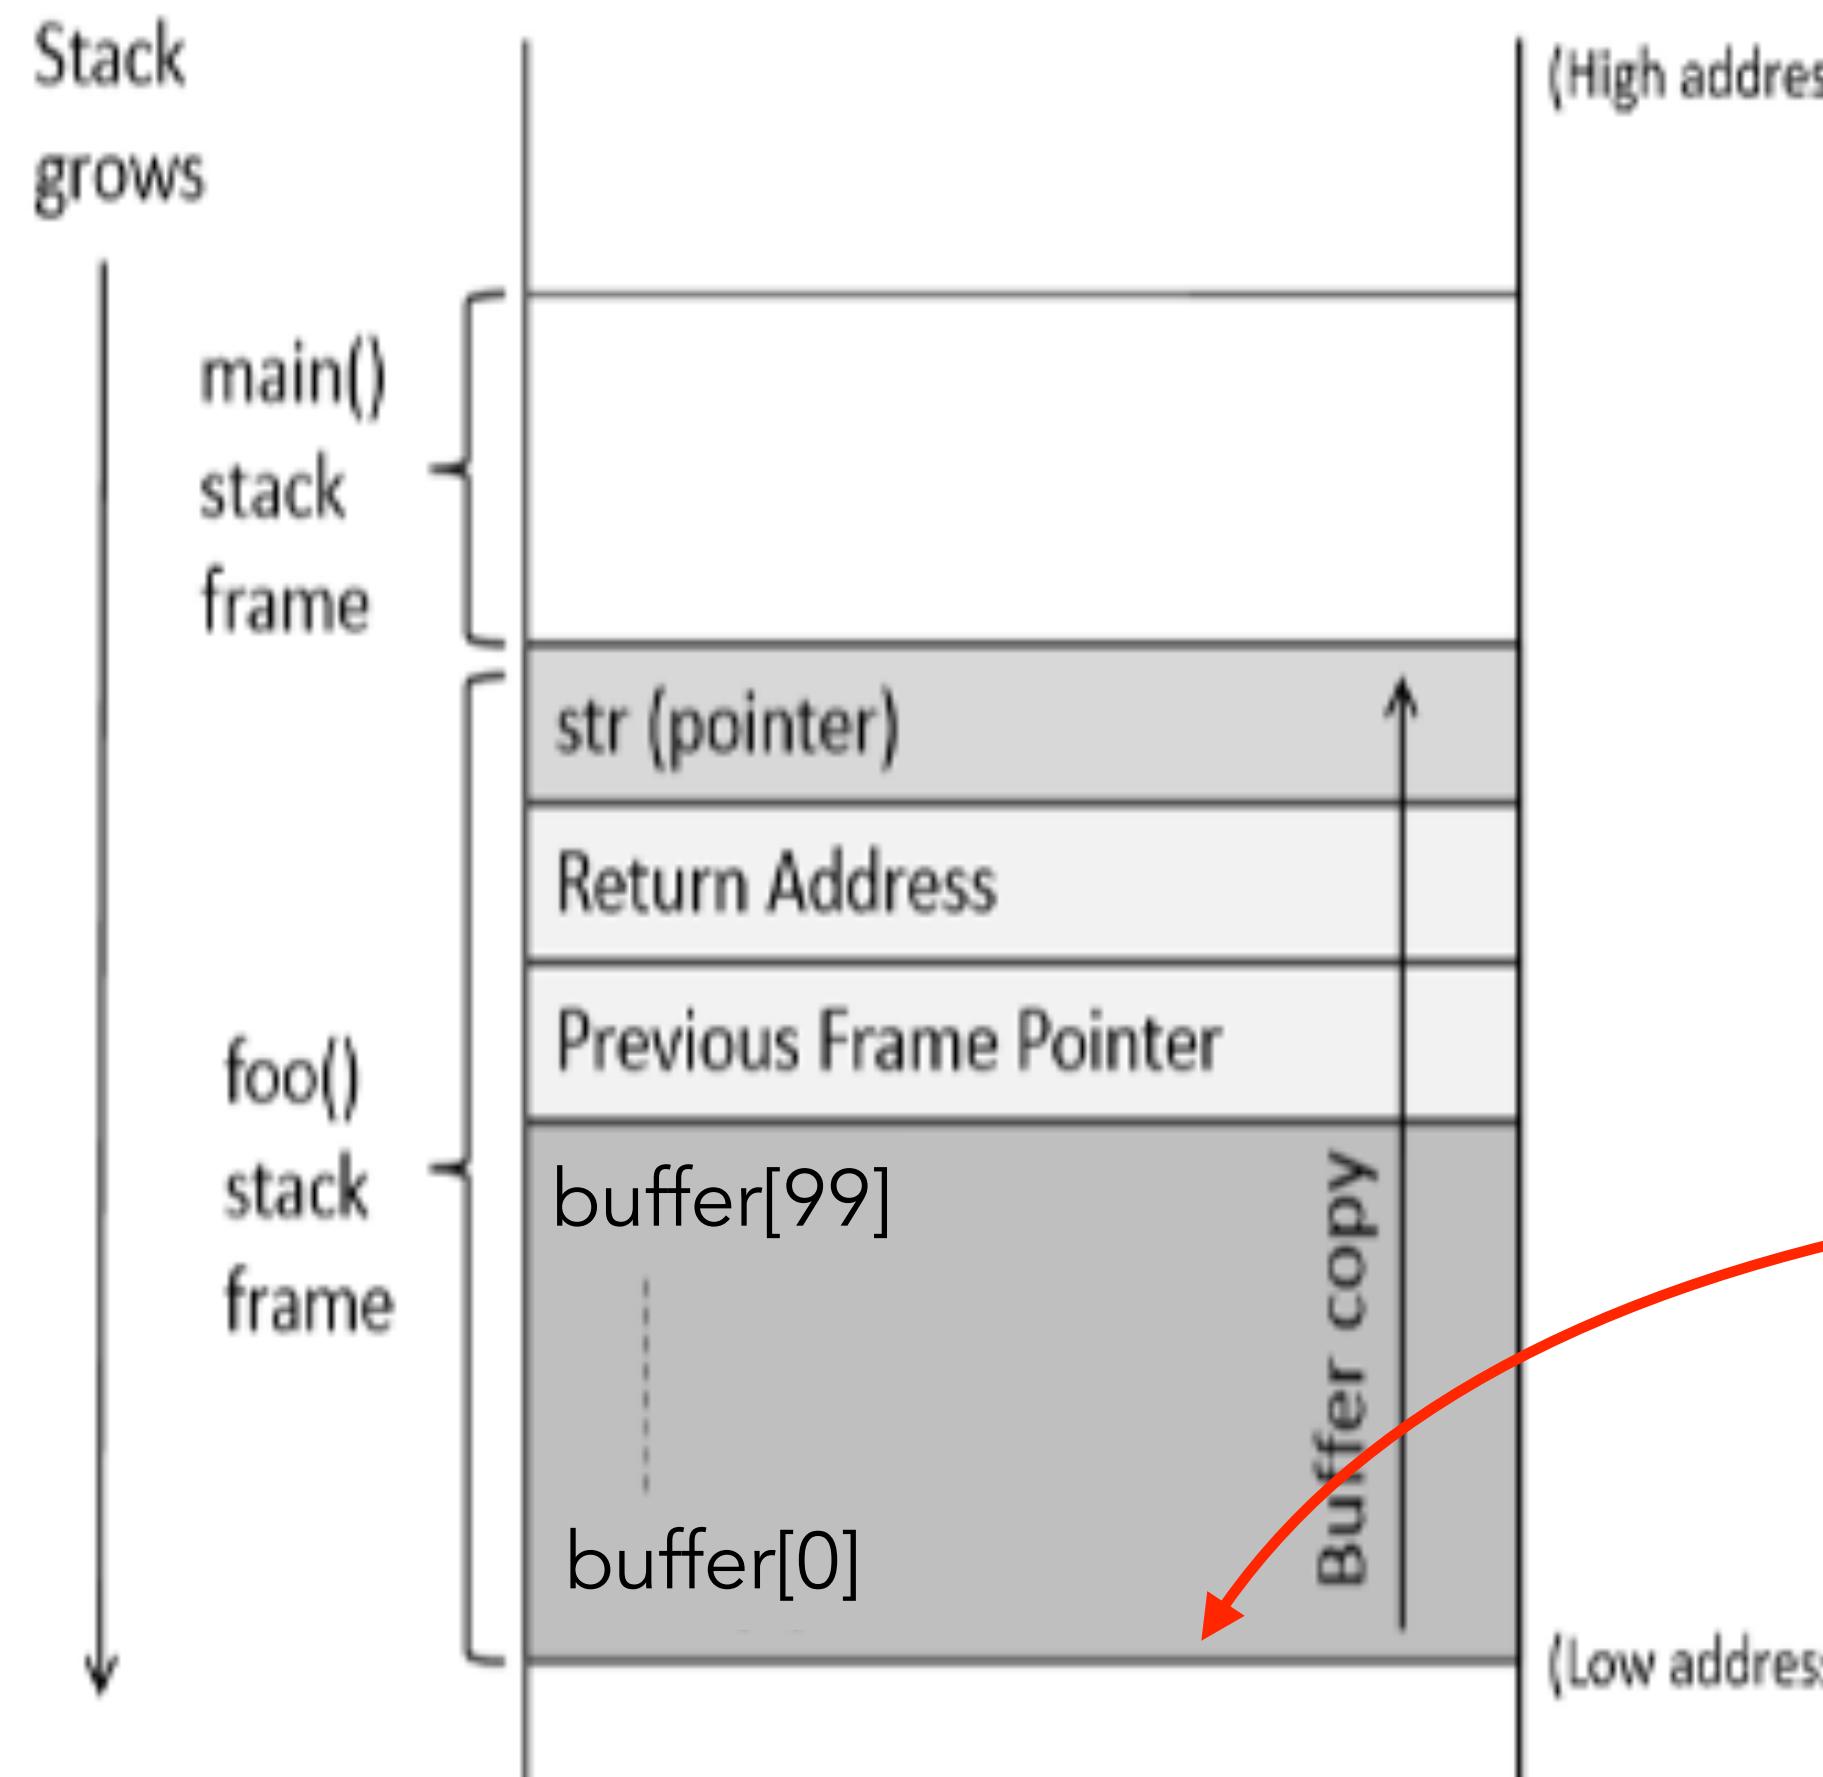
    return 1;
}

int main(int argc, char **argv)
{
    char str[400];
    FILE *badfile;

    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 300, badfile);
    foo(str);

    printf("Returned Properly\n");
    return 1;
}
```

A Vulnerable Program (cont.)



```
/* This program has a buffer overflow vulnerability. */
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str)
{
    char buffer[100];

    /* The following statement has a buffer overflow problem */
    strcpy(buffer, str);

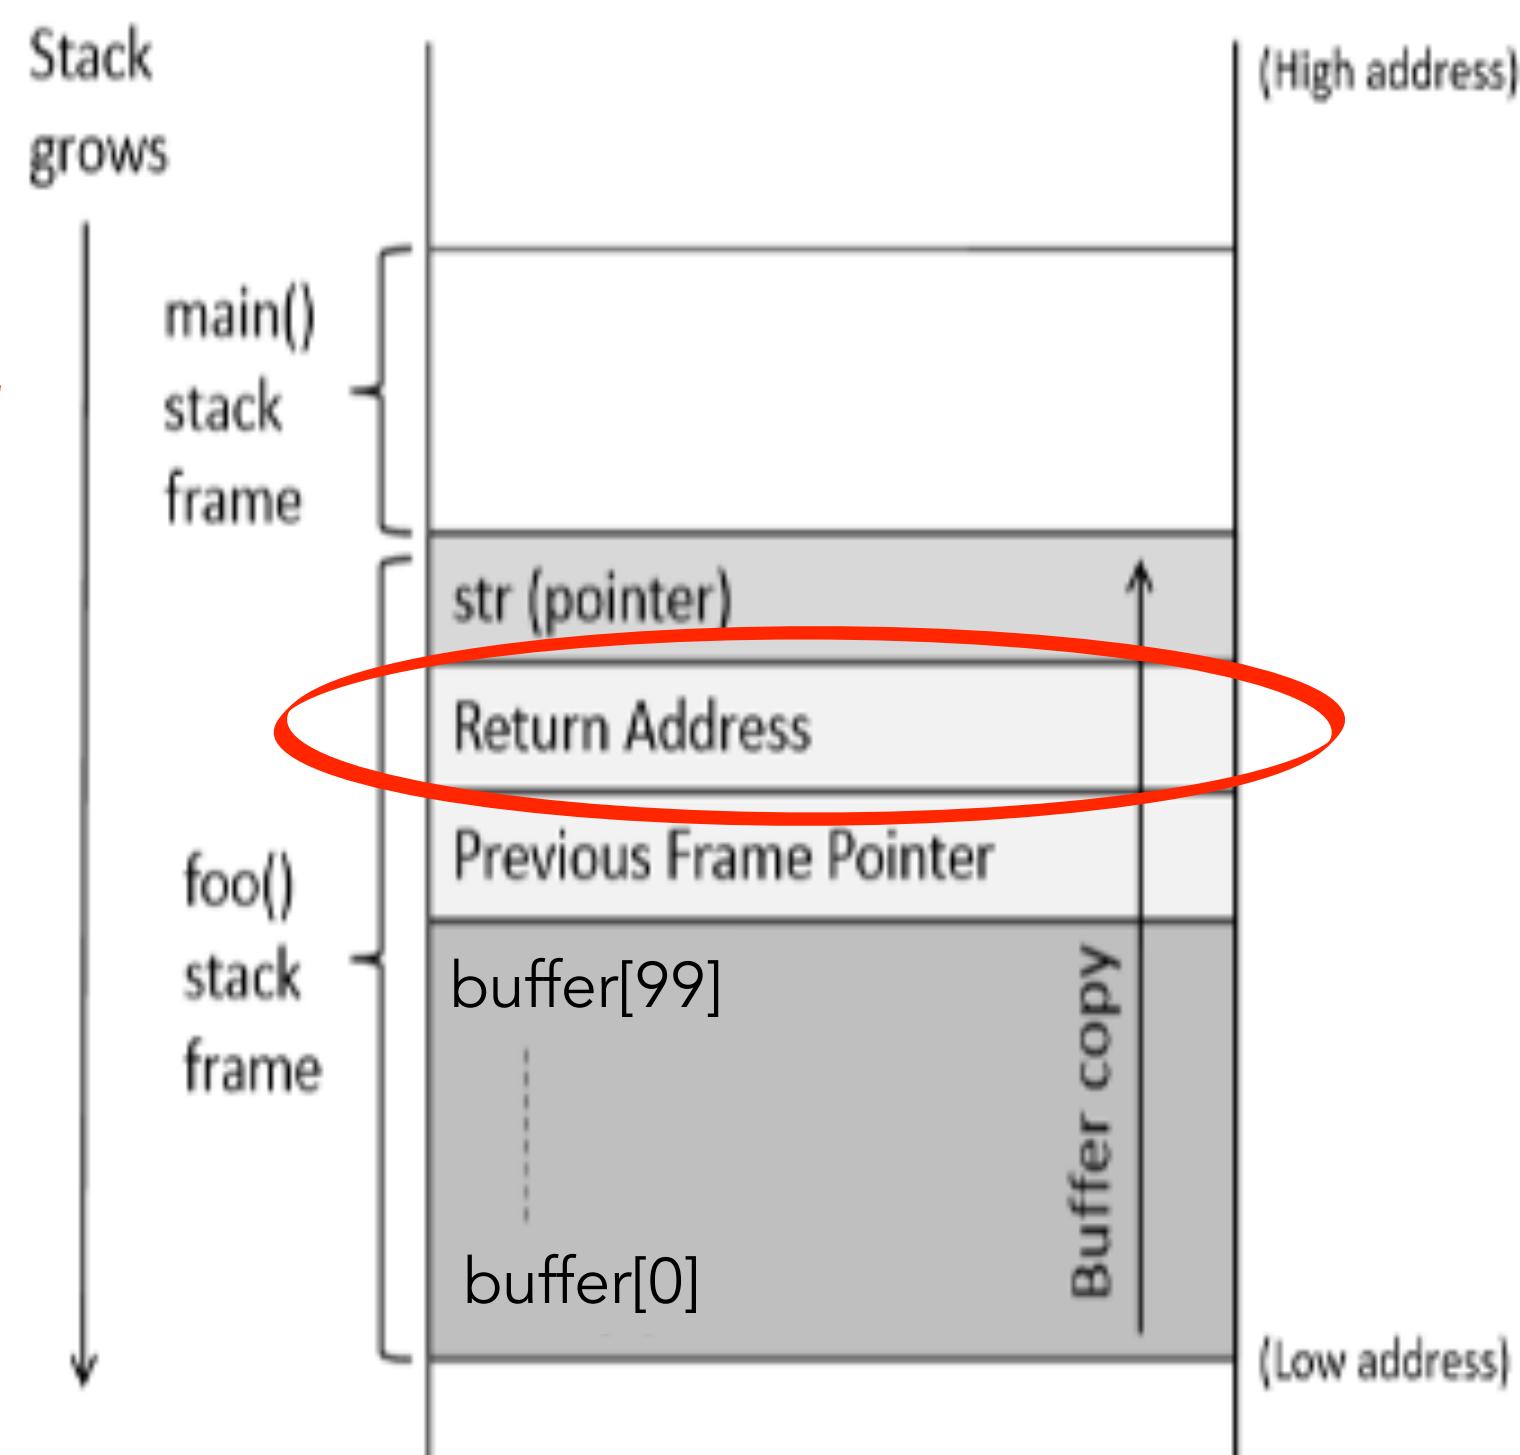
    return 1;
}
```

Consequences of a Buffer Overflow

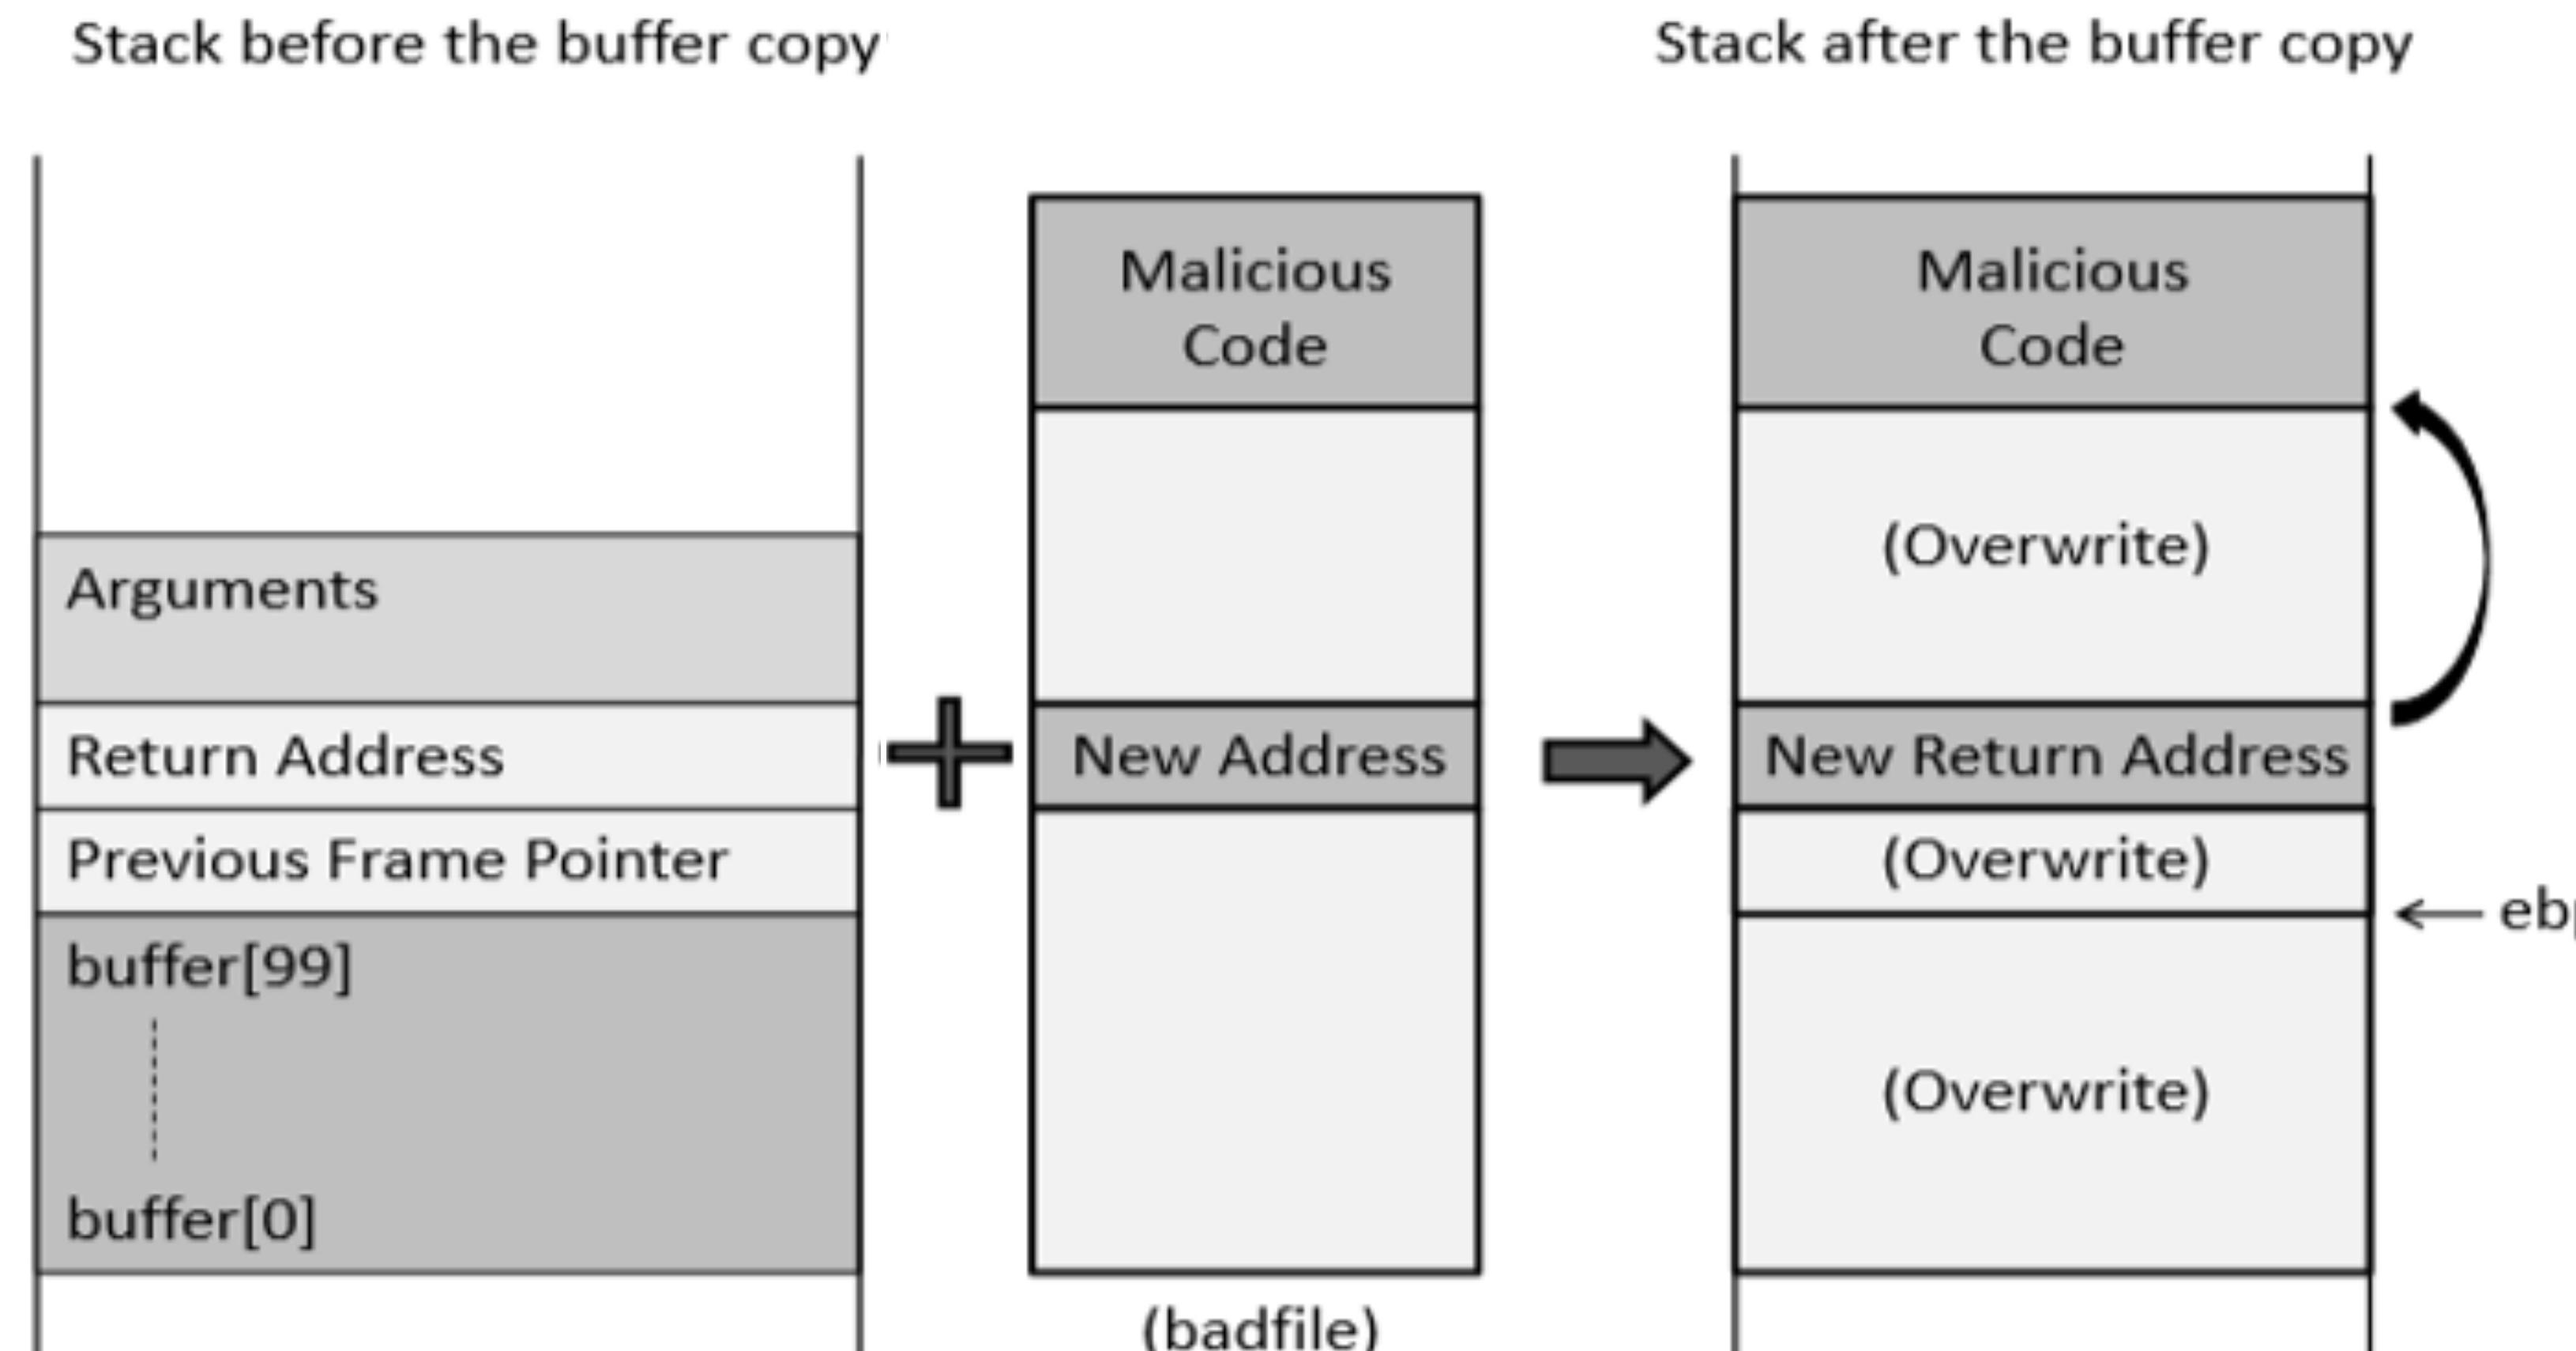
What could really go wrong with a buffer overflow...?

Overwriting the **return address** with some bad data (random address) can lead to:

- Non-existent address >>> access unmapped page; **crash!**
- Access violation >>> access mapped (protected) page (e.g., kernel memory); **crash!**
- Invalid instruction >>> contents of return address is not a valid instruction; **crash!**
- **Execution of attacker's code!**
>>> contents of return address is a valid instruction; program will continue running w/ different logic!



How to Run Malicious Code



OK, so HOW do we carry out the buffer overflow attack?!?!?!!?!



First... Environment Setup

- Turn off **address randomization** (countermeasure)

```
$ sudo sysctl -w kernel.randomize_va_space=0
```

- Compile a set-uid root version of stack.c with **executable stack permitted + no stack guard**

```
$ gcc -o stack -z execstack -fno-stack-protector stack.c
$ sudo chown root stack
$ sudo chmod 4755 stack
```

*Don't worry, we will look at these countermeasures later,
and look at how these can be defeated.....*

Creating the Malicious Input (badfile)

Task A: Find the offset distance between the base of the buffer and the return address.

Task B: Find the address to place the shellcode

